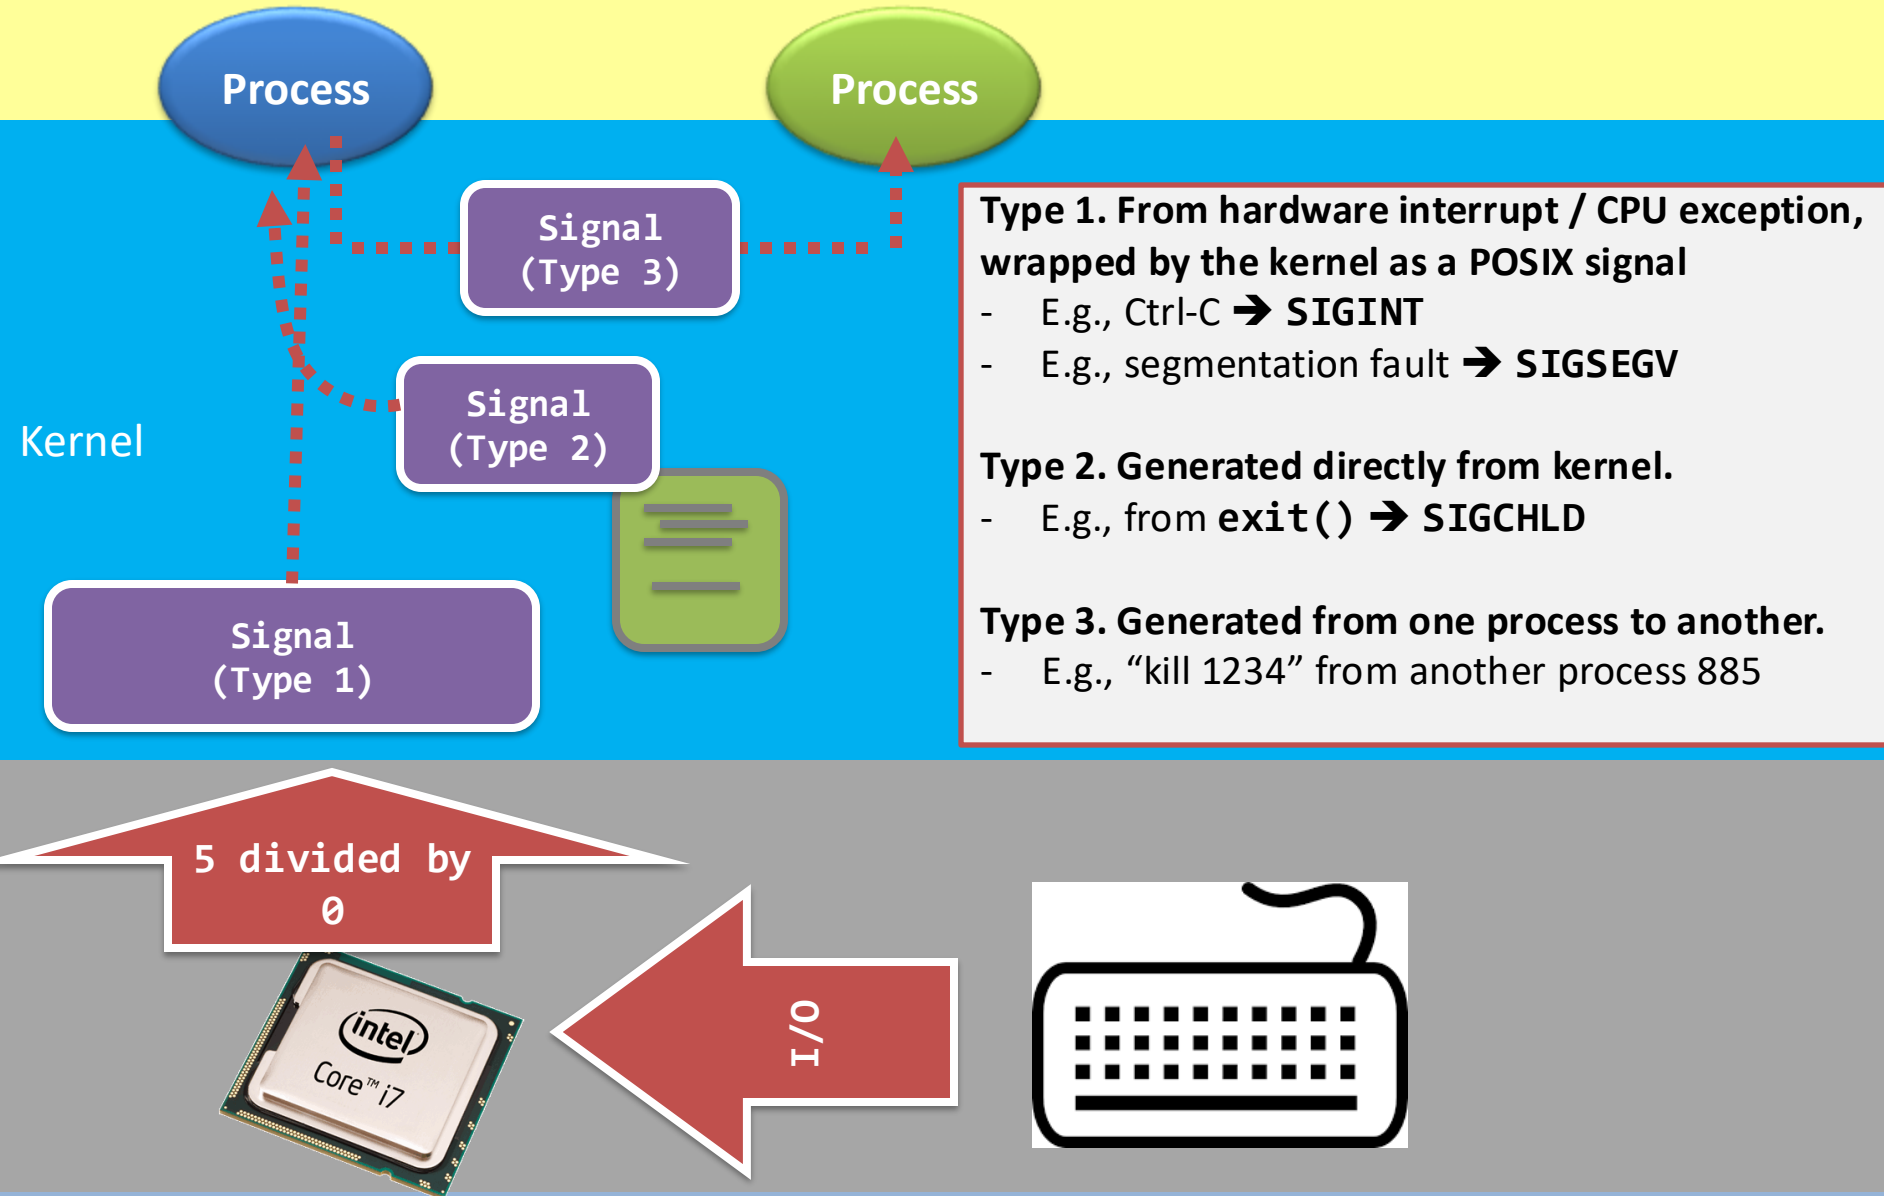


# Operating Systems

Eric Lo

5 Signal

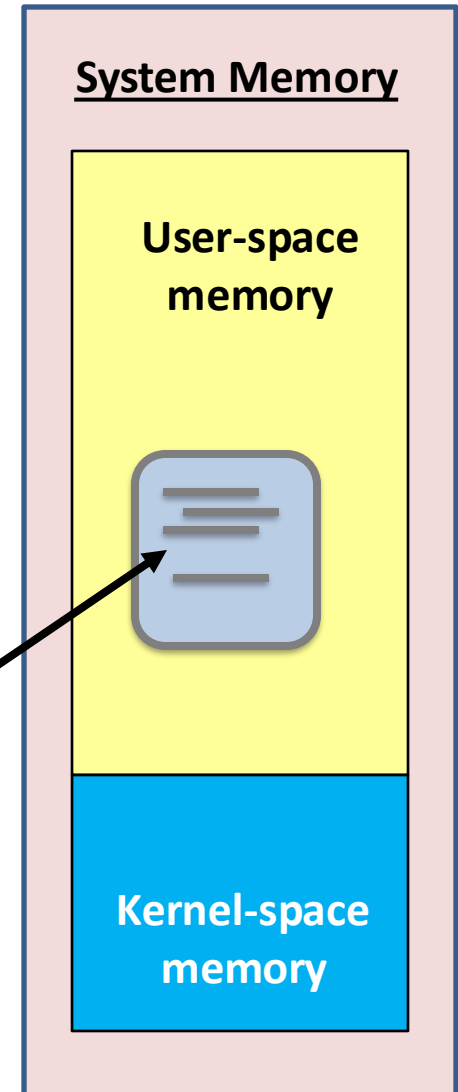
# Signal – a limited form of communications



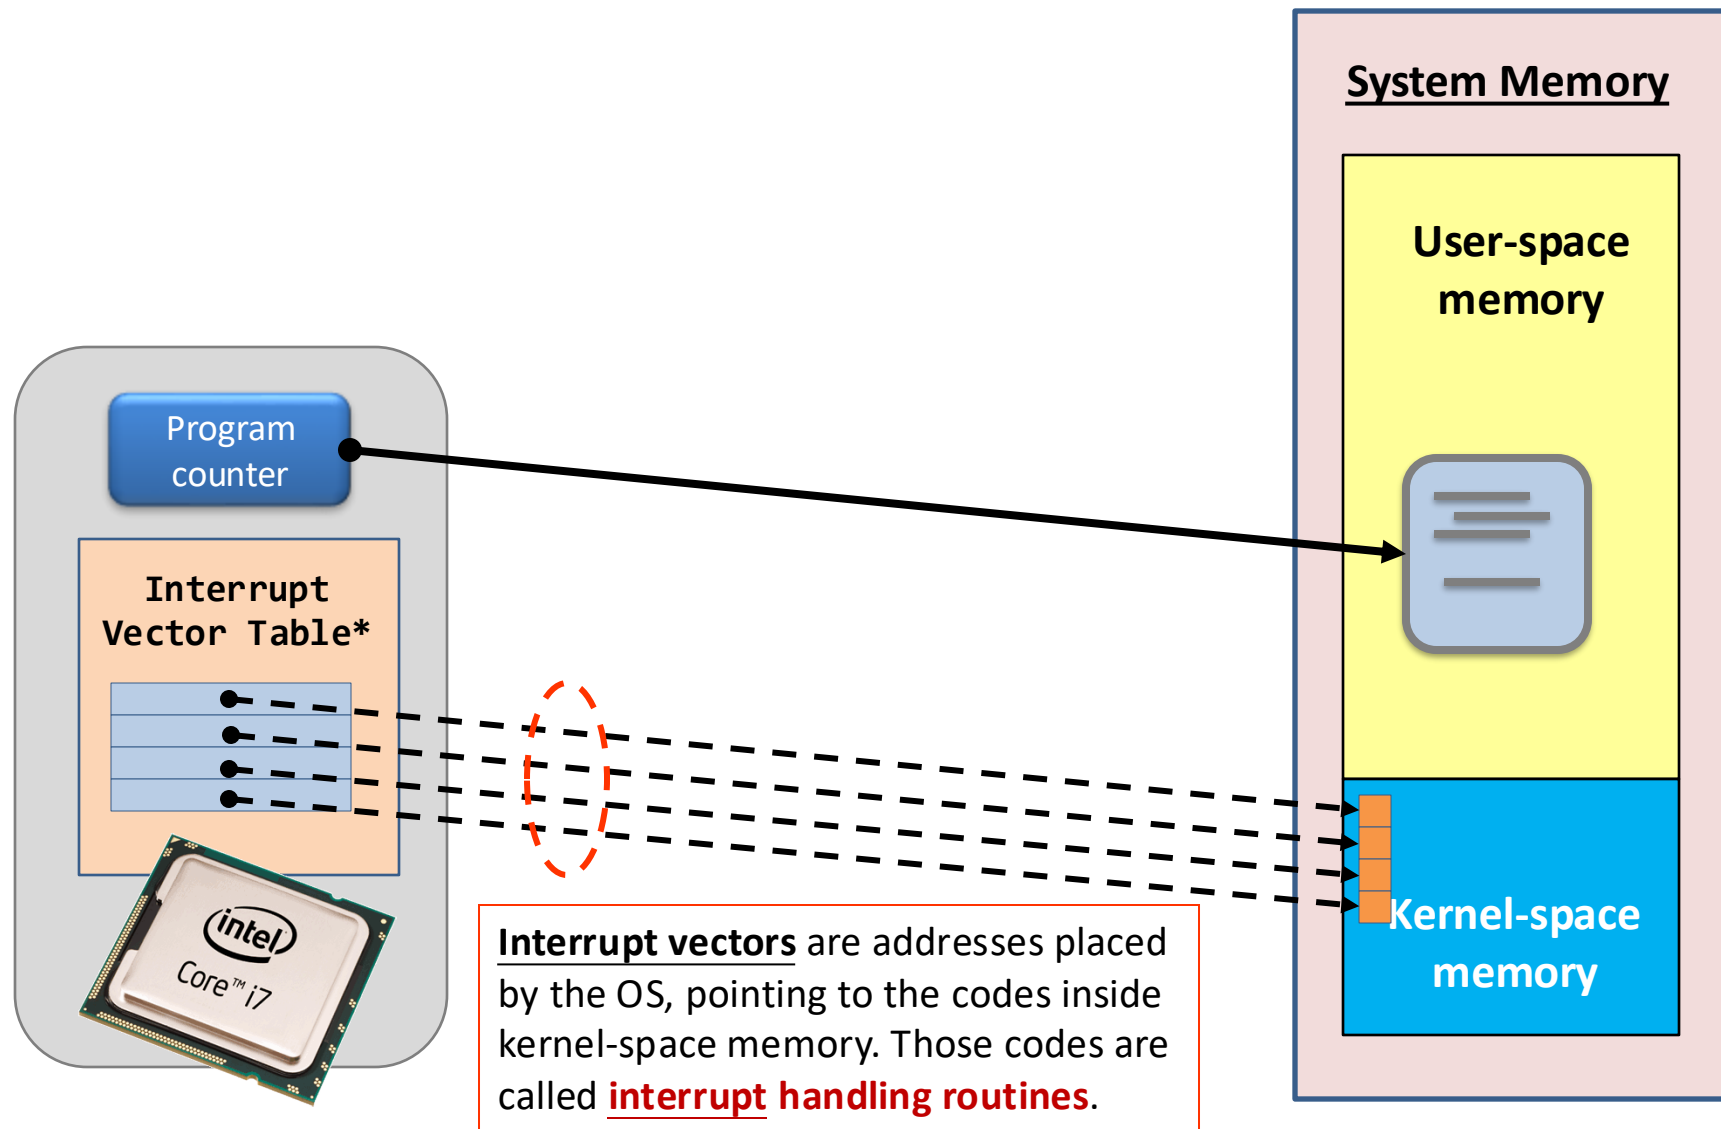
# Hardware Interrupt

- External events that require CPU's attention

Originally, the CPU is working on a program code.

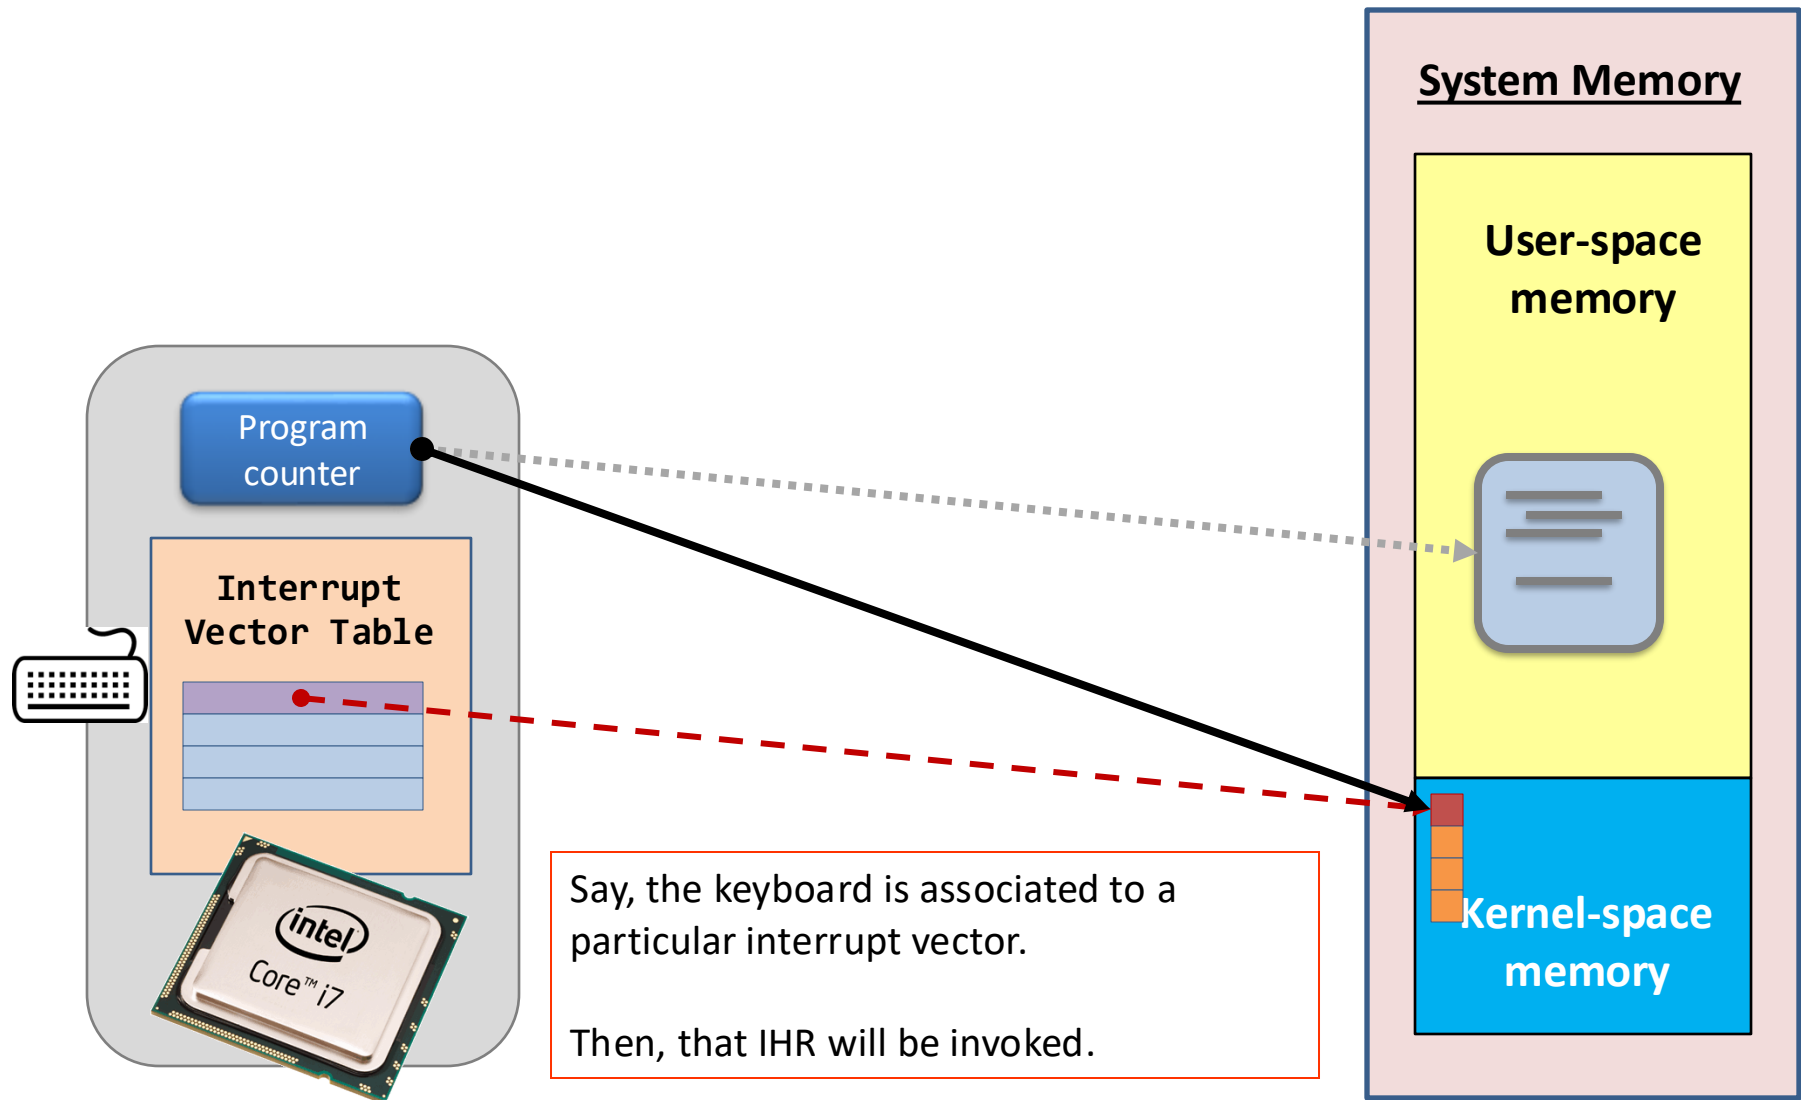


# Interrupt Vector Table



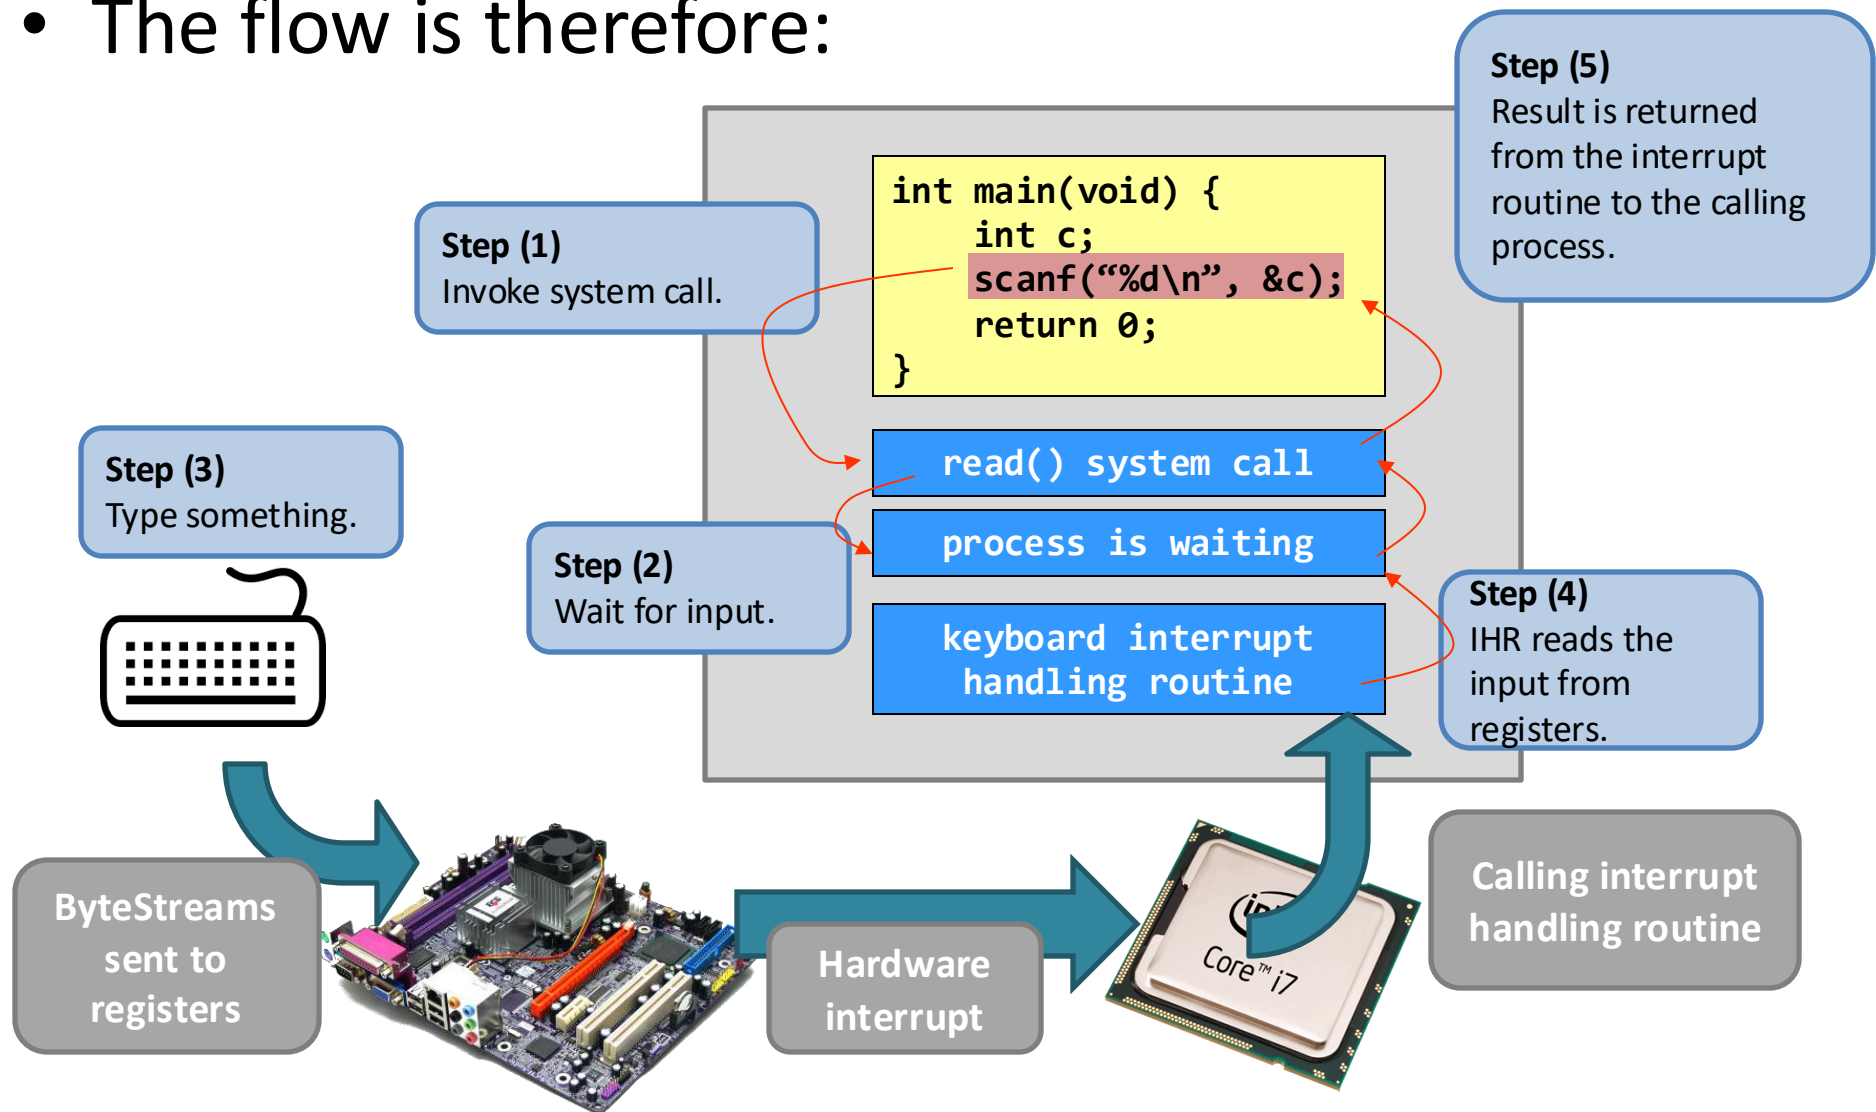
\*IVT's location is architecture-specific: maybe hardcoded in processor/designated space in the kernel-space

# Hardware Interrupt



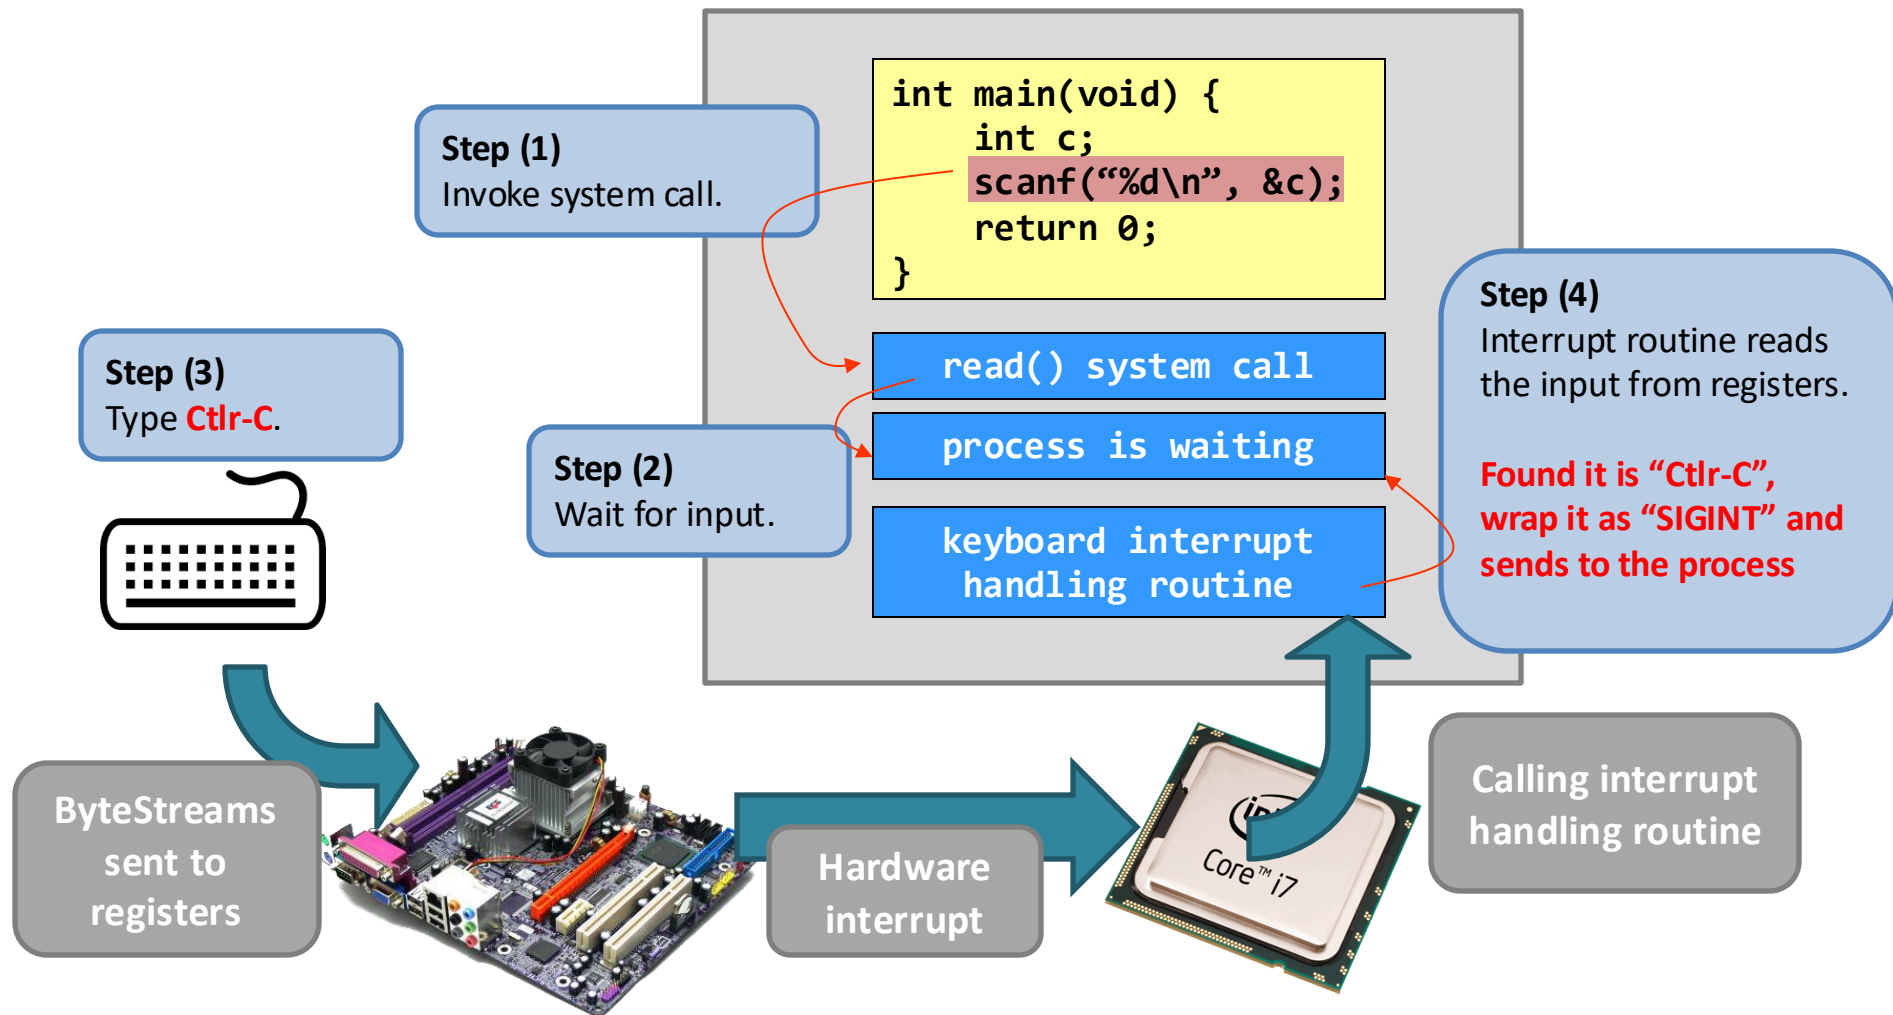
# E.g., scanf

- The flow is therefore:



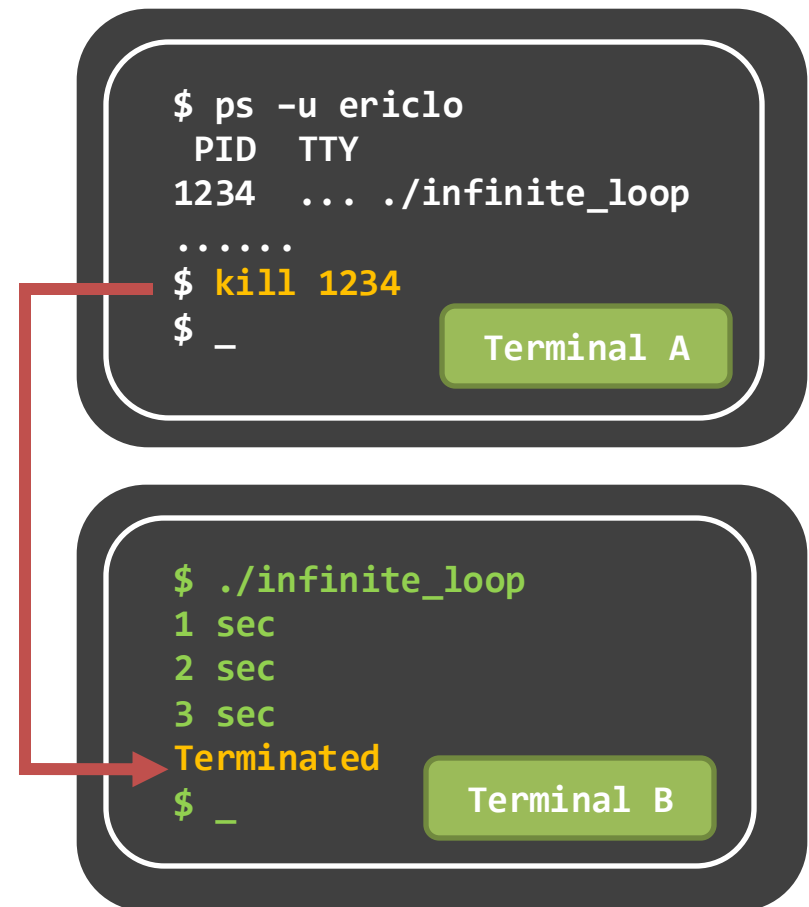
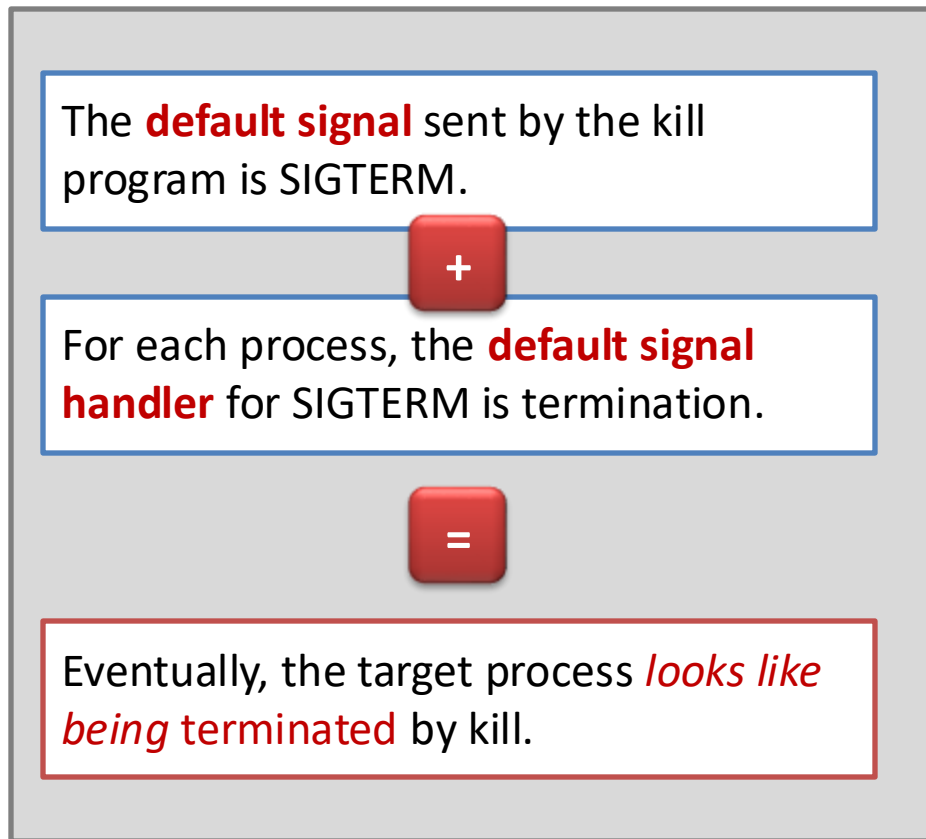
# E.g., Ctrl-C to SIGINT

- The flow is therefore:



# Process2Process: signal sending using “kill” program

- The “kill” program is to **send signals** to target processes.





# POSIX signals

Signal	Description	<u>Default</u> signal handler
<b>SIGINT</b>	Its name is the <b>interrupt signal</b> . Can be generated by <u>“Ctrl + C”</u> .	Target process termination.
<b>SIGTERM</b>	Its name is the <b>termination signal</b> . The default signal sent by the <u>“kill”</u> program. (type 3)	Target process termination.
<b>SIGTSTP</b>	Its name is the <b>terminal stop signal</b> . Can be generated by <u>“Ctrl + Z”</u> .	Target process <i>suspension</i> .
<b>SIGCONT</b>	Its name is the <b>continuation signal</b> .	Target process resumes execution if it is previously suspended.
<b>SIGCHLD</b>	(No special name). It is sent to the parent process from a terminated child via <code>exit()</code> . (type 2)	{ //do nothing by default }
<b>SIGKILL</b>	Its name is the <b>kill signal</b> . If sent, the process <b>MUST DIE</b> .	Target process termination (can't be replaced by other signal handlers)

# Signal: Synchronous vs. Asynchronous

- Asynchronous signal:
  - The signal received is NOT generated by the process itself
    - So, its arrival time is **usually** not deterministic from the process point of view
  - E.g., External hardware interrupt, another process sends **ctrl-c**
- Synchronous signal:
  - The signal is caused by the process itself
    - So, its arrival time is **usually** deterministic from the process point of view
  - E.g., A certain line leads to **SIGFPE**
  - E.g., A certain line accesses an invalid memory region: **SIGSEGV**

# What are SIGTSTP & SIGCONT?

```
$ ps -u ericlo
```

```
PID  TTY
```

```
1234  ... ./infinite_loop
```

```
.....
```

```
$ kill -TSTP 1234
```

```
$ kill -CONT 1234
```

SIGTSTP

SIGCONT

The “terminal stop signal”  
suspends the process.

Shell 2

```
$ ./infinite_loop
```

```
1 sec
```

```
2 sec
```

```
3 sec
```

```
[1] Stopped  ./infinite_loop
```

```
$ 4 sec
```

```
5 sec
```

```
6 sec
```

Try press  
“enter” & then  
“ls”

The “continuation signal”  
resumes the process. But,  
wait....

Shell 1

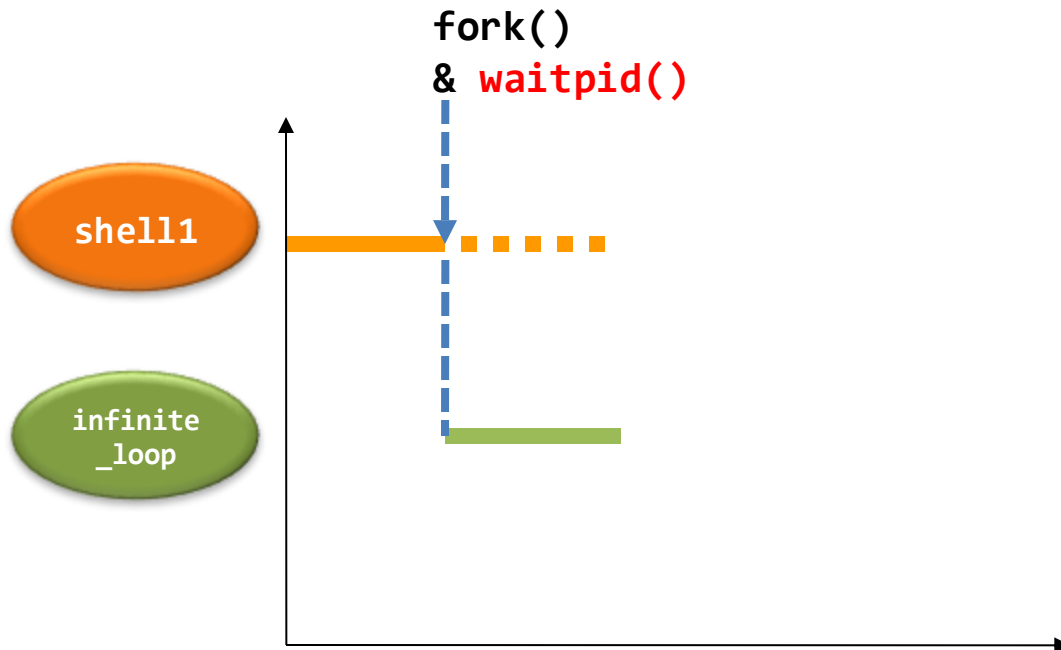
# Foreground and background jobs

A characteristic of a foreground job in a shell is:

**The shell is waiting for the job to change state.**

Foreground job is a concept in a shell:

- read input from stdin
- write output to stdout



```
$ ./infinite_loop
```

```
1 sec
```

```
2 sec
```

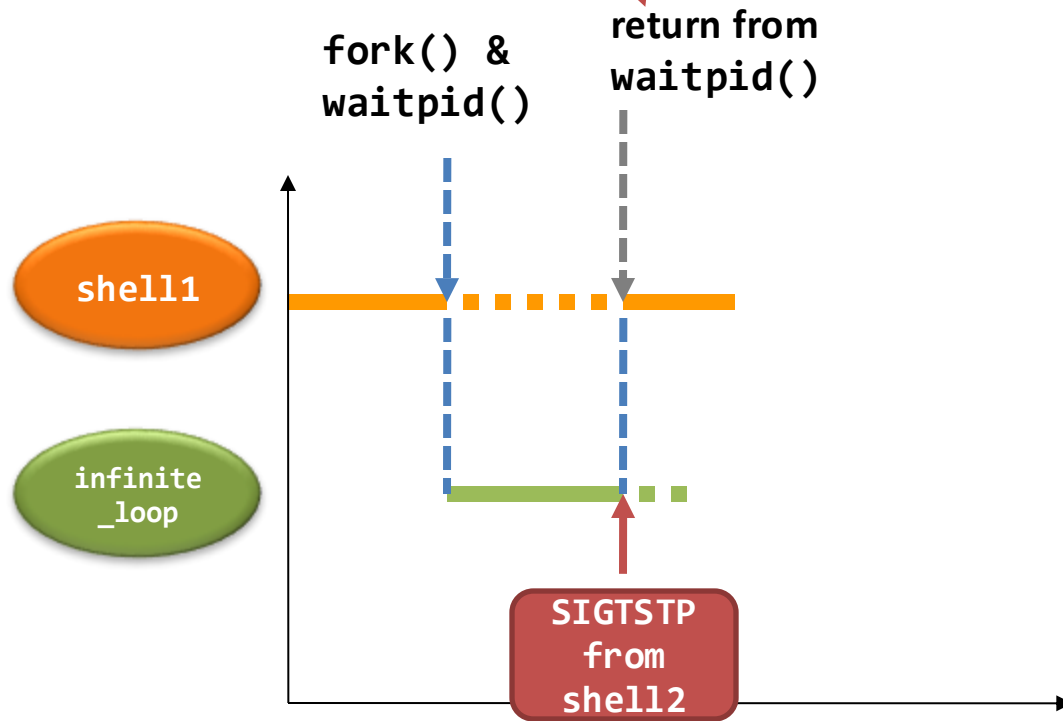
```
3 sec
```

We name this a  
foreground job.

# Foreground and background jobs

This line is actually  
printed by the shell

Now, the shell / parent **wakes** up.



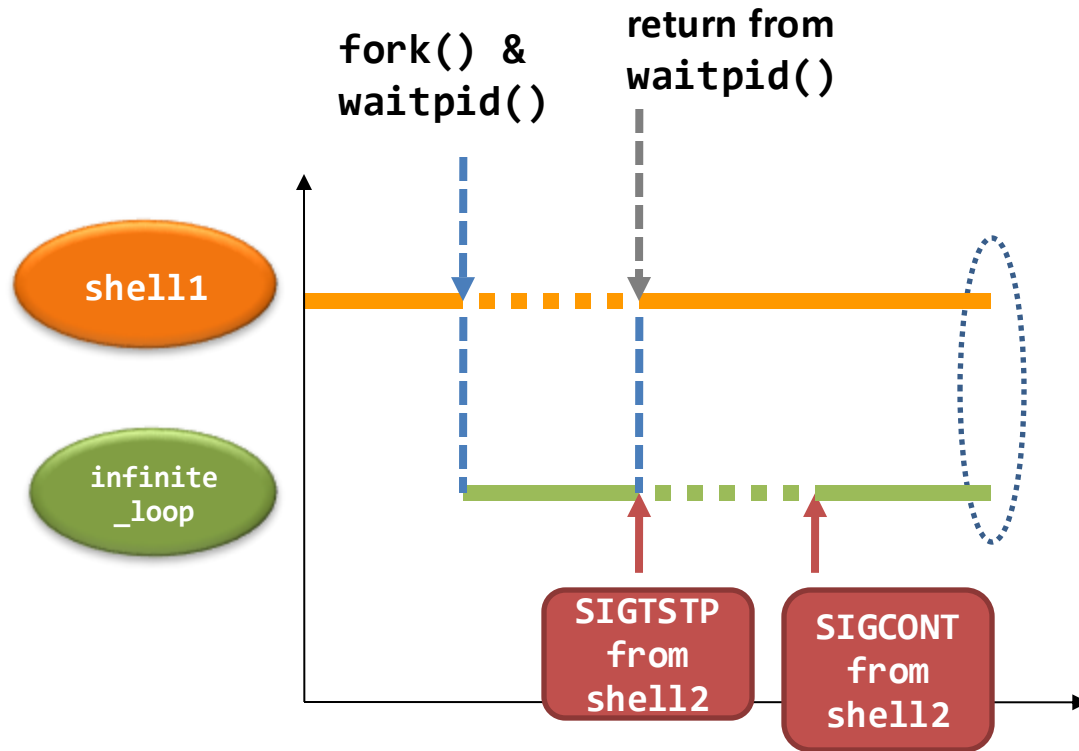
```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ...
$ _
```

# Foreground and background jobs

**SIGCONT** is fired from another shell process, say shell2 to `infinite_loop`

=> so shell1 doesn't know about that,

=> shell1 and `infinite_loop` just run in parallel.



```
$ ./infinite_loop
```

```
1 sec
```

```
2 sec
```

```
3 sec
```

```
[1] Stopped ...
```

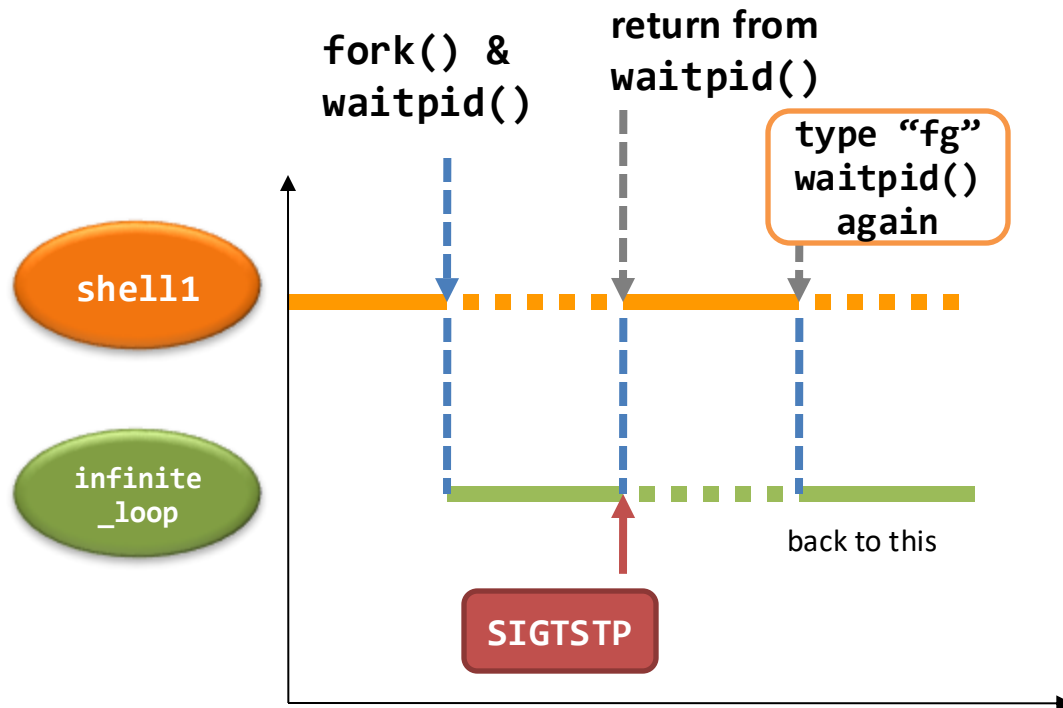
```
$ 4 sec
```

```
5 sec
```

```
6 sec
```

`infinite_loop` inherits parent's (`shell1`) file so outputs to parent's stdout

# Foreground and background jobs



```
$ ./infinite_loop
1 sec
2 sec
3 sec
[1] Stopped ...
$ fg
4 sec
5 sec
6 sec
```

Now, command **fg** turns its most recent suspended/background job (**infinite\_loop**) back as a foreground job

That means shell1 is re-wait() again

# Sending signals in a process

- Remember, **kill()** is intended to send signals.
- So, why not use the name **signal()**?
  - Answer: because **signal()** has already been used for registering a new signal handler

```
int kill(pid_t pid, int sig )
```

PID of  
the target  
process

The desired signal. For the complete  
list, please read the manpages:

On Linux:

**man 7 signal**

On Unix:

**man -s 3HEAD signal**

```
1  int main(void) {
2      int i, sum;
3      srand(time(NULL));
4      while(1) {
5          sum = 0;
6          for(i = 0; i < 3; i++)
7              sum += (rand() % 6) + 1;
8          if(sum == 18)
9              kill(getpid(), SIGTERM);
10     }
11     return 0;
12 }
```

An alternative: **raise(SIGTERM);**

[examples@3150] cat suicide.c



# signal()^

- Register a new signal handler
  - The process no longer executes the default handler...
  - A signal is handled by a user-defined/given function

```
signal( int signum, void (* function)(int) )
```

The signal number, e.g., **SIGINT** or **SIGTERM**.

A function pointer to the new handler. This means your signal handler shall have:

- return type of “**void**” and
- one argument with type “**int**”

Two special values:

- **SIG\_DFL**: set to the default handler.
- **SIG\_IGN**: ignore this signal completely.

# Registering new signal handlers

```
1 void sig_handler(int sig) {  
2     if(sig == SIGINT) ←  
3         printf("\nCtrl + C\n");  
4 }  
5  
6 int main(void) {  
7     signal(SIGINT, sig_handler);  
8     printf("Press enter\n");  
9     getchar();  
10    printf("End of program\n");  
11 }
```

Why put an 'if' here if we know it "always" handles SIGINT here?  
Ans: yes, you are correct.  
It isn't necessary here. But, check out handle\_int.c to find the reason.

Line 7 registers the signal handler for **SIGINT**.

Lines 1-4 together define the customized signal handler.

[examples@3150] cat handle\_int.c

# Registering new signal handlers

```
1 void sig_handler(int sig) {  
2     if(sig == SIGINT)  
3         printf("\nCtrl + C\n");  
4 }  
5  
6 int main(void) {  
7     signal(SIGINT, sig_handler);  
8     printf("Press enter\n");  
9     getchar();  
10    printf("End of program\n");  
11 }
```

```
$ ./handle_int  
Press enter  
^C  
Ctrl + C
```

```
[examples@3150] cat handle_int.c
```

# Registering new signal handlers

- IMPORTANT:
  - Apparently, when a signal handler returns, **the process should go back to where it was executing.**
  - But...

```
1 void sig_handler(int sig) {  
2     printf("\nSignal received.\n");  
3 }  
4  
5 int main(void) {  
6     signal(SIGINT, sig_handler);  
7     printf("Sleep for 24 hours\n");  
8     sleep(24 * 60 * 60);  
9     printf("Wake up and die.\n");  
10 }
```

```
$ ./break_sleep  
Sleep for 24 hours  
^C  
Signal received.  
Wake up and die  
$_
```



[examples@3150] cat break\_sleep.c

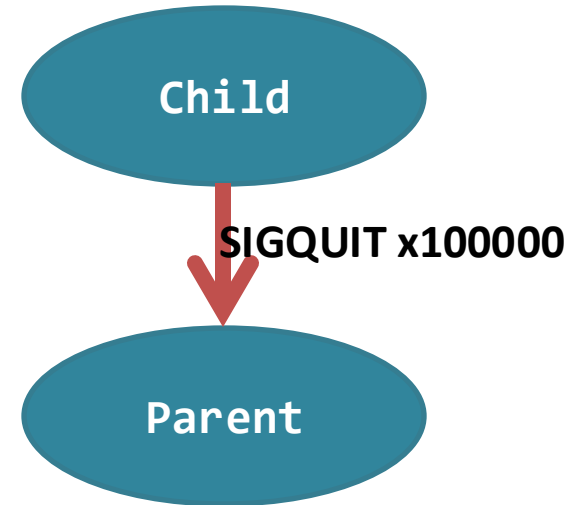
# Registering new signal handlers

- IMPORTANT:
  - Apparently, when a signal handler returns, **the process should go back to where it was executing.**
  - But...this only happens when the involved system/library call can be resumed/restarted.

Can be resumed/restarted	Cannot be resumed/restarted
<p>[file related] <code>open(), read(), write();</code></p> <p>[process related] <code>wait(), waitpid();</code></p>	<p><u><code>sleep()</code></u>; <code>pause();</code></p> <p>With dozens of calls that you may not meet before...</p>

# How to implement “Signal” indeed?

```
1 void handler(int sig) {
2     static int count = 0;
3     printf("count = %d\n", ++count);
4 }
5
6 int main(void) {
7     int i;
8     if( fork() == 0 ) {
9         printf("Press Enter...\n");
10        while(getchar() != '\n');
11        for(i = 0; i < 100000; i++)
12            kill(getppid(), SIGQUIT);
13    }
14    else {
15        signal(SIGQUIT, handler);
16        sleep(1000);
17        wait(NULL); //wait until Child terminates
18    }
19    return 0;
20 }
```



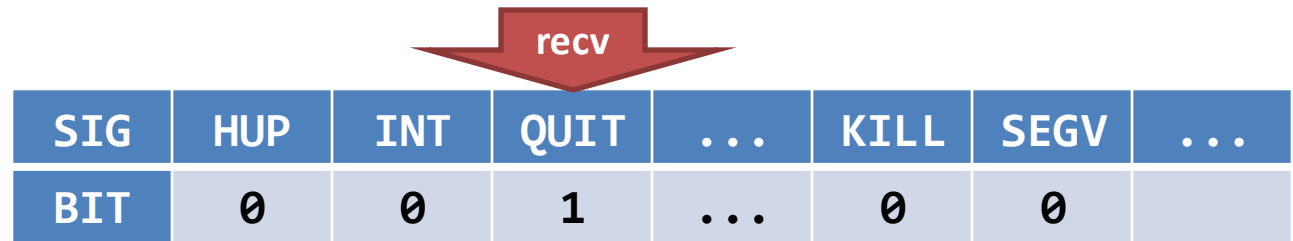
```
$ ./many_signal
Press Enter...
[ENTER]
count = 1
$_
```



[examples@3150] cat many\_signal.c

# How to implement “Signal” indeed?

- Standard signals are not implemented as a queue
- A bit array is used for indicating a signal has received or not.



SIG	HUP	INT	QUIT	...	KILL	SEGV	...
BIT	0	0	1	...	0	0	

- In the previous example, the bit (or **mask**) is set to 1 once **SIGQUIT** is received.
  - The mask will be set to 0 once the signal is handled.

# Signal Implementations

- Check your knowledge, read this:
  - <http://www.linuxjournal.com/article/3985?page=0,0>
  - In the bottom, you'll read 2 kernel data structures:
    - current-signal – *“contains a bitmask of pending signals”*
    - sigqueue – *“double-linked list of pending signals”*
- Tip:
  - Read man 7 signal
    - <http://man7.org/linux/man-pages/man7/signal.7.html>



## Misc. Topics

- Waiting for signals;
- Cleanup when facing Ctrl-C;
- Timers and periodic signals;

# (1) - Waiting for a signal

- The **pause()** system call suspends the calling process until a signal is received.
- Vs. **wait()** :
  - **wait()** doesn't mean "I wait"
  - **wait()** means "I wait **for** ..."
  - It suspends only when it has a child process

It suspends the execution of the program until a signal is caught...

```
1 void sig_handler(int sig) {  
2 }  
3  
4 int main(void) {  
5     signal(SIGINT, sig_handler);  
6     pause();  
7     printf("Ctrl+C received. Bye!\n");  
8     return 0;  
9 }
```



## (2) – Cleanup when facing **Ctrl+C**

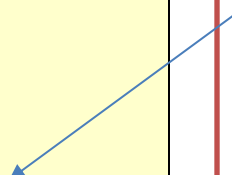
```
int stop = 0;

void sig_handler(int sig) {
    stop = 1;
}

int main(void) {
    unsigned int i = 0;
    signal(SIGINT, sig_handler);

    while( !stop ) {
        sleep(1);
        printf("%d sec\n", ++i);
        fflush(stdout);
    }

    printf("Exit peacefully\n");
    return 0;
}
```



“**Ctrl+C**” terminates a program too ugly.

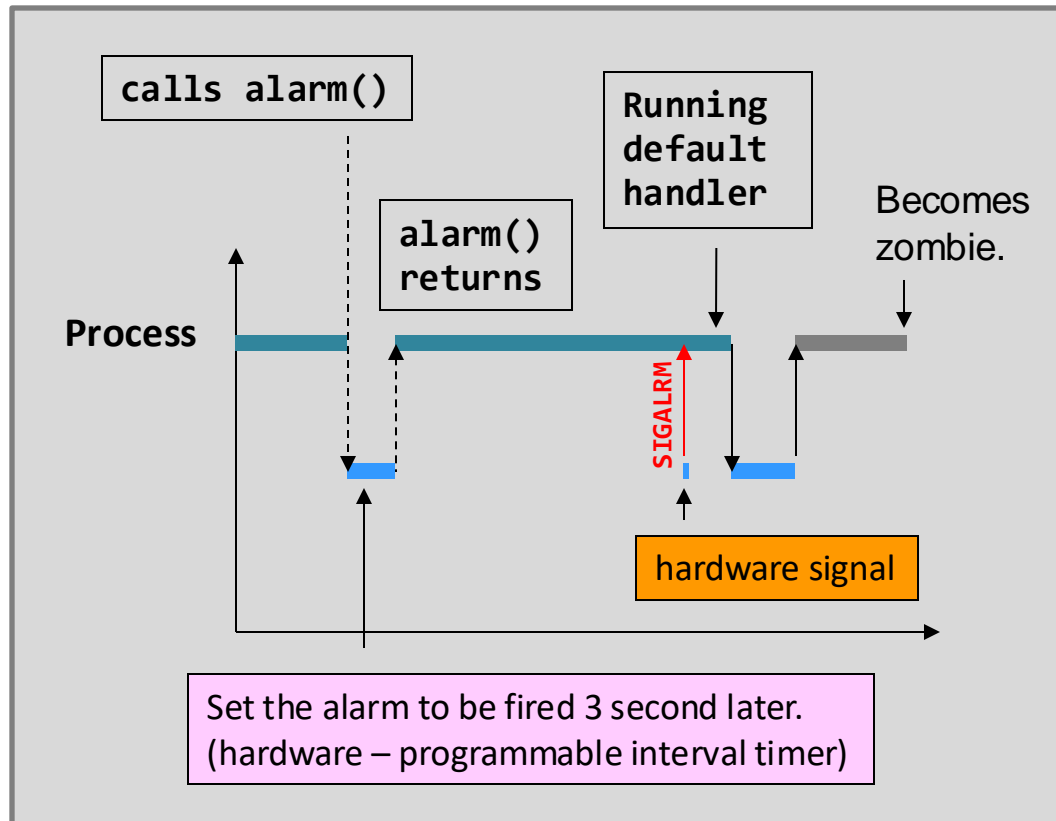
It allows the program to **exit gracefully when facing “Ctrl+C”**, e.g., you can:

- close network connections,
- commit database changes,
- etc.

[examples@3150] cat infinite\_loop.c

### (3) – Timer and alarm()

- **alarm()** is a system call that allows **asynchronous timing** for a process.



```
int main(void) {  
    alarm(3);  
    while(1);  
    return 0;  
}
```

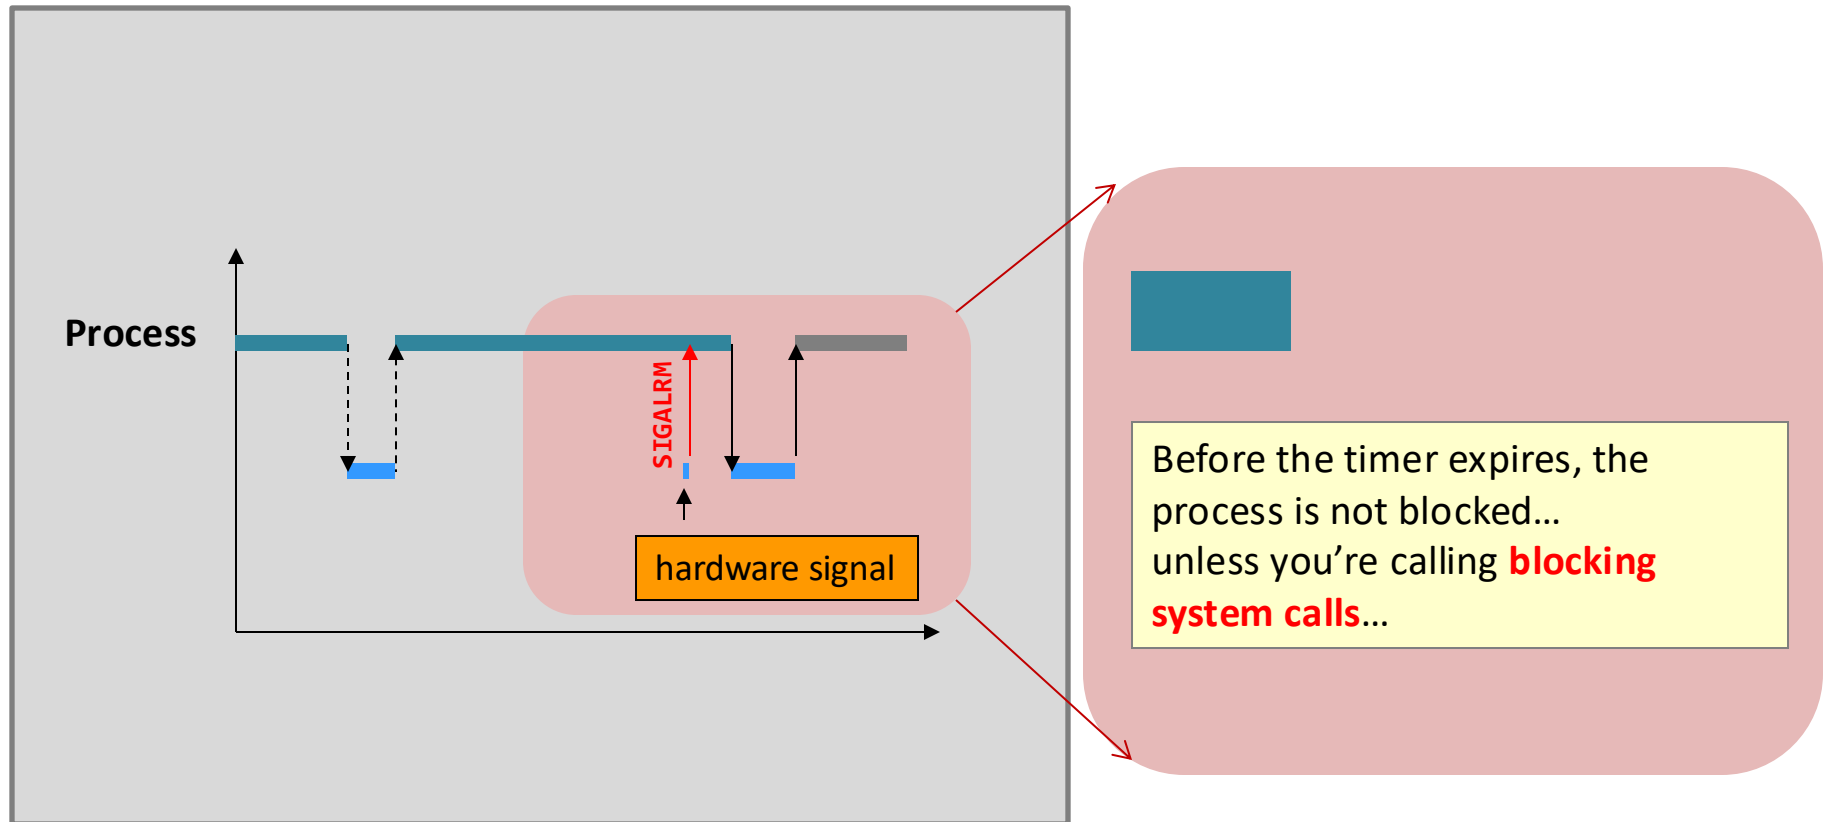
```
$ ./alarm  
Alarm clock  
$ _
```

3 sec later

[examples@3150] cat alarm.c

### (3) – Timer and alarm()

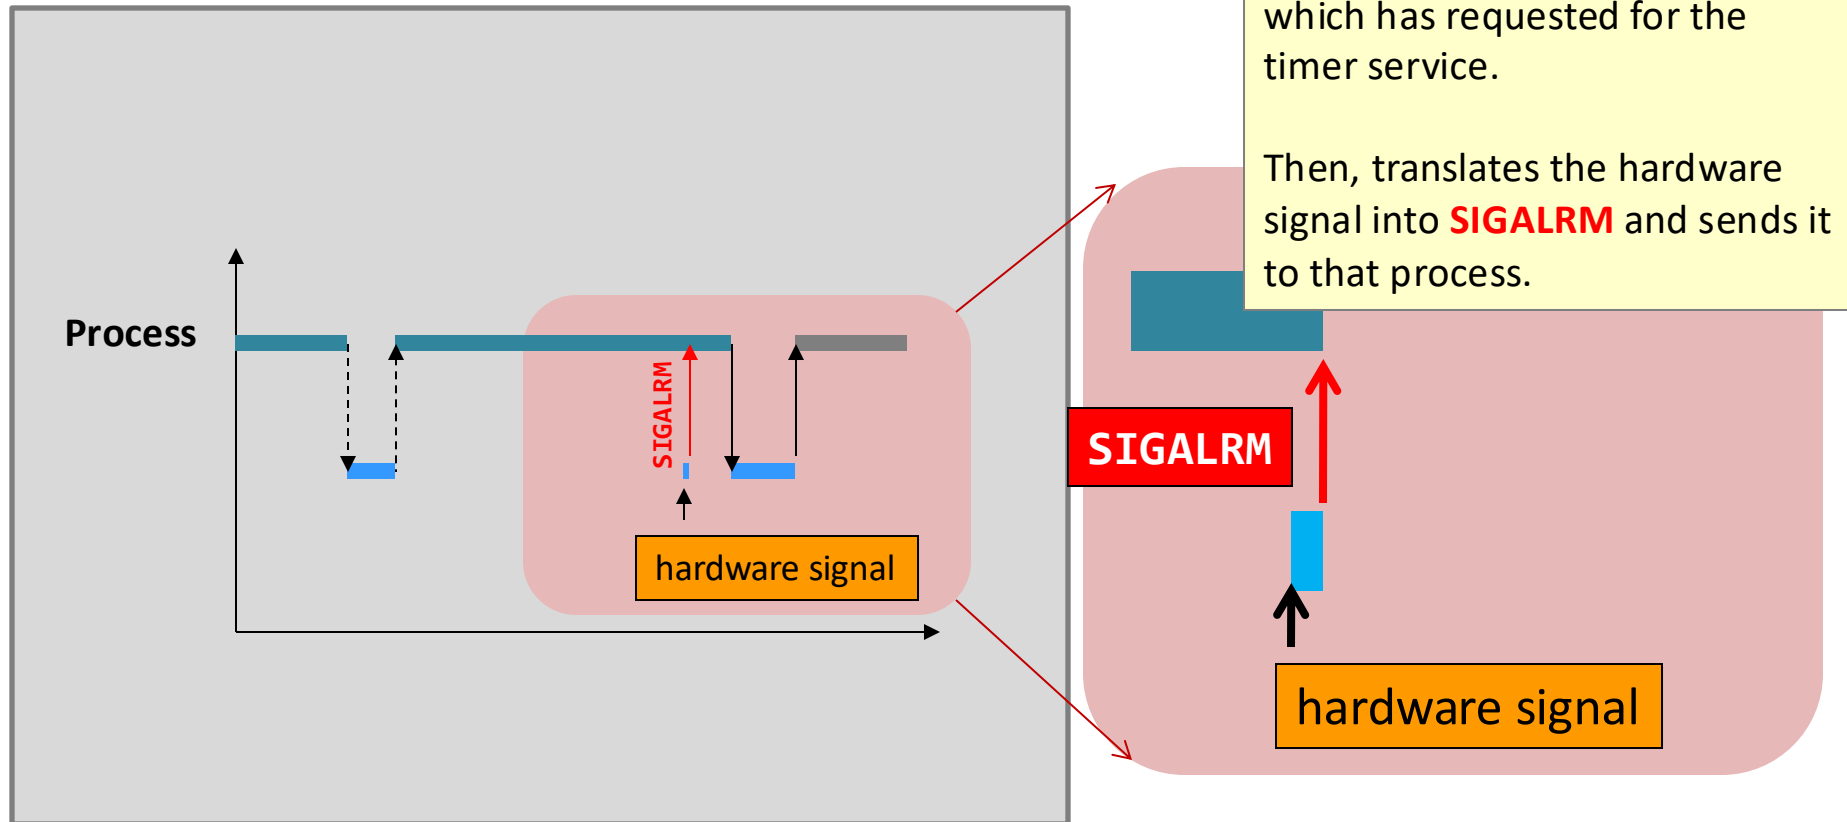
- **alarm()** is a system call that allows **asynchronous timing** for a process.



```
[examples@3150] cat alarm.c
```

### (3) – Timer and alarm()

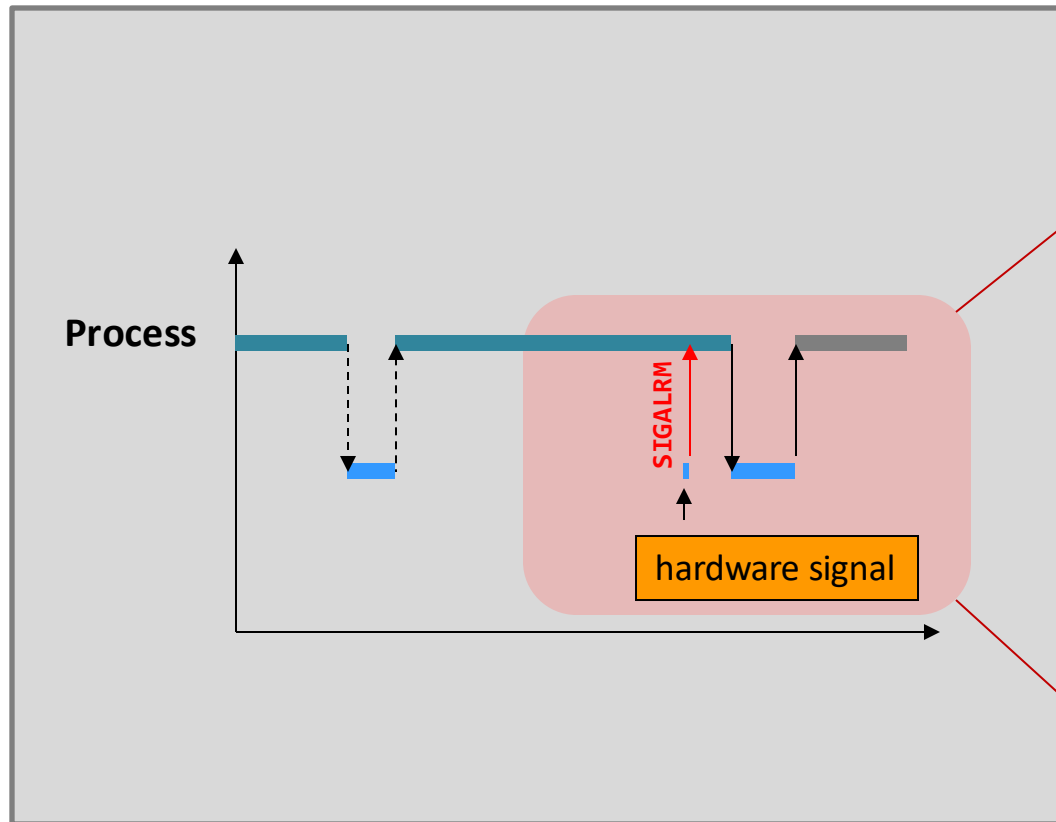
- **alarm()** is a system call that allows **asynchronous timing** for a process.



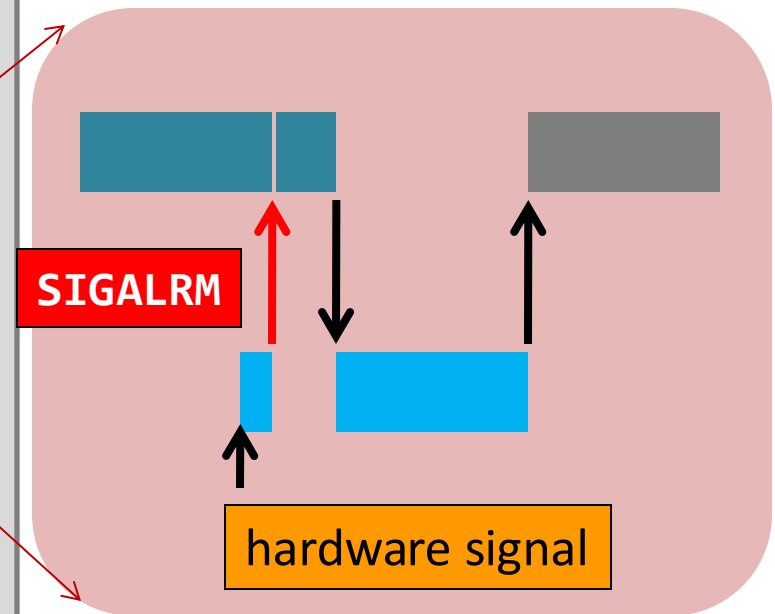
```
[examples@3150] cat alarm.c
```

### (3) – Timer and alarm()

- **alarm()** is a system call that allows **asynchronous timing** for a process.



Upon the arrival of the signal, the default signal handling routine is called, which is **termination**!



[examples@3150] cat alarm.c

### (3) – Timer and alarm()

- Of course, you can implement something more meaningful.

Guess: what will happen with this `exit()` call.

Listen! You've only 5 seconds to finish your typing!

This **cancels** the scheduled clock interrupt!

```
void sig_handler(int sig) {
    printf("\nTimeout! Goodbye!\n");
    exit(0);
}

int main(void) {
    char buf[1024];
    signal(SIGALRM, sig_handler);
    alarm(5);
    if(fgets(buf, 1024, stdin) == NULL) {
        printf("No input. Goodbye!\n");
        exit(0);
    }
    alarm(0);
    printf("Your input: %s", buf);
}
```

[examples@3150] cat alarm\_fgets.c



## (3) – Timer and alarm()

- Remember, “**alarm()**” only fires once!
  - What if I want **periodic signals**?
  - “**setitimer()**” (set interval timer) can help you.
    - Her sibling is “**getitimer()**”.
    - Read the manpage by yourself.

# Further reading

- *Advanced Programming Environment in UNIX;*
- “**man signal**” is a vast resource.

# When will a process check for an arrival of signal?

- A. Decode phase of every instruction
- B. When context switch back to the process
- C. A signal arrival will update the program counter to the signal handler
- D. None of the above



Join at:  
[ahaslides.com/  
0LQ2X](https://ahaslides.com/0LQ2X)