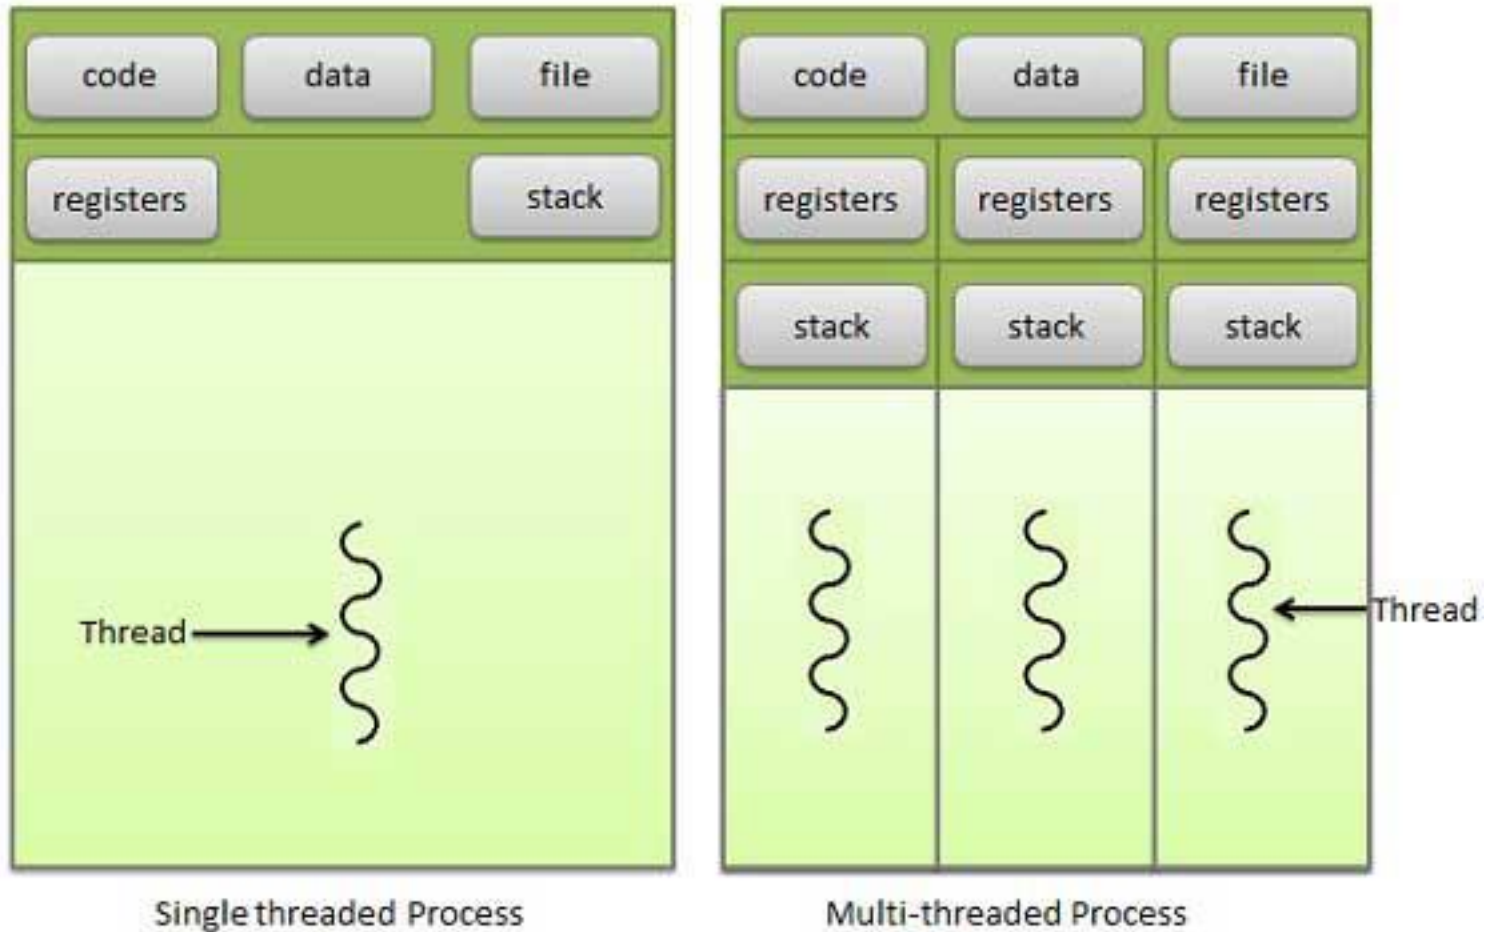


# Operating Systems


Eric Lo

## 7 – Multi-threading

# Single-threaded process vs. multi-threaded

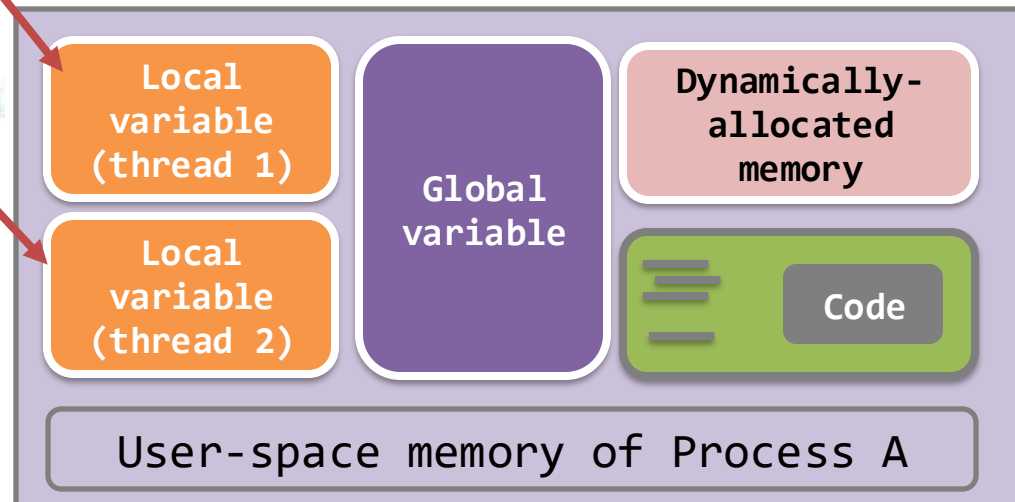


# Multi-threading



Multi-threaded Process

- Threads share the same code
  - From the same program
- Threads share the same address space
  - Threads can then share via global variables
- Stacks of non-parent threads are located somewhere in the memory-map segment



# Process vs. Thread

Process	Thread
Heavy weight	Light weight
Costly creation	Cheap creation
Process can't access the memory area of other processes	Threads share the same memory area
Costly process switching	Cheap thread switching
One process multiple threads	
Different programs	Same program (different thread functions)


While thread “sounds” much better than a process technically

**Don't put the cart before the horse**

You won't imagine the same program “Is” has one thread doing the job of “Is” and another thread doing the job of “less” → they should be 2 different programs

# Multi-thread examples: pros & cons

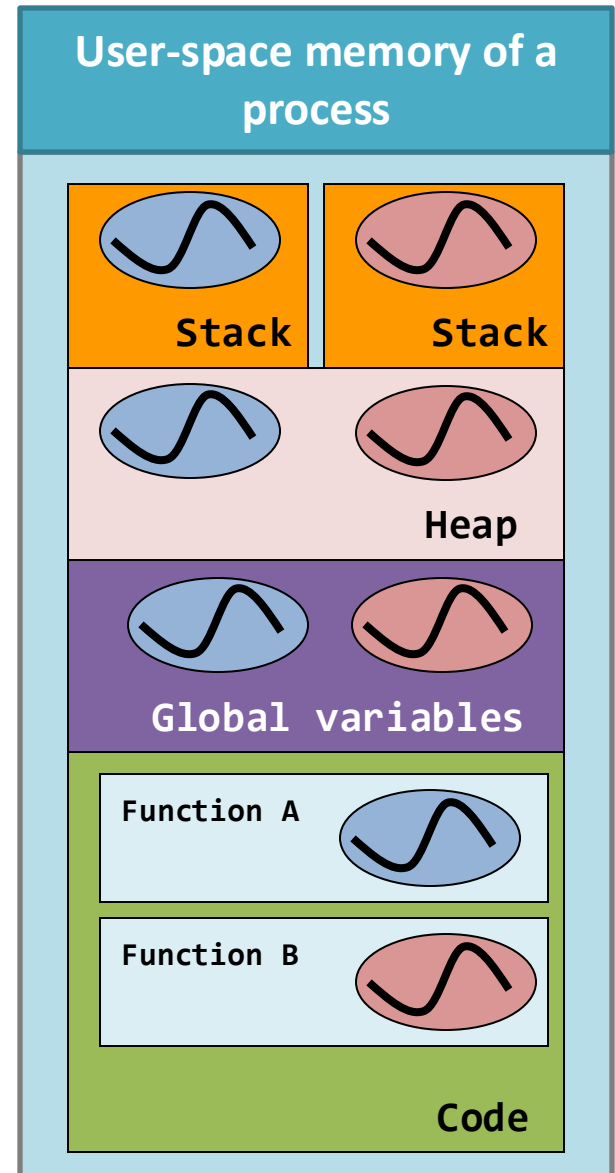
- Old Firefox
  - Single-process
  - Multi-threaded
    - Thread 1: Tab1
    - Thread 2: Tab2
  - Faster
  - Crash one? Crash all!
- Chrome
  - 1 tab 1 process



# Multi-thread – Code and thread function

## Code


- All threads share the same code
- A thread starts with **one specific function**
  - the **thread function**
  - E.g., Functions A & B in the diagram.
- Of course, the thread function can invoke other functions or system calls.
- But, a thread **never returns to the caller of the thread function**



# Multi-thread – Local Variables

## Local variables

- Each thread has its own stack for the local variables.
- i.e., thread function's () {  
    int x;  
}
- Thread 1 updates her x != updating Thread 2's x



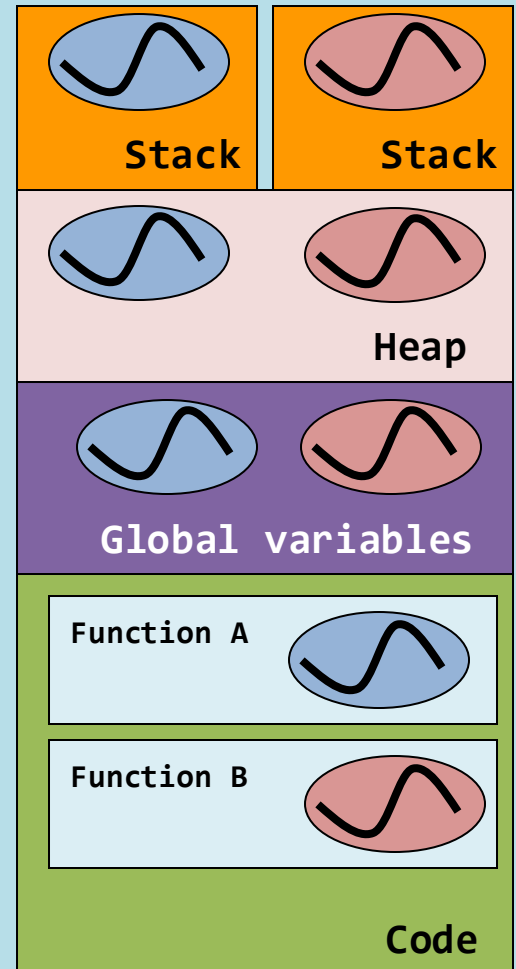
# Multi-thread – Global Variables + Heap

## Heap

## Global variables

- All threads share the global variable zone and the heap
- i.e.,
- `int g;`
- thread function's () {  
    `int x;`  
}
- Thread 1 updates g → Thread 2 see that change

## User-space memory of a process





# Multi-threading

- Introduction.
- **Basic Programming.**



# The Pthread library

- POSIX thread library

	Process	Thread
Creation	<code>fork()</code>	<code>pthread_create()</code>
I.D. Type	PID, an integer	“pthread_t”, a structure
Who am I?	<code>getpid()</code>	<code>pthread_self()</code>
Termination	<code>exit()</code>	<code>pthread_exit()</code>
Wait for child termination	<code>wait()</code> or <code>waitpid()</code>	<code>pthread_join()</code>
Kill	<code>kill()</code>	<code>pthread_kill()</code>

# Thread creation – `pthread_create()`

## Thread Function

```
1 void * hello( void *input ) {  
2     printf("%s\n", (char *) input);  
3     pthread_exit(NULL);  
4 }
```

## Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, "hello world");  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

This sets the thread function of the to-be-created thread as: **hello()**.

[examples@3150] cat pthread\_hello\_world.c

# Thread creation – pthread\_create()


## Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

## Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

At the beginning,  
there is only one  
thread running: **the  
main thread.**



Main Thread

```
[examples@3150] cat pthread_hello_world.c
```

# Thread creation – pthread\_create()

## Thread Function


```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

## Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```


The hello thread is created!

It is running *“together”* with the main thread.



Main Thread

pthread\_create()



Hello Thread

[examples@3150] cat pthread\_hello\_world.c

# Thread creation – `pthread_create()`

## Thread Function


```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

## Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```


Remember `wait()`  
and `waitpid()`?

`pthread_join()`  
performs similarly.



Blocked

Main Thread



Hello Thread

[examples@3150] cat pthread\_hello\_world.c

# Thread creation – pthread\_create()


## Thread Function

```
1 void * hello( void *input ) {  
2     printf(“%s\n”, (char *) input);  
3     pthread_exit(NULL);  
4 }
```

## Main Function

```
5 int main(void) {  
6     pthread_t tid;  
7     pthread_create(&tid, NULL, hello, “hello world”);  
8     pthread_join(tid, NULL);  
9     return 0;  
10 }
```

**Termination of the target thread** causes `pthread_join()` to return.



[examples@3150] cat pthread\_hello\_world.c

# Shared address space

## Thread Function

```
1 void * do_your_job( void *input ) {  
2     printf("child = %d\n", *( (int *) input) );  
3     *((int *) input) = 20;  
4     printf("child = %d\n", *( (int *) input) );  
5     pthread_exit(NULL);  
6 }
```

## Main Function

```
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10;  
10    printf("main = %d\n", input);  
11    pthread_create(&tid, NULL, do_your_job, &input);  
12    pthread_join(tid, NULL);  
13    printf("main = %d\n", input);  
14    return 0;  
15 }
```

```
$ ./pthread_sharing  
main = 10  
child = 10  
child = 20  
main = 20  
$
```

[examples@3150] cat pthread\_sharing.c



# Shared address space

The local variable “**input**” is in the stack of the main thread.

```
1 void * do_your_job( void *input ) {
2     printf("child = %d\n", *( (int *) input) );
3     *((int *) input) = 20;
4     printf("child = %d\n", *( (int *) input) );
5     pthread_exit(NULL);
6 }

7 int main(void) {
8     pthread_t tid;
9     int input = 10;
10    printf("main = %d\n", input);
11    pthread_create(&tid, NULL, do_your_job, &input);
12    pthread_join(tid, NULL);
13    printf("main = %d\n", input);
13    return 0;
14 }
```

Local  
(main thread)

Dynamic

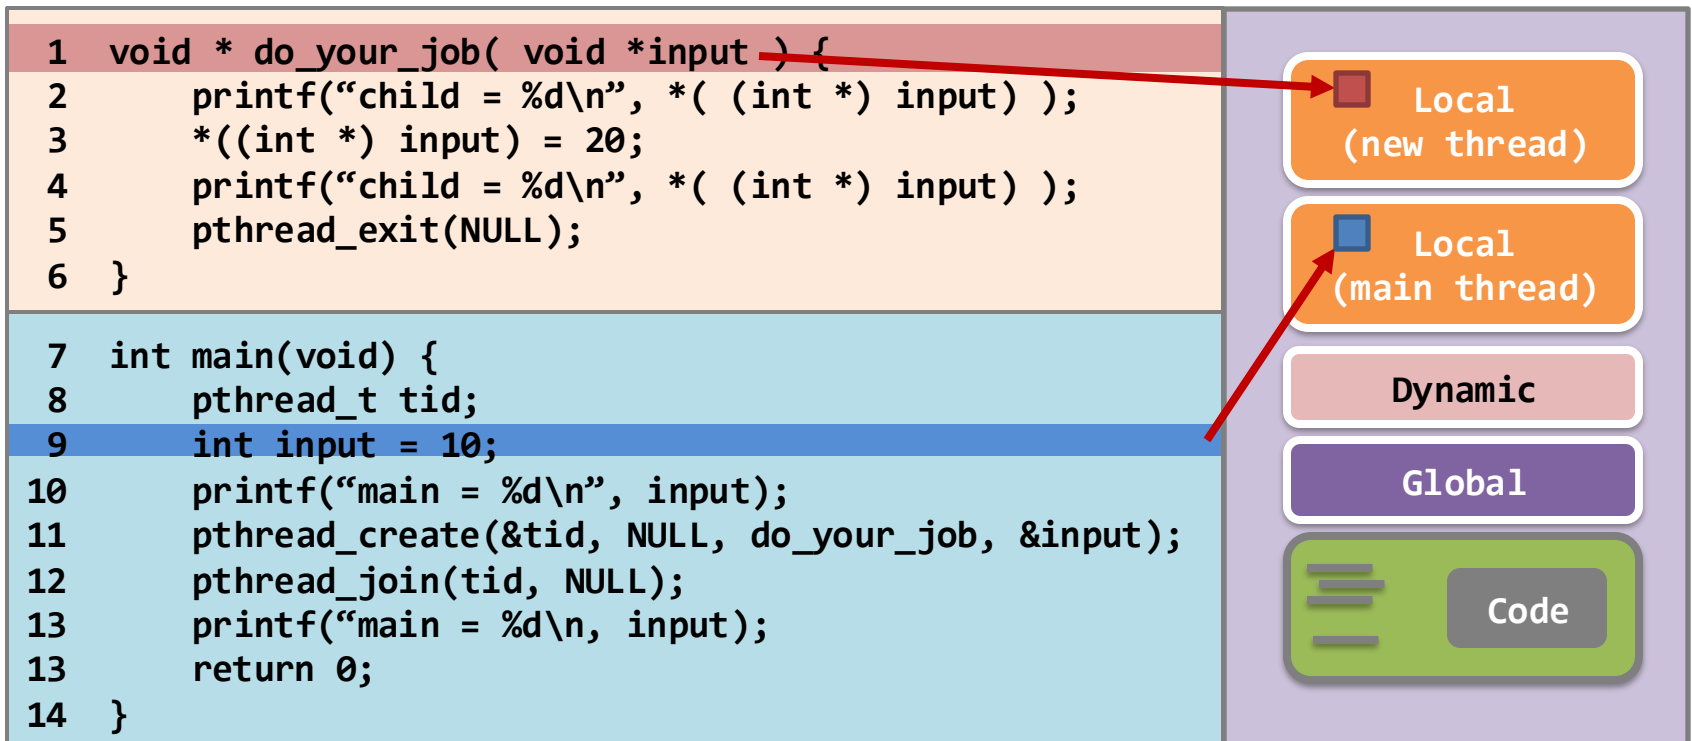
Global

Code

[examples@3150] cat pthread\_sharing.c

# Shared address space

Yet...the stack for the new thread is not on the another process, but is on the same piece of user-space memory as the main thread.



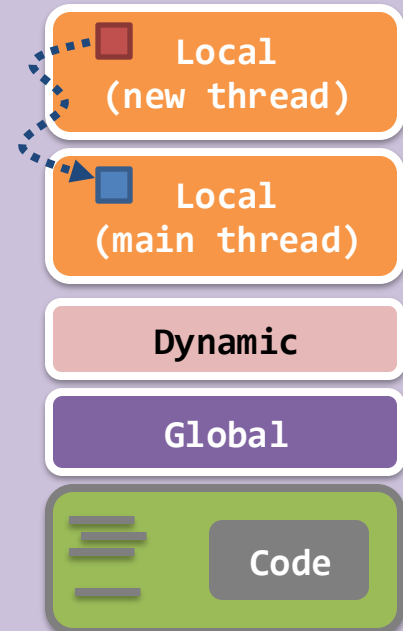
[examples@3150] cat pthread\_sharing.c

# Shared address space

The `pthread_create()` function passes an address to the new thread in the same process space

```
1 void * do_your_job( void *input ) {
2     printf("child = %d\n", *( (int *) input) );
3     *((int *) input) = 20;
4     printf("child = %d\n", *( (int *) input) );
5     pthread_exit(NULL);
6 }

7 int main(void) {
8     pthread_t tid;
9     int input = 10;
10    printf("main = %d\n", input);
11    pthread_create(&tid, NULL, do_your_job, &input);
12    pthread_join(tid, NULL);
13    printf("main = %d\n", input);
13    return 0;
14 }
```




[examples@3150] cat pthread\_sharing.c

# Shared address space

Therefore, the new thread can change the value in the main thread, and vice versa.

```
1 void * do_your_job( void *input ) {  
2     printf("child = %d\n", *( (int *) input) );  
3     *((int *) input) = 20;  
4     printf("child = %d\n", *( (int *) input) );  
5     pthread_exit(NULL);  
6 }  
  
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10;  
10    printf("main = %d\n", input);  
11    pthread_create(&tid, NULL, do_your_job, &input);  
12    pthread_join(tid, NULL);  
13    printf("main = %d\n", input);  
14    return 0;  
15 }
```



[examples@3150] cat pthread\_sharing.c

# Thread termination – passing return value?

## Thread Function

```
1 void * do_your_job(void *input) {  
2     int *output = (int *) malloc(sizeof(int));  
3     srand(time(NULL));  
4     *output = ((rand() % 10) + 1) * (((int *) input));  
5     pthread_exit(output);  
6 }
```

On thread exit, give pthread library the address of where the return value stored

## Main Function

```
7 int main(void) {  
8     pthread_t tid;  
9     int input = 10, *t-output;  
10    pthread_create(&tid, NULL, do_your_job, &input);  
11    pthread_join(tid, (void **) &t-output );  
12    printf("output = %d\n", *t-output);  
13    return 0;  
14 }
```

Pass to the library the address of a space that the main set aside to receive the thread's return value

[examples@3150] cat pthread\_exit.c

# Kernel

- The Kernel itself is a multi-threaded program
- The threads created by Kernel are named
  - kernel threads
  - OS threads

# Debunking SMT / Hyper-threading

- Simultaneous Multi-Threading (SMT)
  - It refers to **hardware** threading
  - **Not** pThread/OS level **software** threads
- Each core simulates two (virtual) cores
  - So, when a CPU sells you as 8 cores, look at the specification more carefully what 8 cores really are!
  - Virtual cores alternating between them on a physical cycle-by-cycle basis
    - When one hardware thread waits for memory
    - Another hardware thread can use the physical core to do computation