

Operating Systems

Eric Lo

8 - Scheduling

What is scheduling?

- Scheduling is required because the number of computing resource – the CPU – is **limited**.

CPU-bound Process	I/O-bound process
Spends most of its running time on the CPU, i.e., user-time > sys-time	Spends most of its running time on I/O, i.e., sys-time > user-time
<u>Examples</u> - CSCI2100 assignments, AI programs.	<u>Examples</u> - /bin/l s, networking programs.

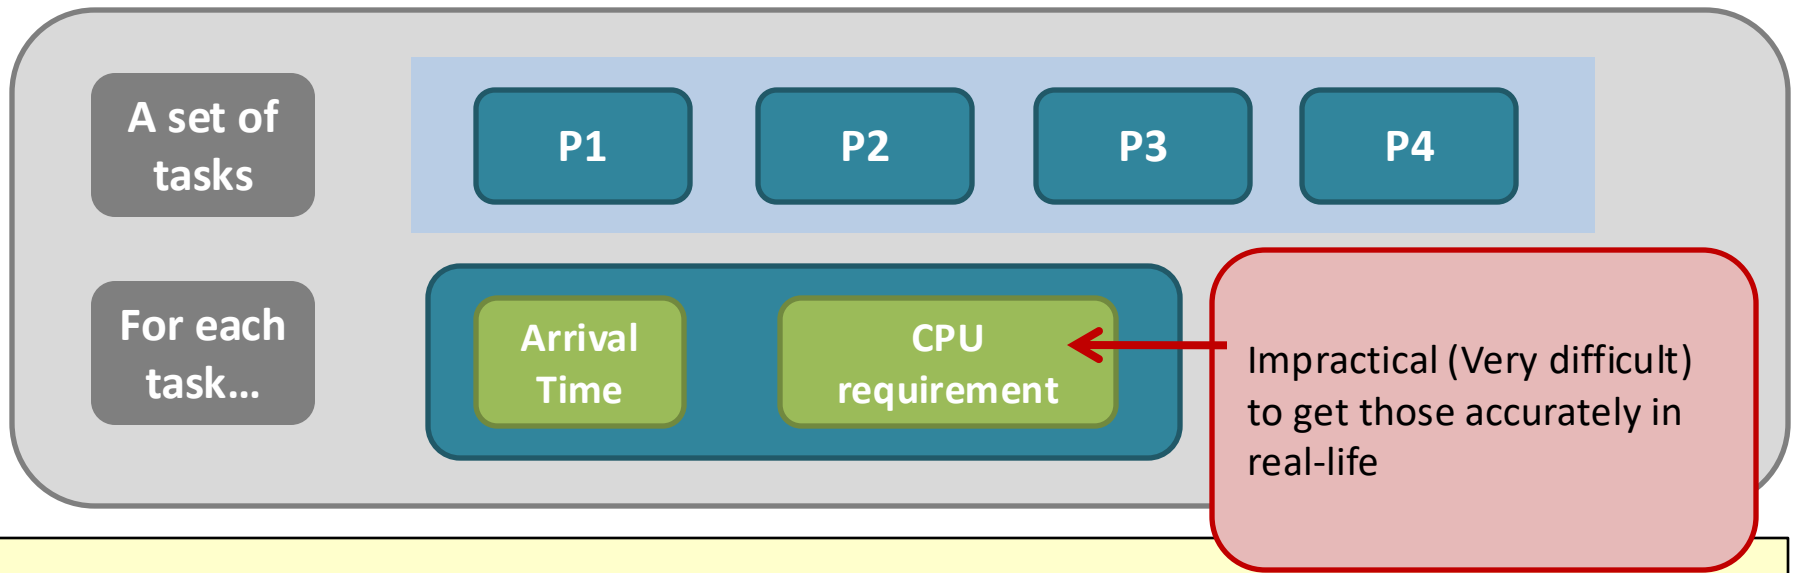
Classes of scheduling

- Preemptive scheduling (Non-preemptive is out)

What is it?	<p>When a process/thread is chosen by the scheduler, the process would have the CPU until...</p> <ul style="list-style-type: none">-the process voluntarily waits for I/O, or-the process voluntarily releases the CPU, e.g., <code>exit()</code>, <code>yield()</code>.-particular kinds of interrupts (e.g., periodic clock interrupt, a new process steps in) are detected.
History	<p>In old days, it was called “time-sharing”</p> <p>Nowadays, all systems are time-sharing</p>
Pros	<p>Good for systems that emphasize interactiveness.</p> <ul style="list-style-type: none">- Because every task will receive attentions from the CPU.
Cons	<p>Bad for systems that emphasize the time in finishing tasks.</p>

Scheduling algorithms

- Inputs to the algorithms.



Online VS Offline

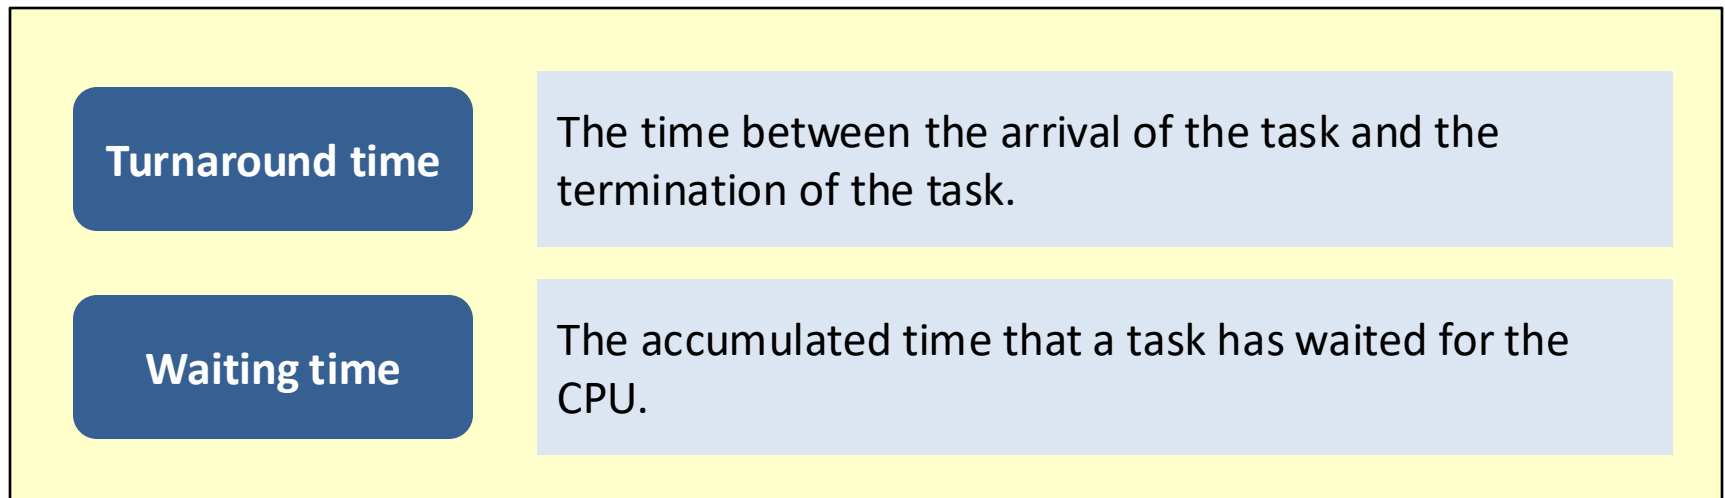
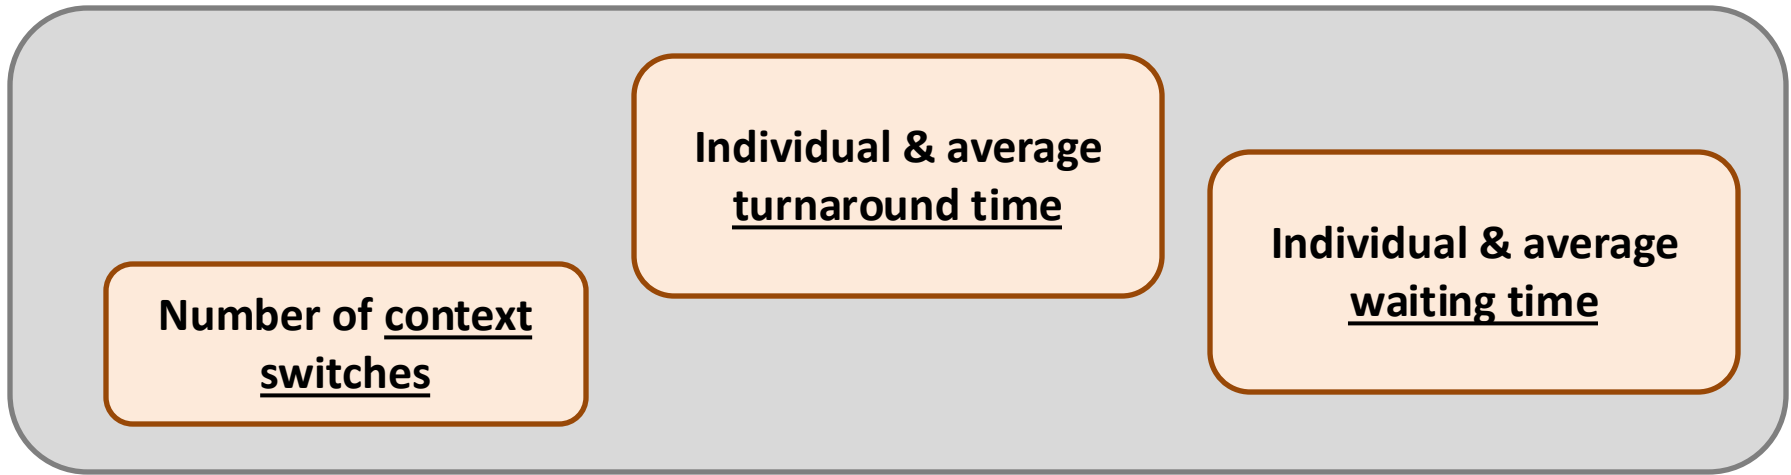
An **offline scheduling algorithm** assumes that you know the sequence of processes that a scheduler will face

- Theoretical baseline

An **online scheduling algorithm** does not have such an assumption.

- Practical use

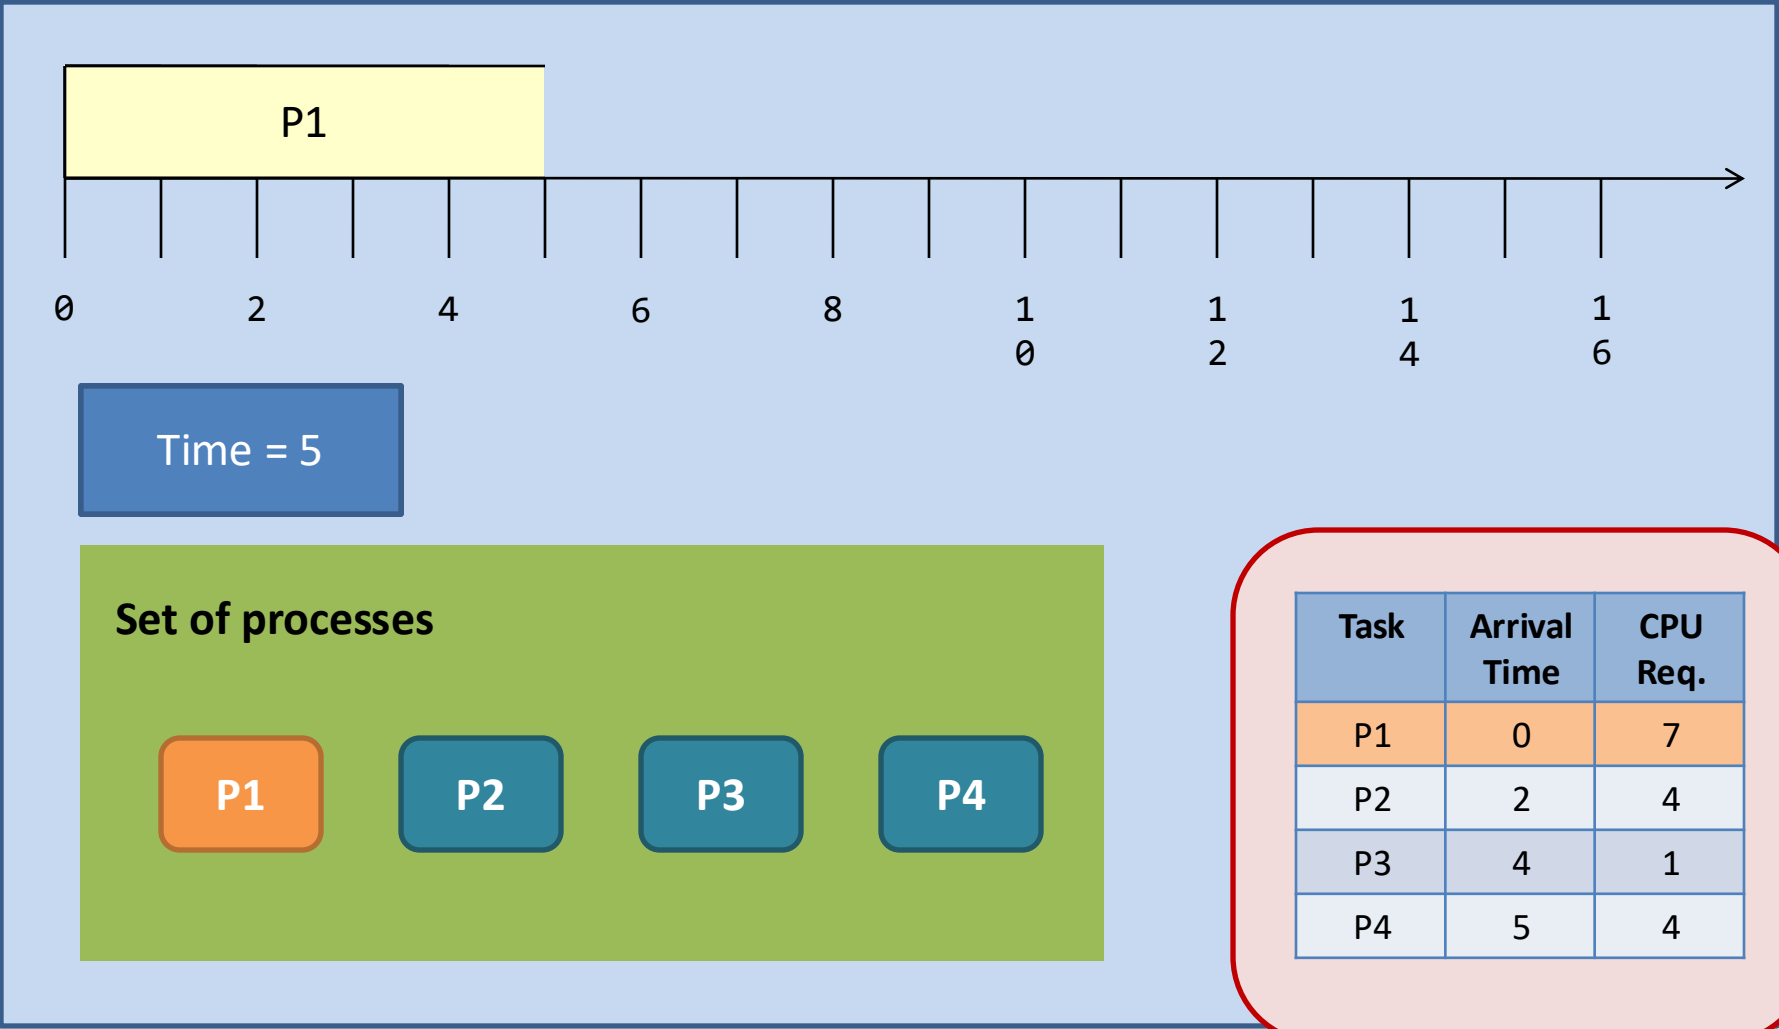
Algorithm evaluation



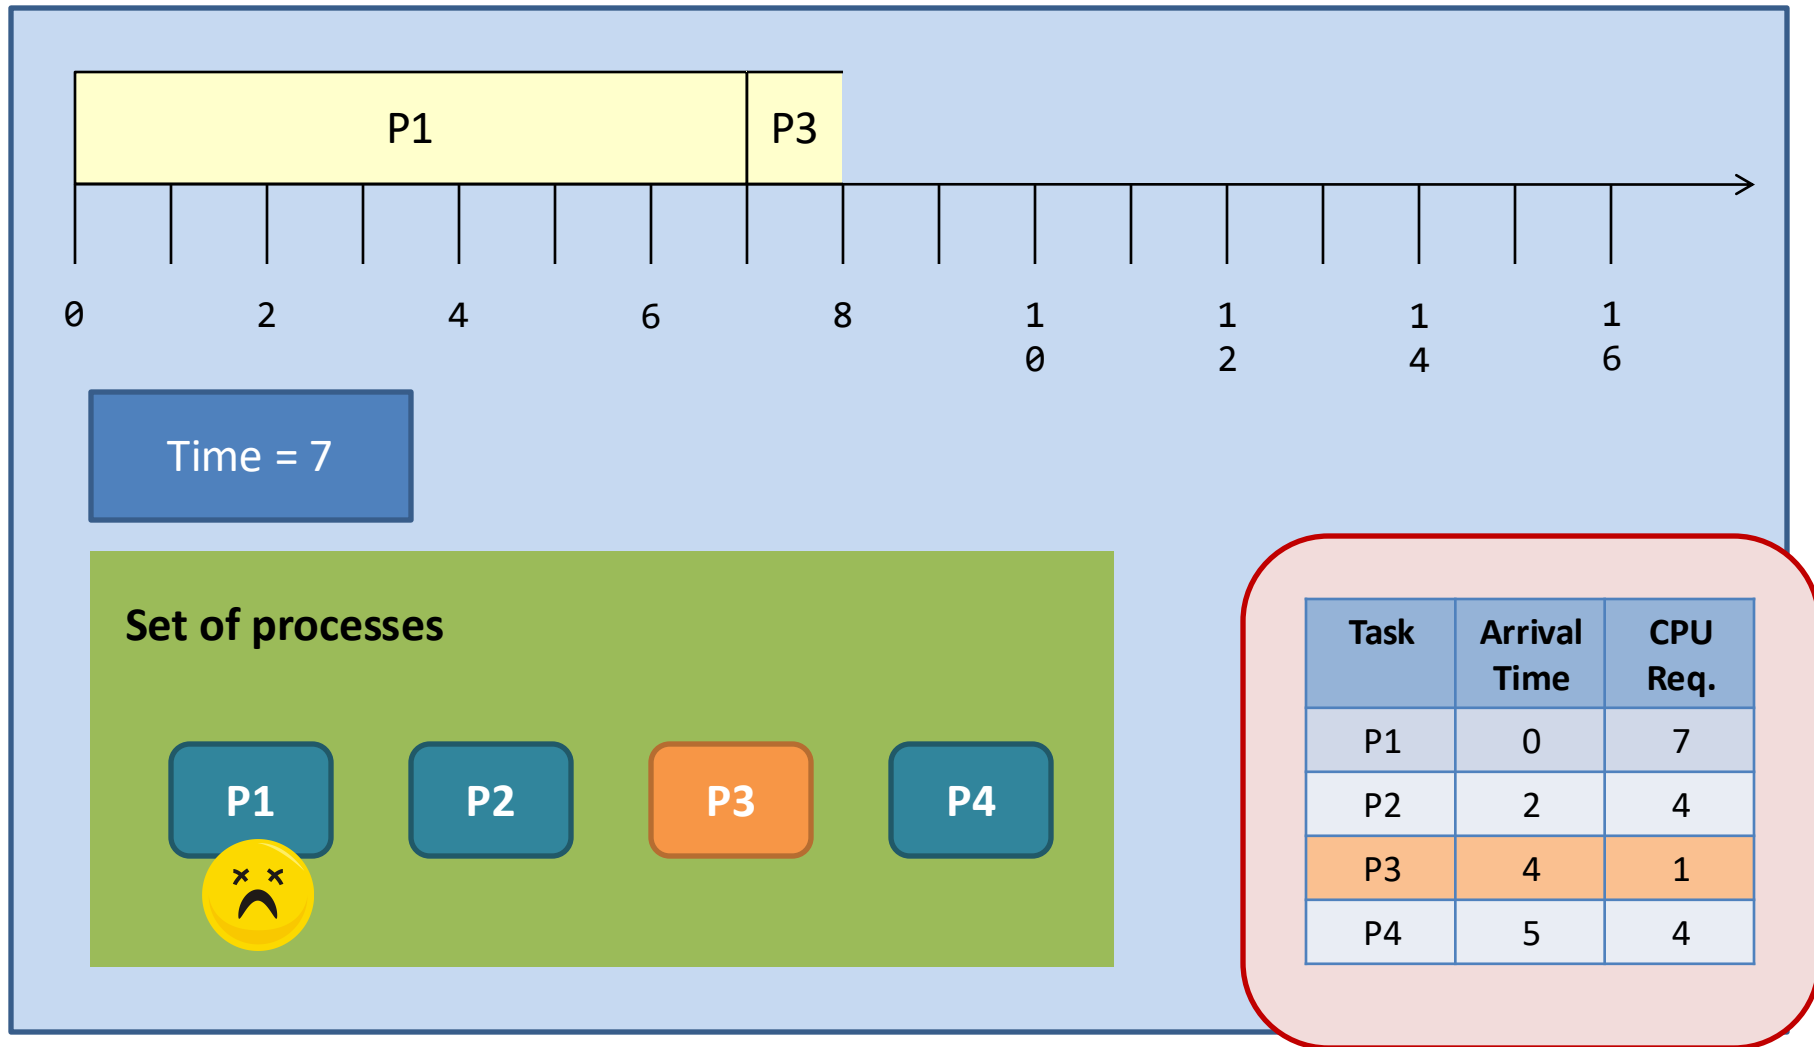
Different algorithms

Algorithms
Shortest-job-first (SJF)
Round-robin (RR)
Priority scheduling with multiple queues

Non-preemptive SJF (assume context switch is free)




Non-preemptive SJF



Non-preemptive SJF

In this example, we use **FIFO** to break the tie.



Time = 8

Set of processes

P1


P2

P3

P4


Task	Arrival Time	CPU Req.
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Non-preemptive SJF




Time = 16

Set of processes



Task	Arrival Time	CPU Req.
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Non-preemptive SJF



Waiting time:

$P1 = 0; P2 = 6; P3 = 3; P4 = 7;$

$Average = (0 + 6 + 3 + 7) / 4 = 4.$

Turnaround time:


$P1 = 7; P2 = 10; P3 = 4; P4 = 11;$

$Average = (7 + 10 + 4 + 11) / 4 = 8.$

Task	Arrival Time	CPU Req.
P1	0	7
P2	2	4
P3	4	1
P4	5	4

- Problem:
 - What if tasks arrive after P2 **all have CPU requirement < 3** ?
 - Problem persists even for its preemptive version

Preemptive SJF




-Whenever a new process arrives at the system, the scheduler steps in and selects the next task based on **their remaining CPU requirements**.

Task	Arrival Time	CPU Req.	
		Initial	Remain
P1	0	7	7
P2	2	4	4
P3	4	1	1
P4	5	4	4


Preemptive SJF

Animation; don't print




Preemptive SJF

Animation; don't print



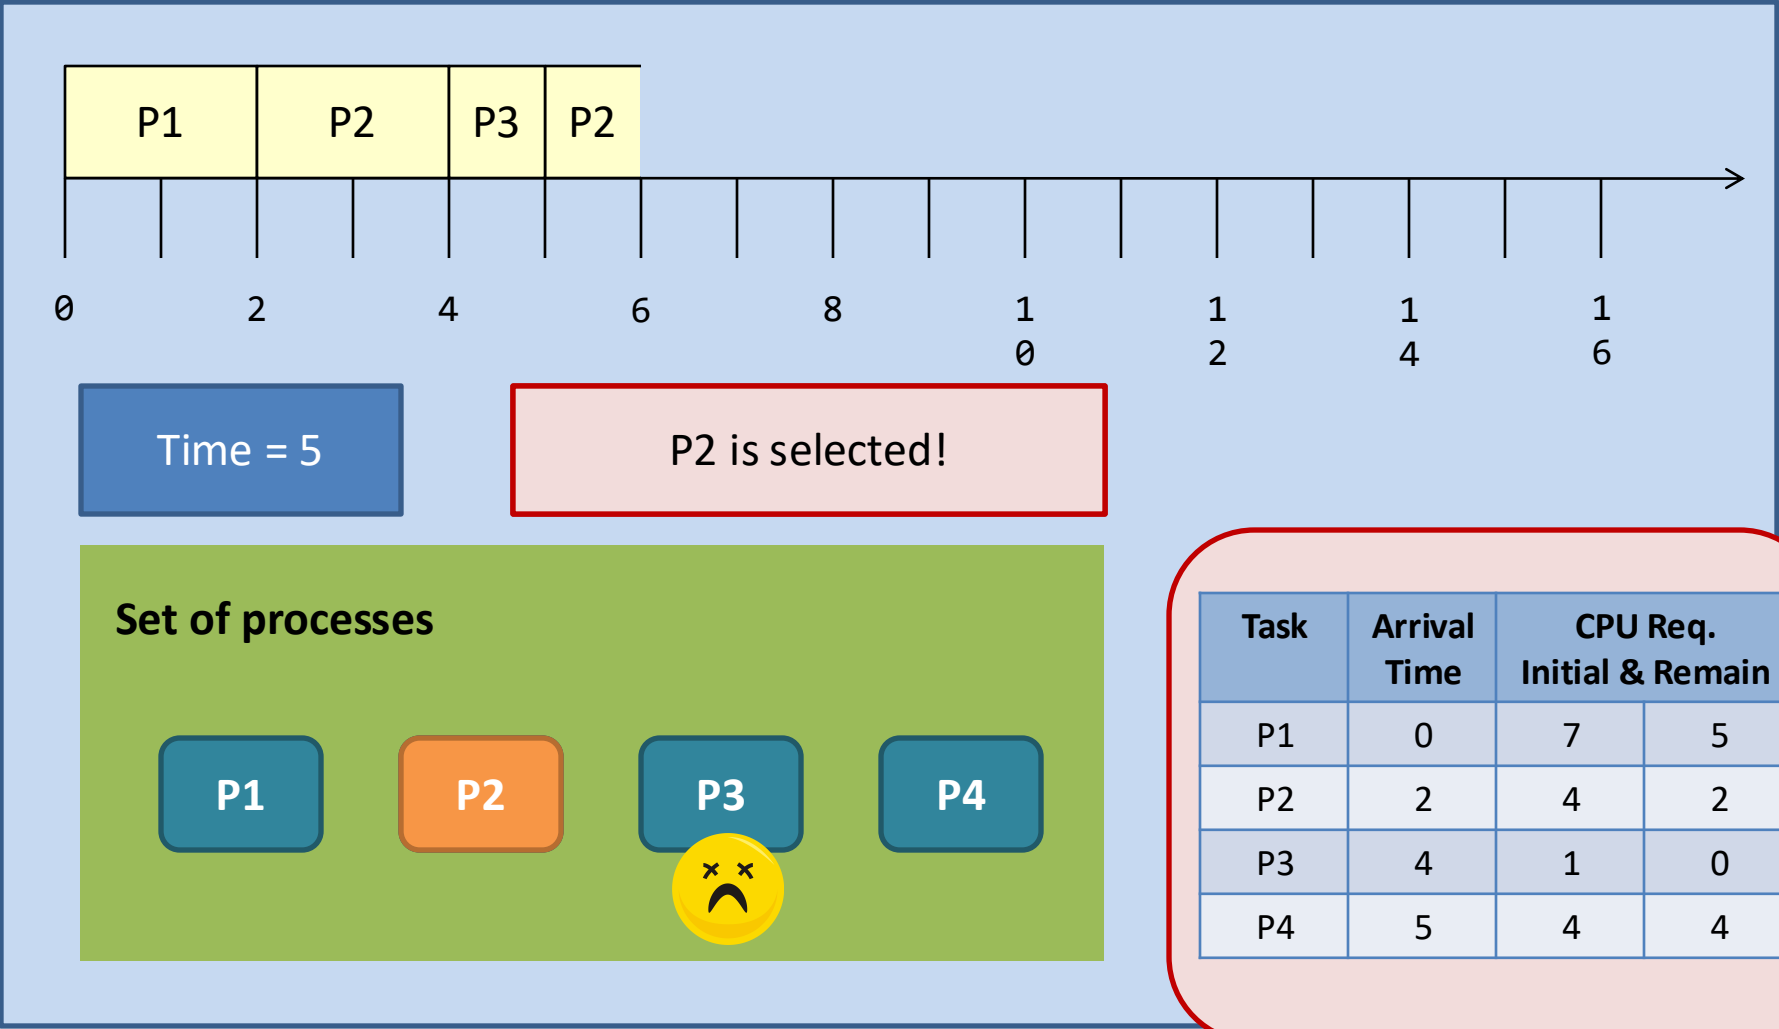
Preemptive SJF

Animation; don't print



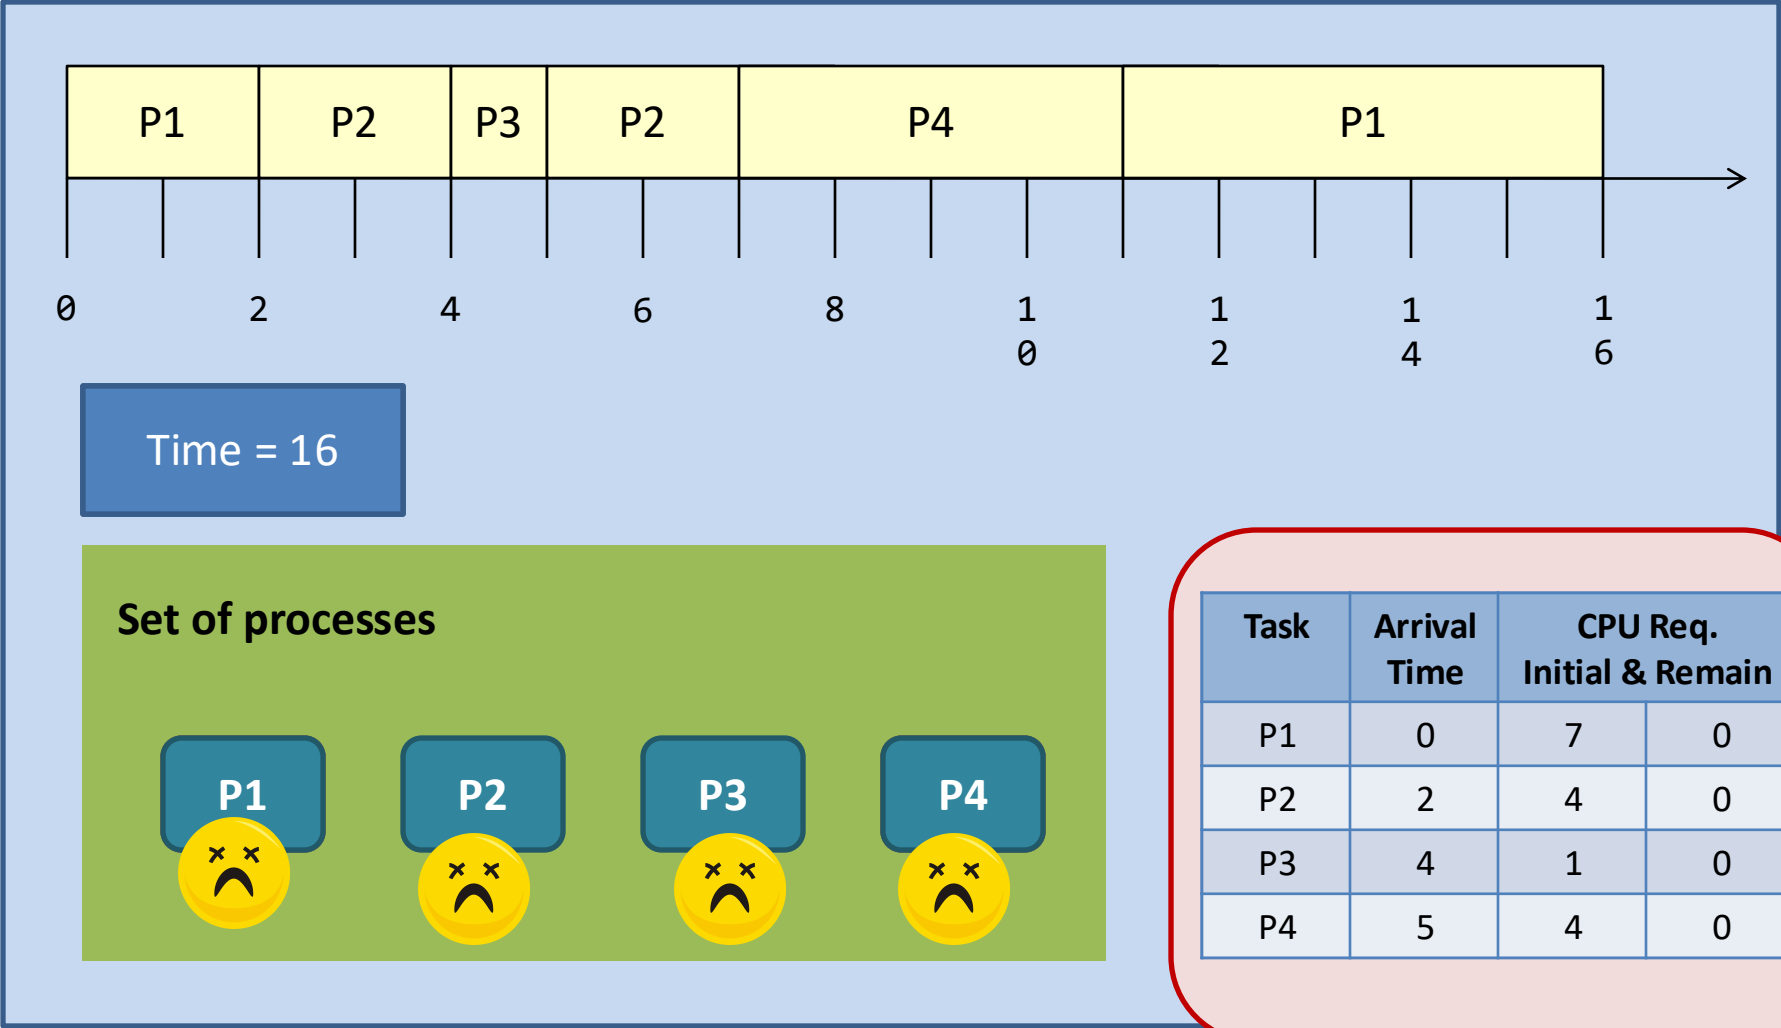
Preemptive SJF

Animation; don't print




Preemptive SJF

Animation; don't print



Preemptive SJF



Waiting time:

$P1 = 9; P2 = 1; P3 = 0; P4 = 2;$

$Average = (9 + 1 + 0 + 2) / 4 = 3.$

Turnaround time:


$P1 = 16; P2 = 5; P3 = 1; P4 = 6;$

$Average = (16 + 5 + 1 + 6) / 4 = 7.$

Task	Arrival Time	CPU Req. Initial & Remain	
P1	0	7	0
P2	2	4	0
P3	4	1	0
P4	5	4	0

SJF: Preemptive or not?

	Non-preemptive SJF	Preemptive SJF
Average waiting time	4	3 (smallest)
Average turnaround time	8	7 (smallest)
# of context switching	3	5 (largest)



The waiting time and the turnaround time decrease at the expense of the **increased number of context switches**.

Context switch is expensive. (That's why we shall minimize the # of sys calls as well; on a syscall, the program switch from user-process to kernel-"process".)

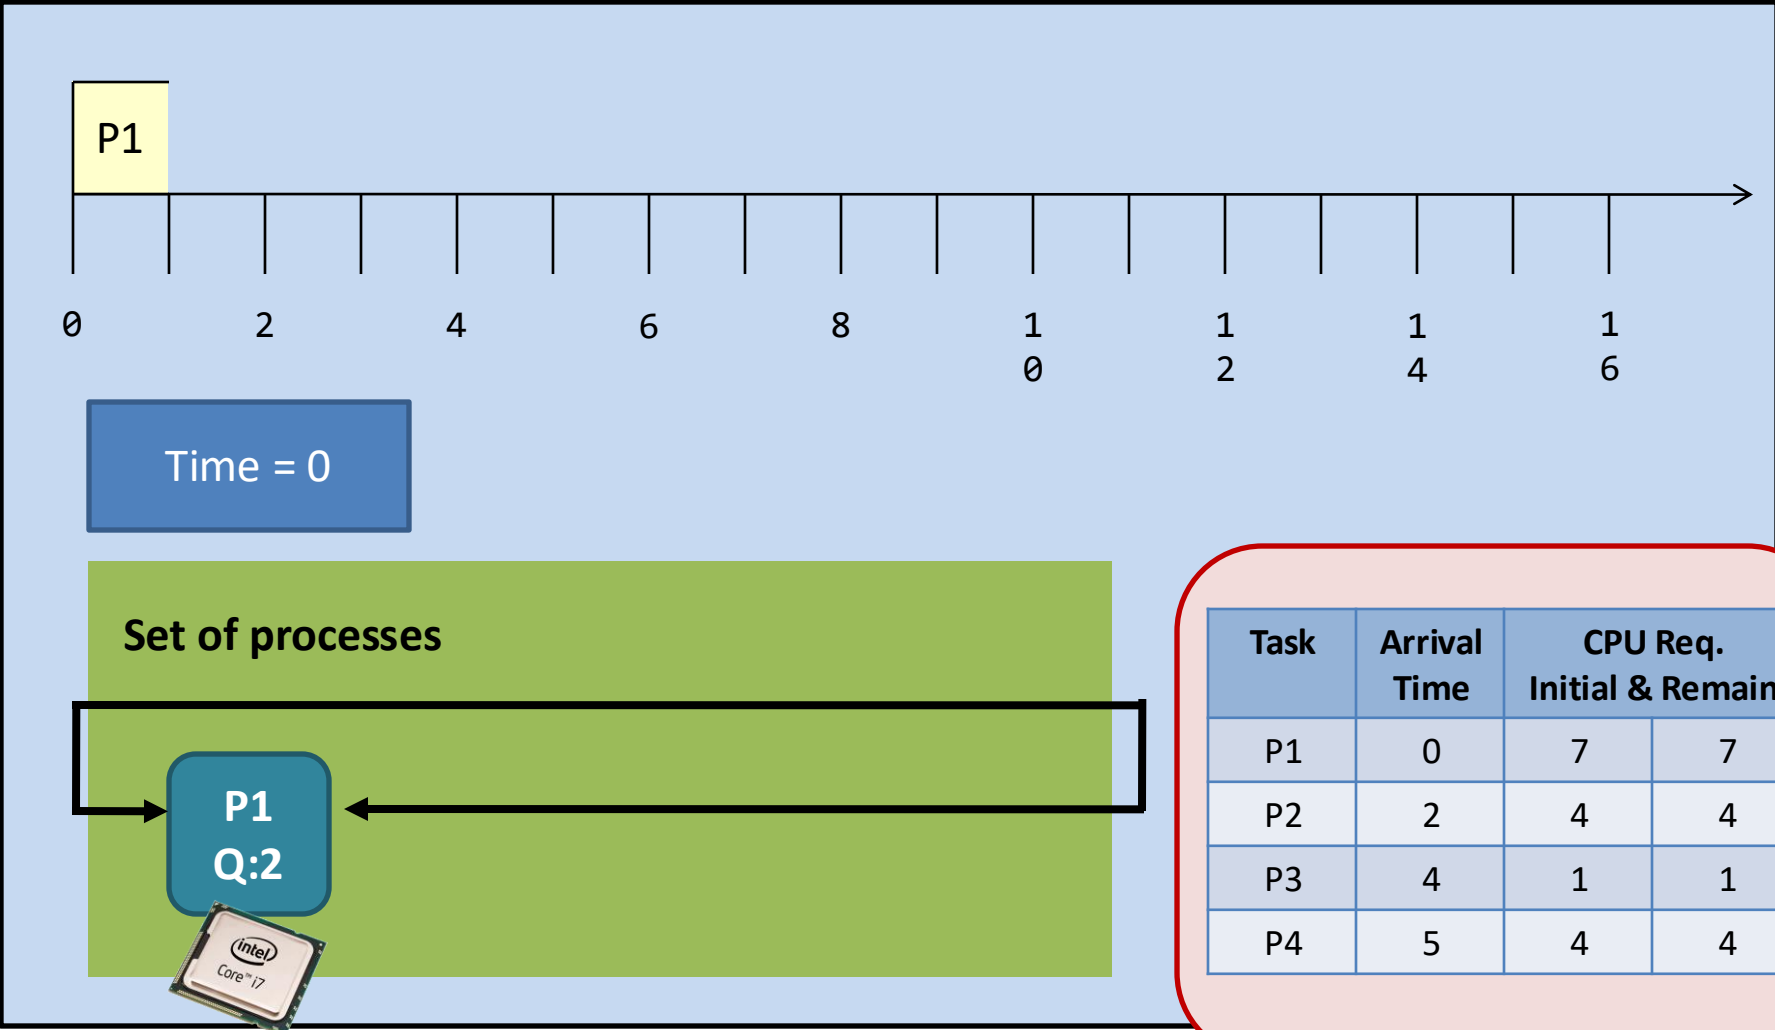
Task	Arrival Time	CPU Req.
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Round-robin

- Round-Robin (RR) scheduling is preemptive.
 - Every process is given a **quantum**, or the amount of time allowed to execute.
 - Whenever the quantum of a process is used up (i.e., 0), the process releases the CPU and **this is the preemption**.
 - Then, the scheduler steps in and it chooses **the next process which has a non-zero quantum** to run.
 - If all processes in the system have used up the quantum, they will be re-charged to their initial values.
 - Processes are therefore running one-by-one as a **circular queue**, for the basic version (i.e., no priority)
 - New processes are added to the tail of the ready queue
 - New process arrives won't trigger a new selection decision

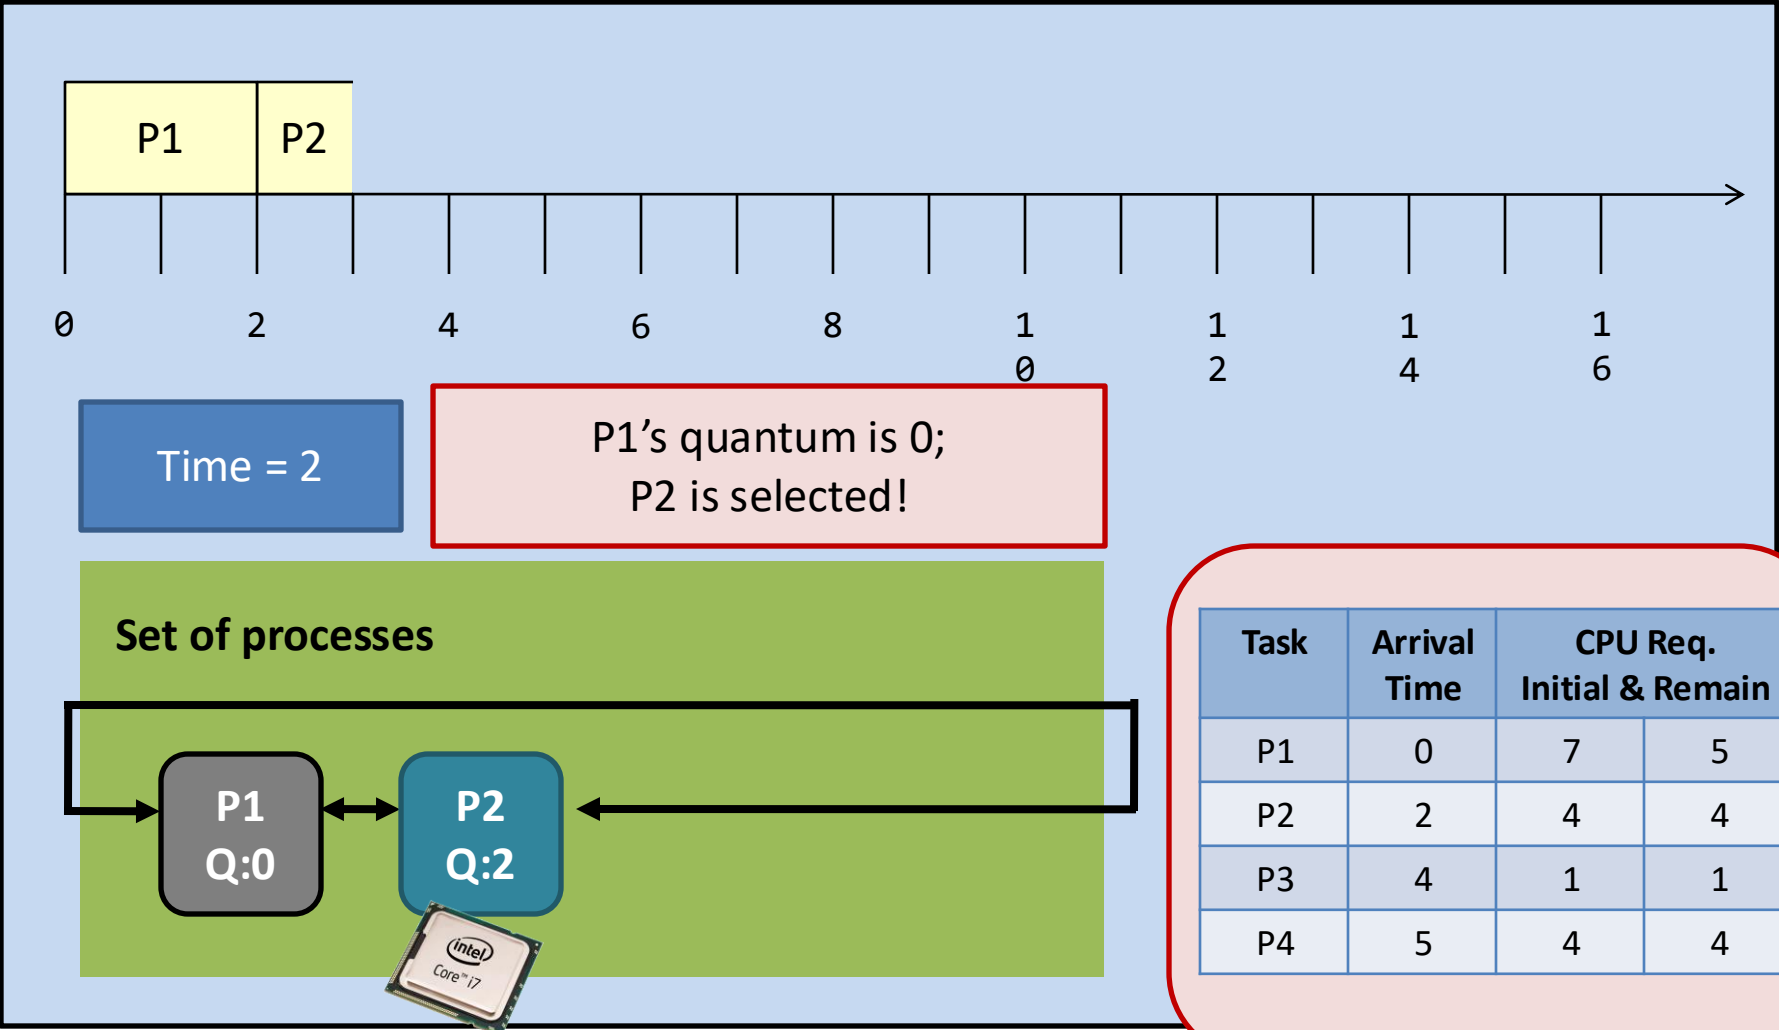
Round-robin (quantum =2)

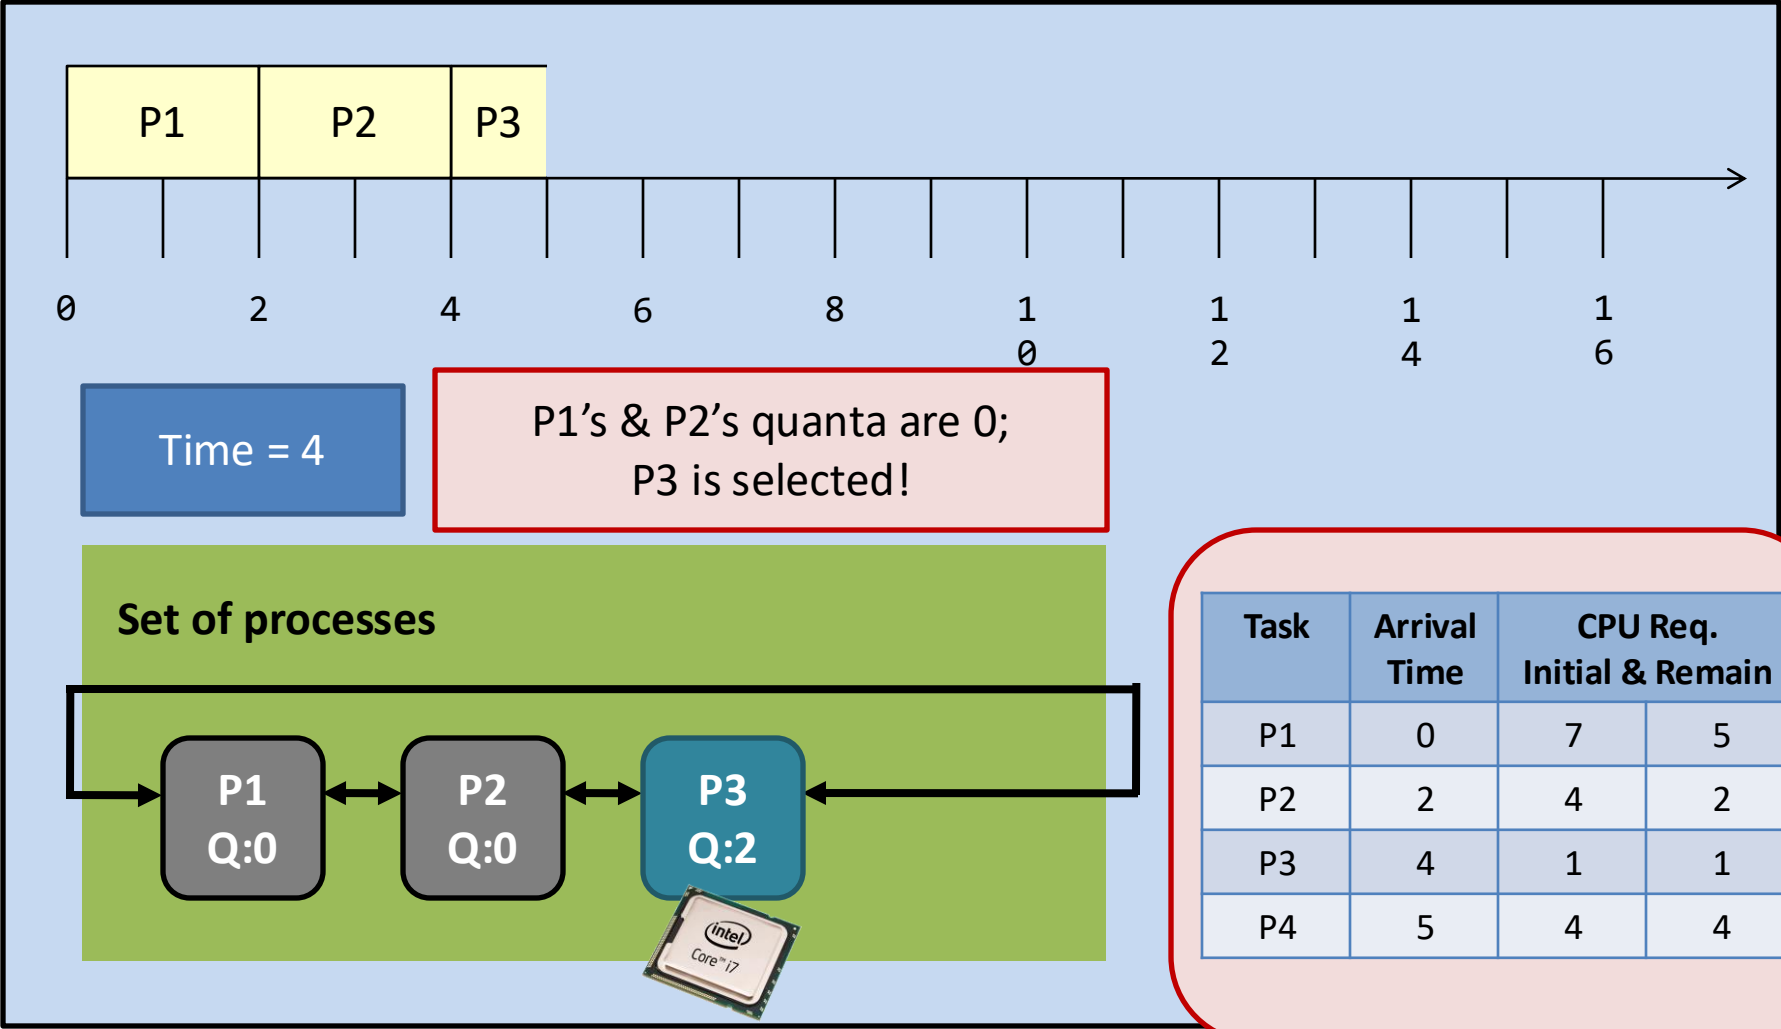
Animation; don't print



Round-robin

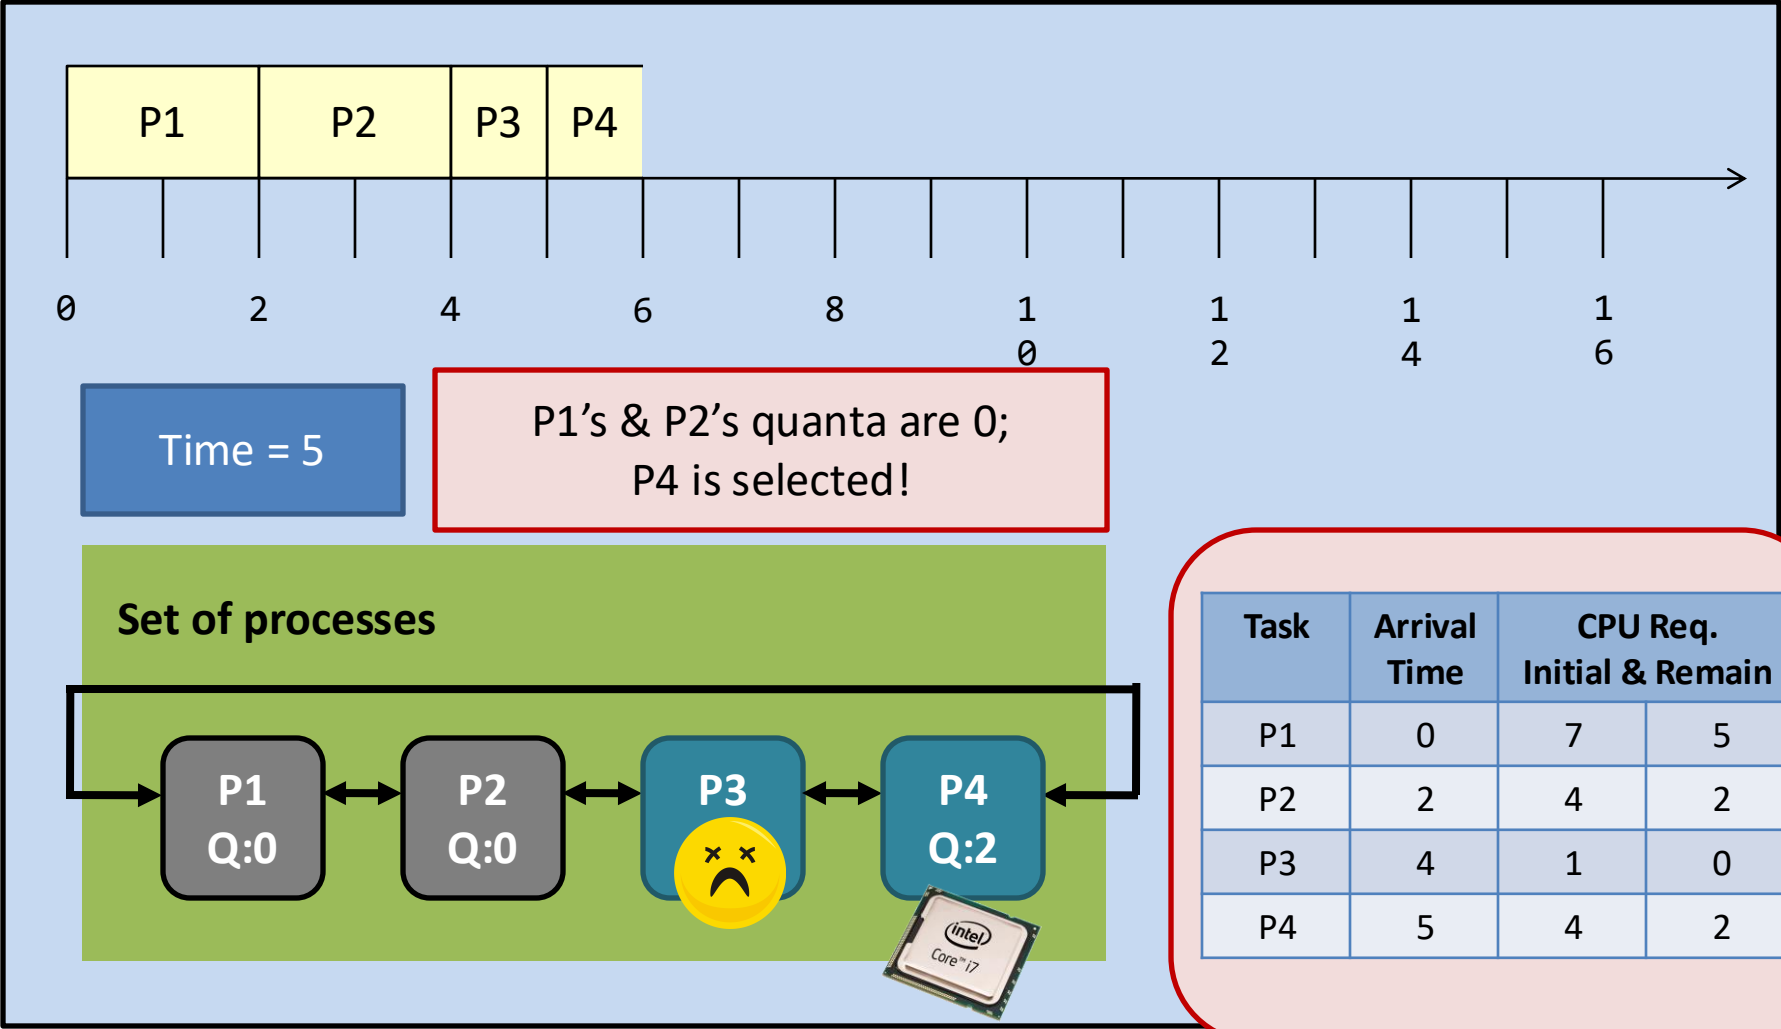
Animation; don't print





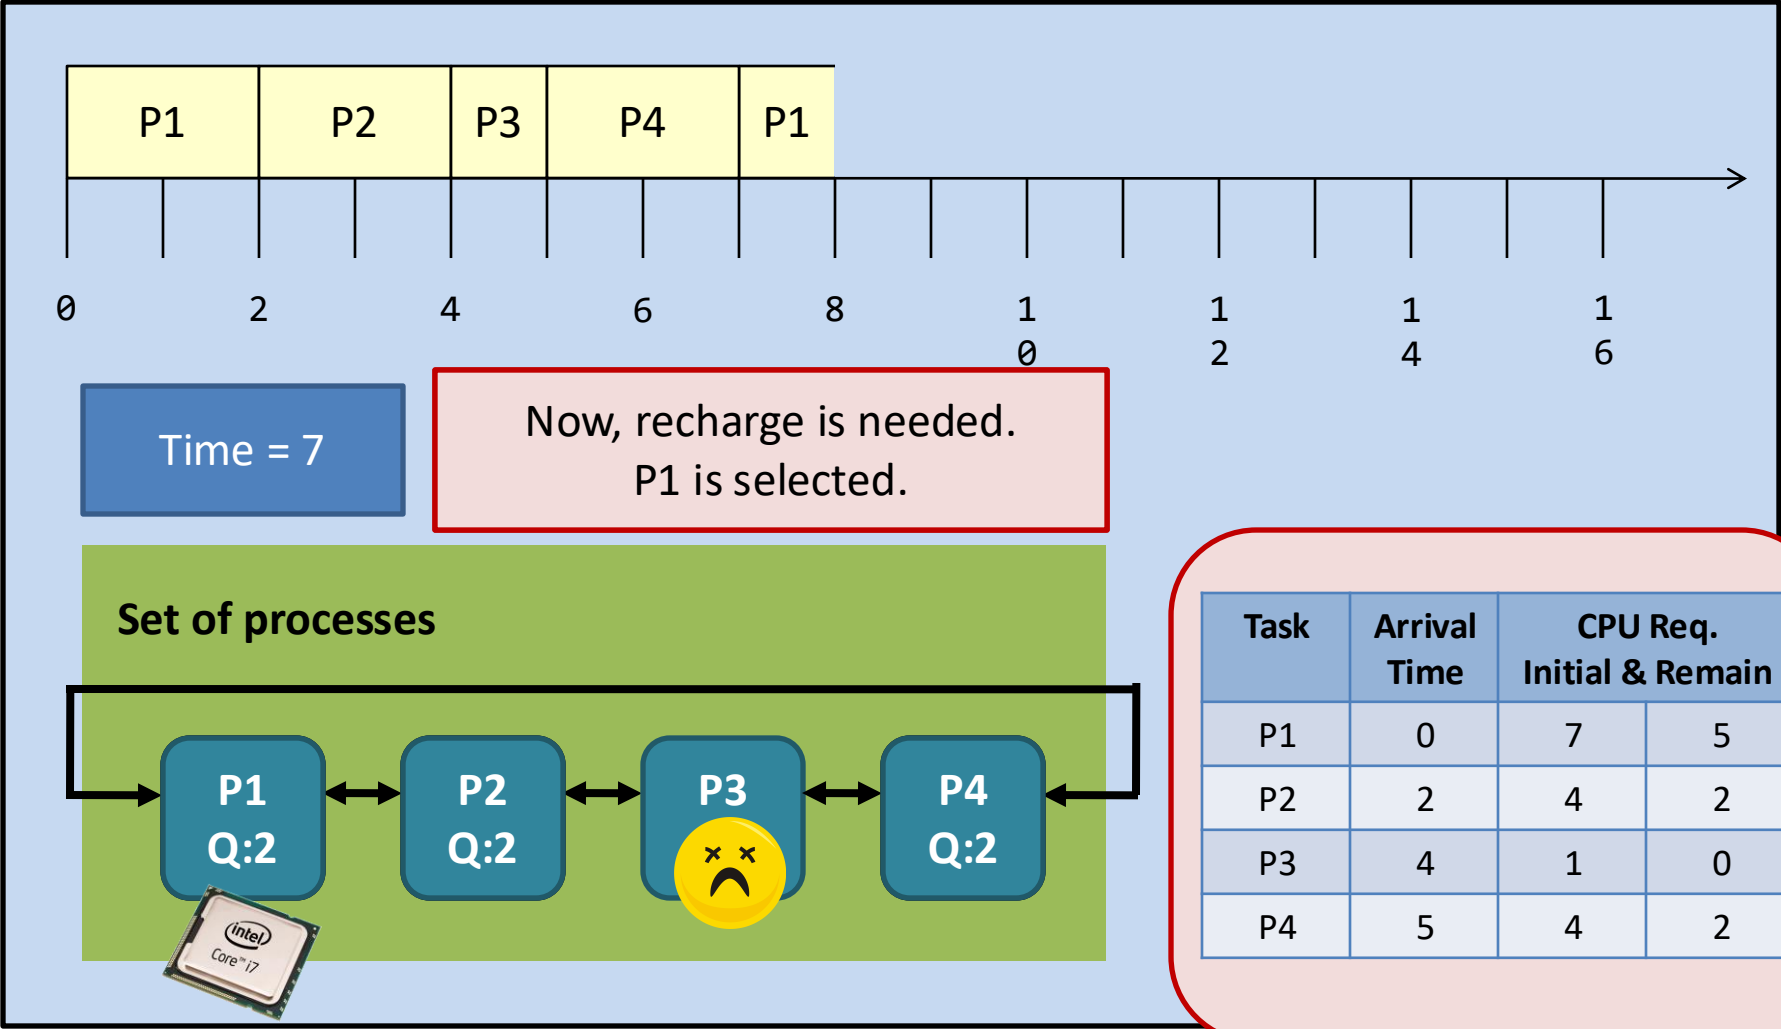
Round-robin

Animation; don't print



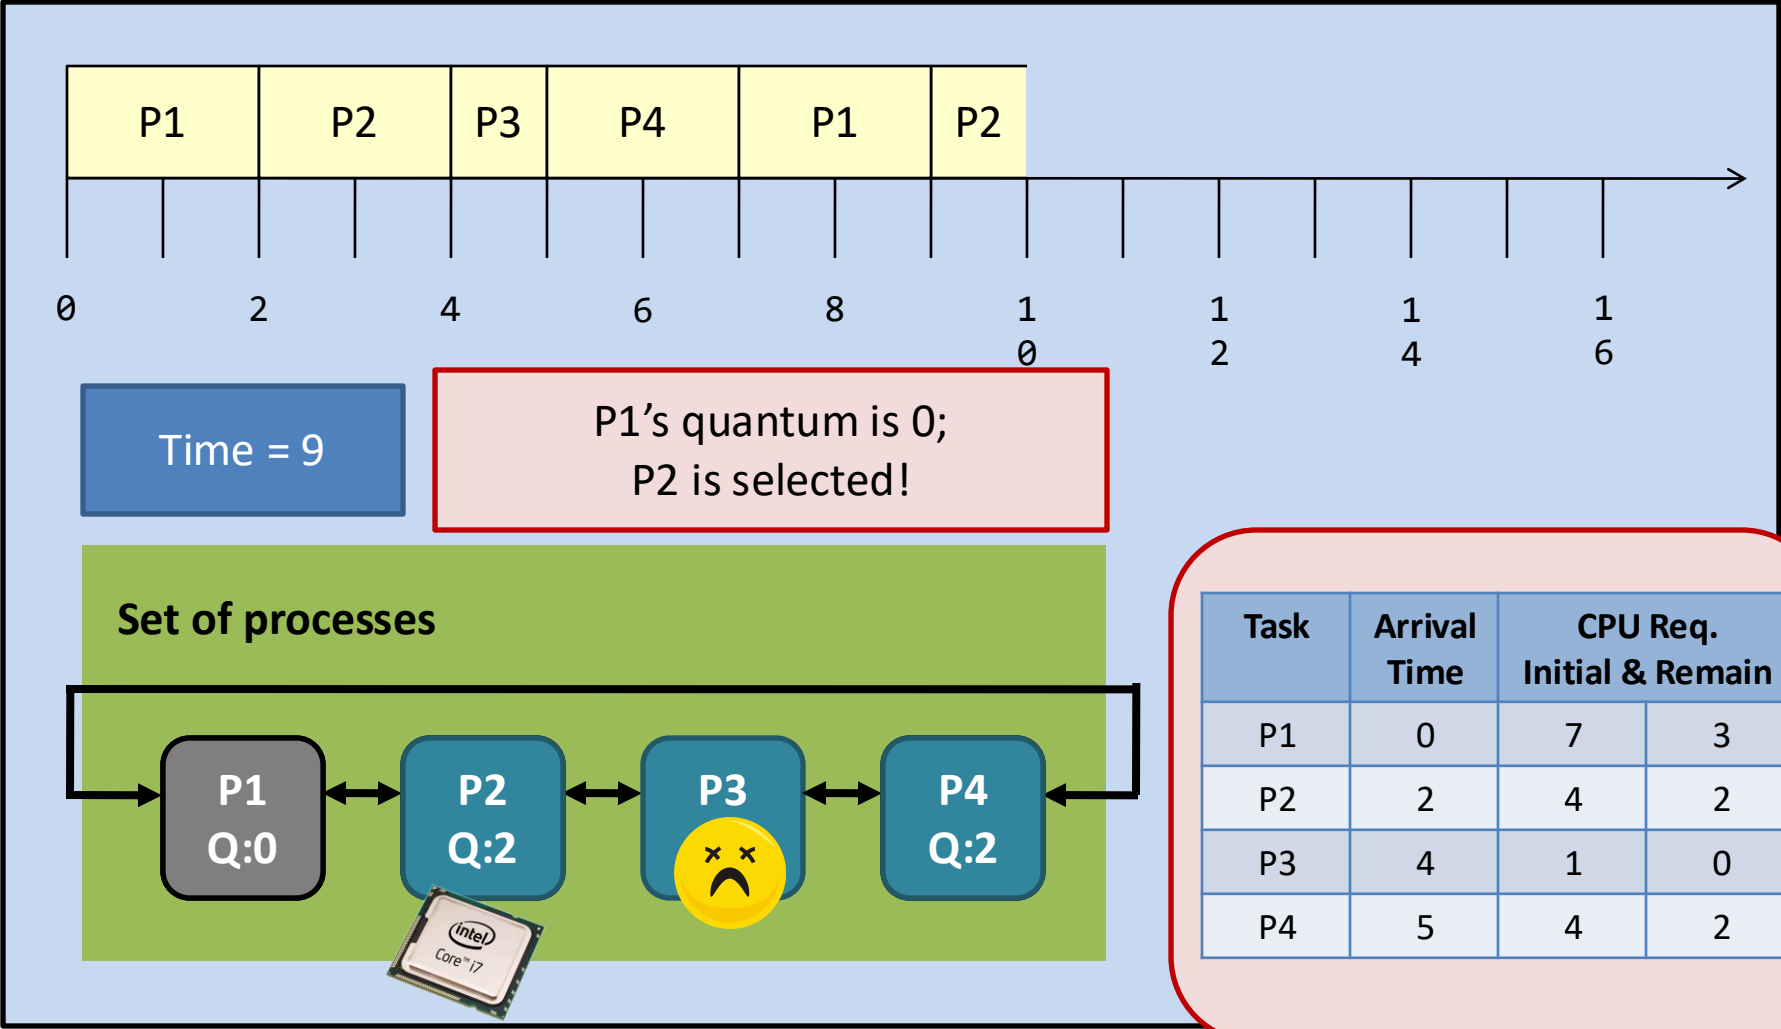
Round-robin

Animation; don't print



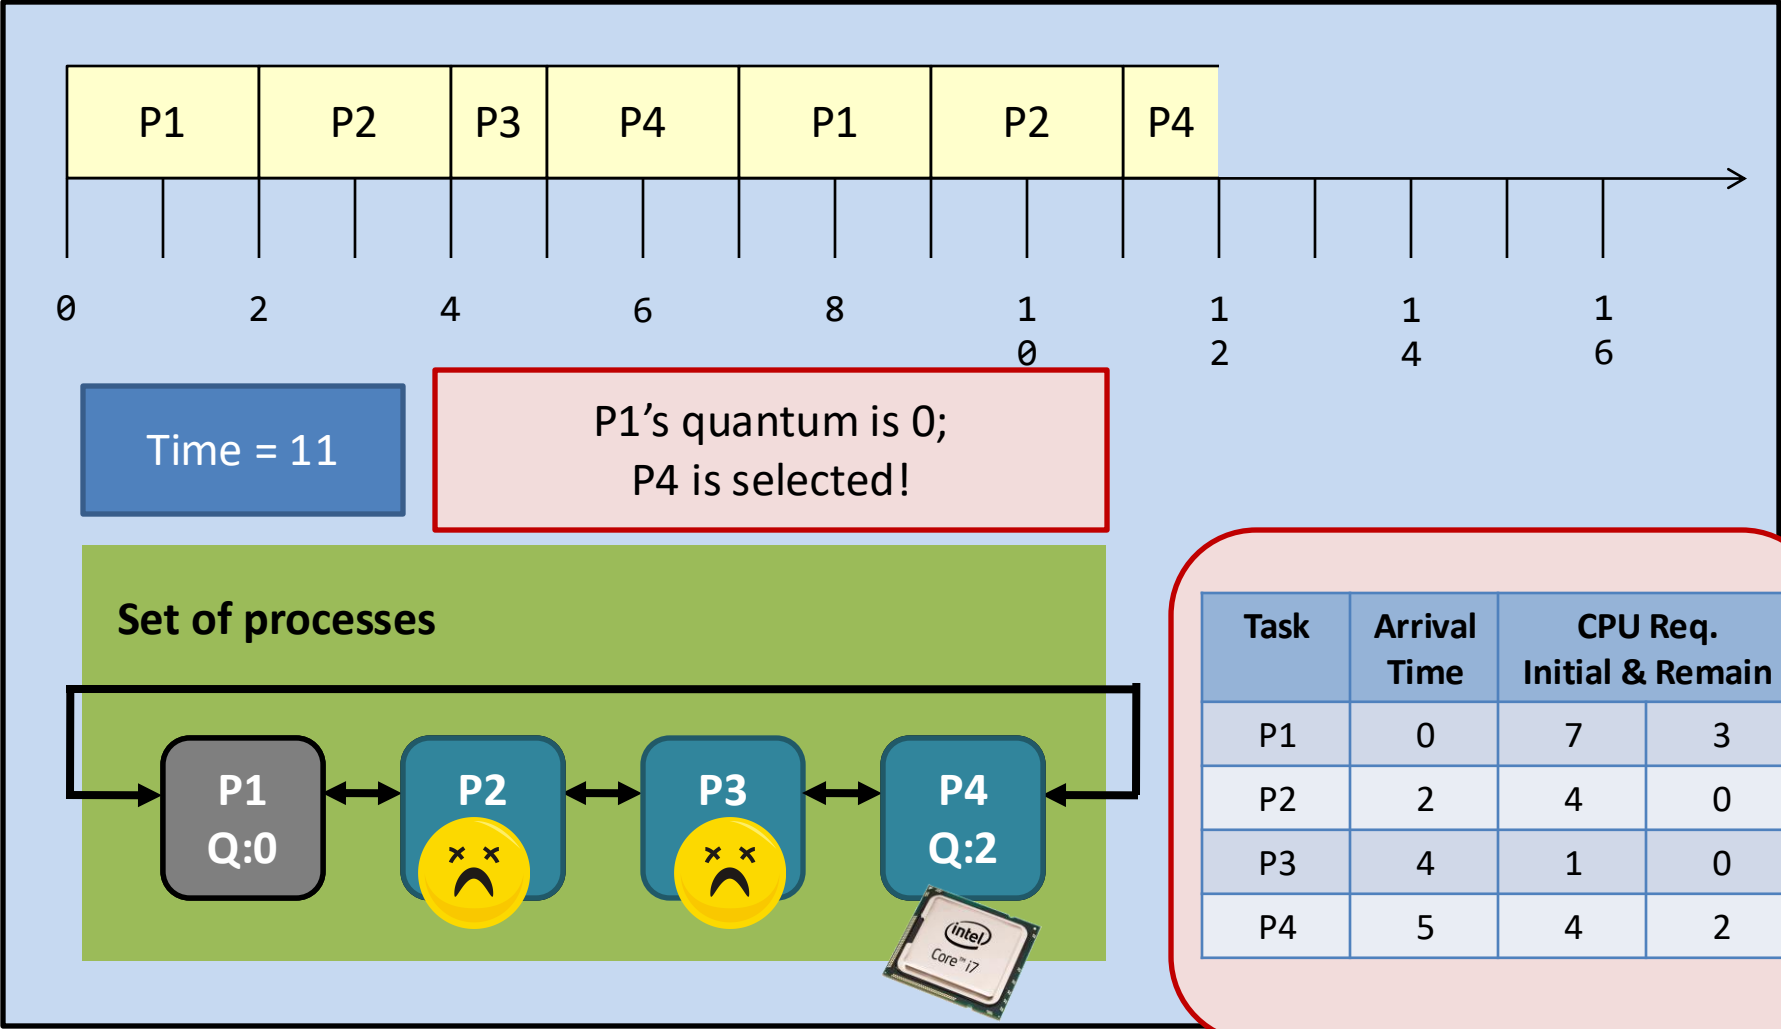
Round-robin

Animation; don't print



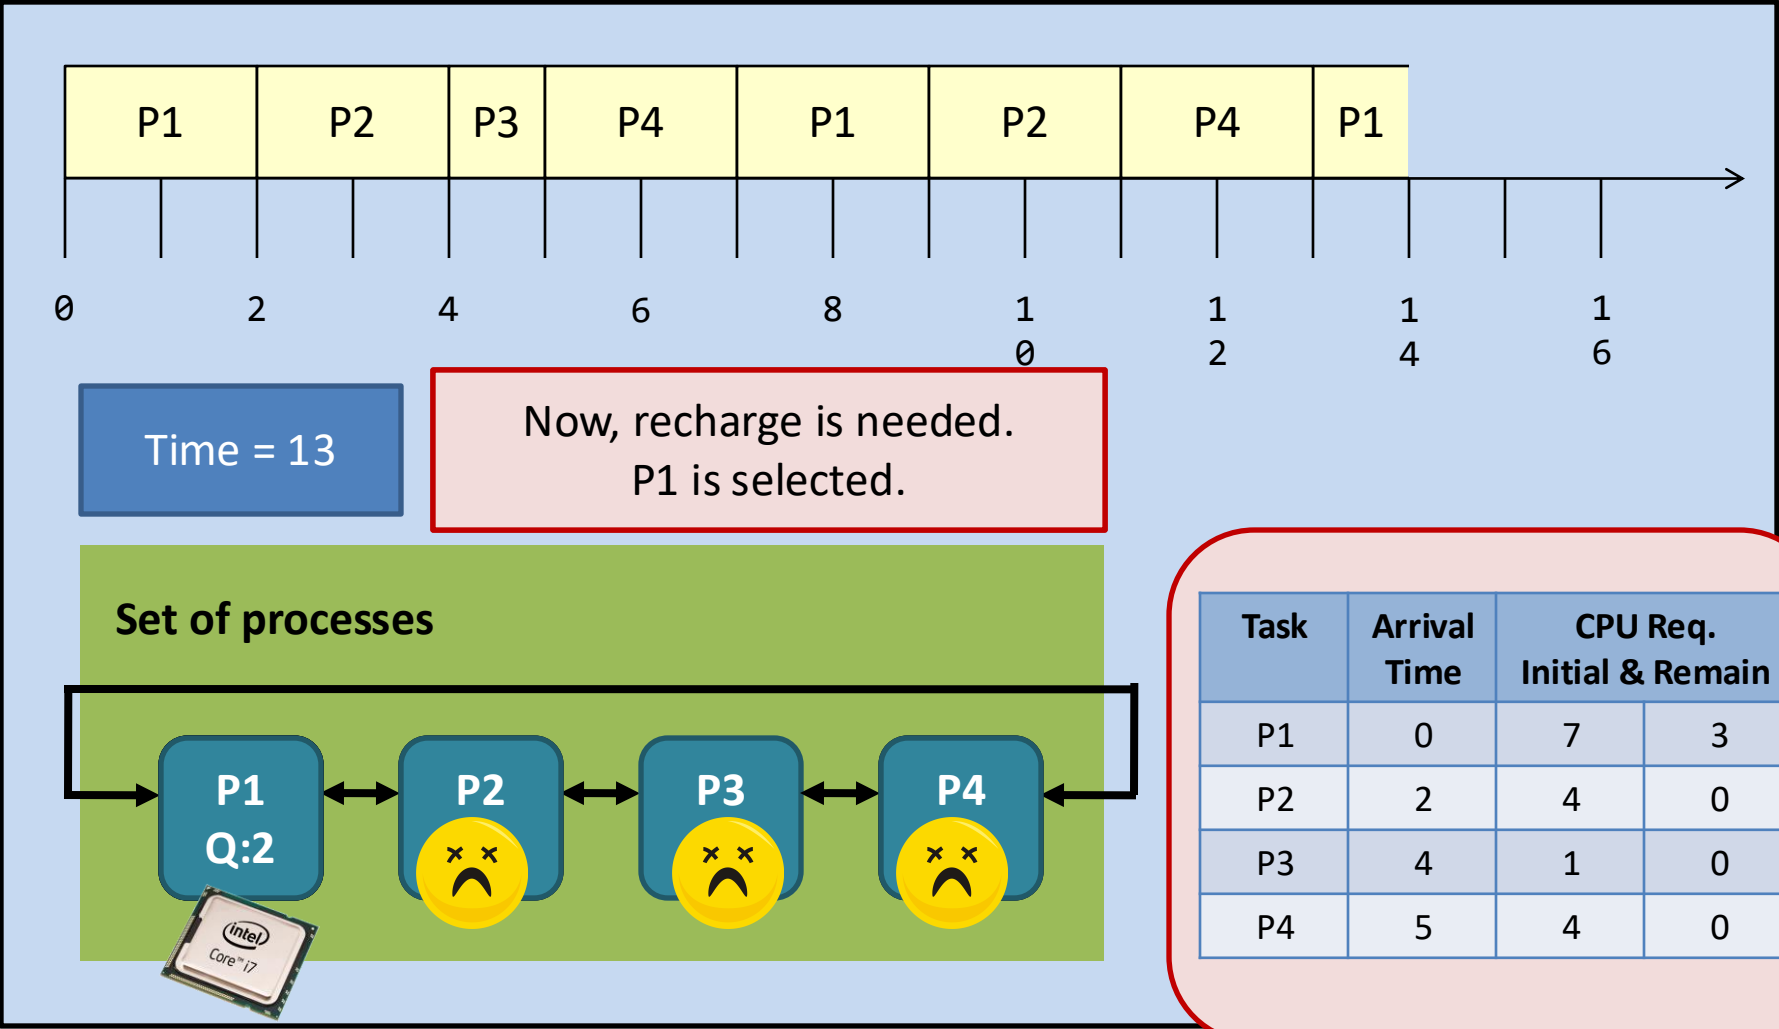
Round-robin

Animation; don't print



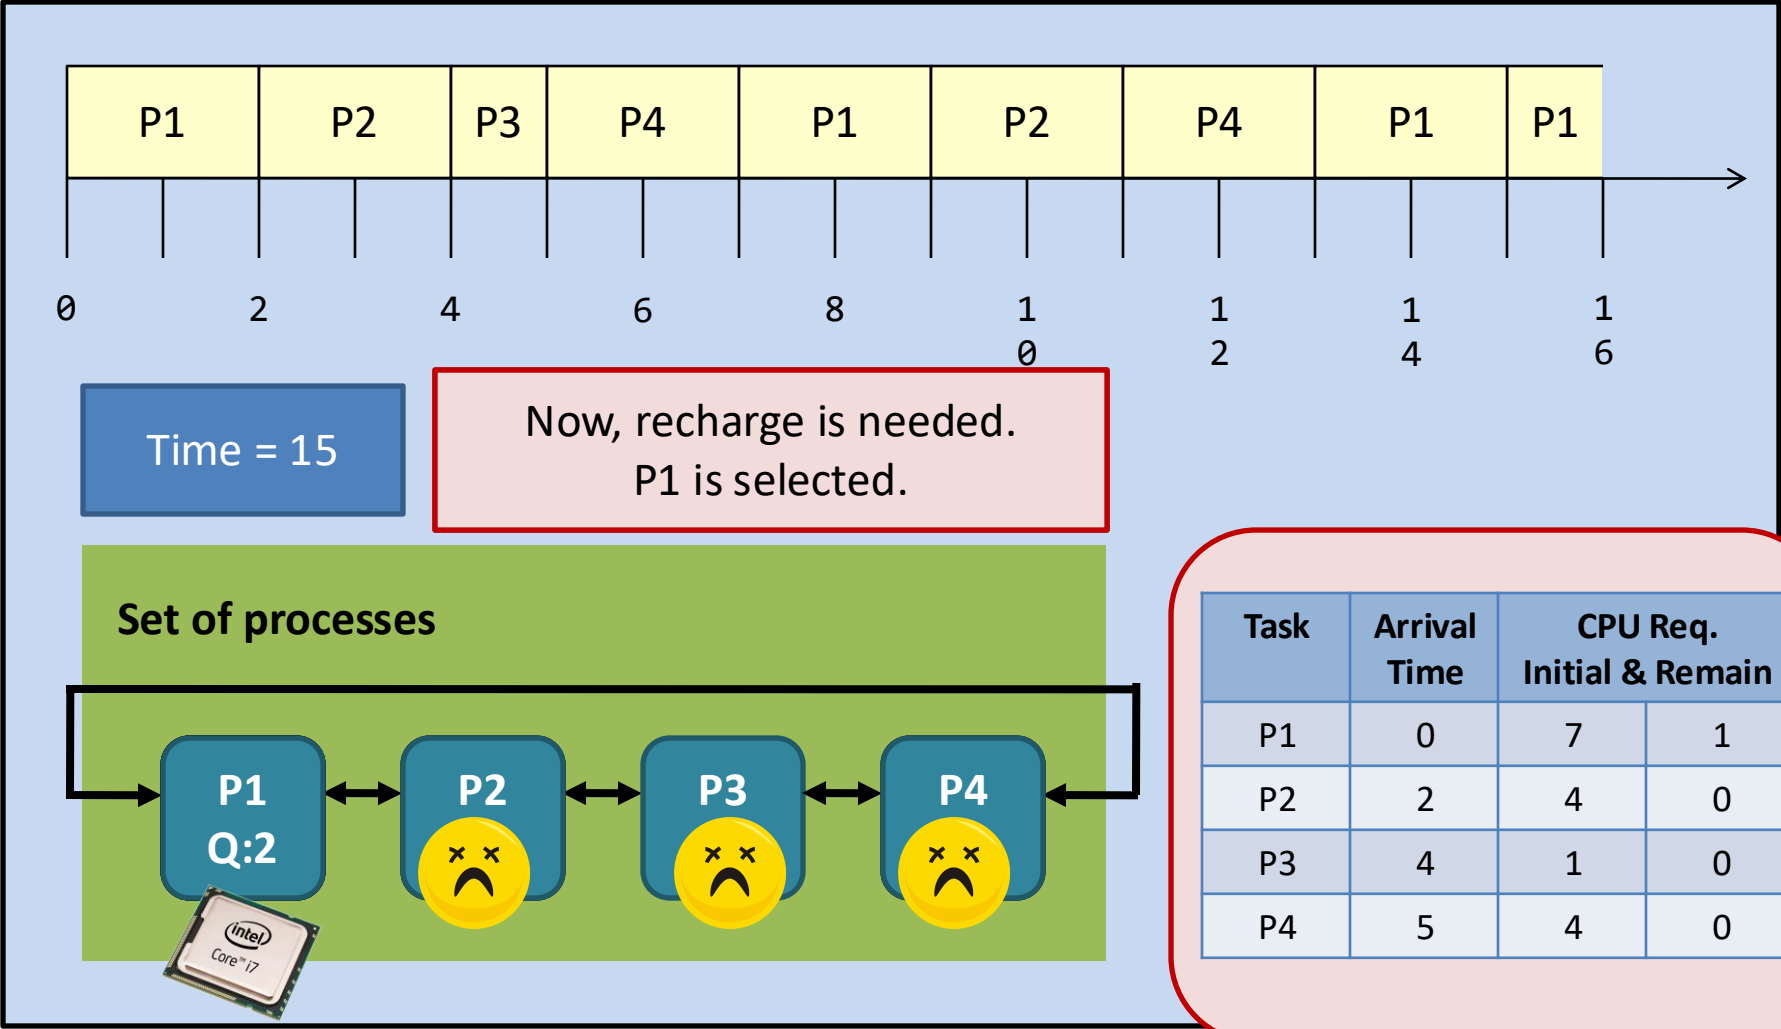
Round-robin

Animation; don't print




Round-robin

Animation; don't print



Round-robin



Waiting time:

$P1 = 9; P2 = 5; P3 = 0; P4 = 4;$

$Average = (9 + 5 + 0 + 4) / 4 = 4.5$

Turnaround time:

$P1 = 16; P2 = 9; P3 = 1; P4 = 8;$

$Average = (16 + 9 + 1 + 8) / 4 = 8.5$

Task	Arrival Time	CPU Req. Initial & Remain	
P1	0	7	0
P2	2	4	0
P3	4	1	0
P4	5	4	0

RR VS SJF

	Non-preemptive SJF	Preemptive SJF	RR
Average waiting time	4	3	4.5 (largest)
Average turnaround time	8	7	8.5 (largest)
# of context switching	3	5	8 (largest)



So, the RR algorithm gets all the bad! Why do we still need it?

The responsiveness of the processes is great under the RR algorithm. E.g., you won't feel a job is "frozen" because every job gets the CPU from time to time!

Priority Scheduling

- A task is given a priority (and is usually an integer).
- A scheduler selects the next process based on the priority
- New process arrival triggers a new selection
 - Same priority? RR, SJF, etc.

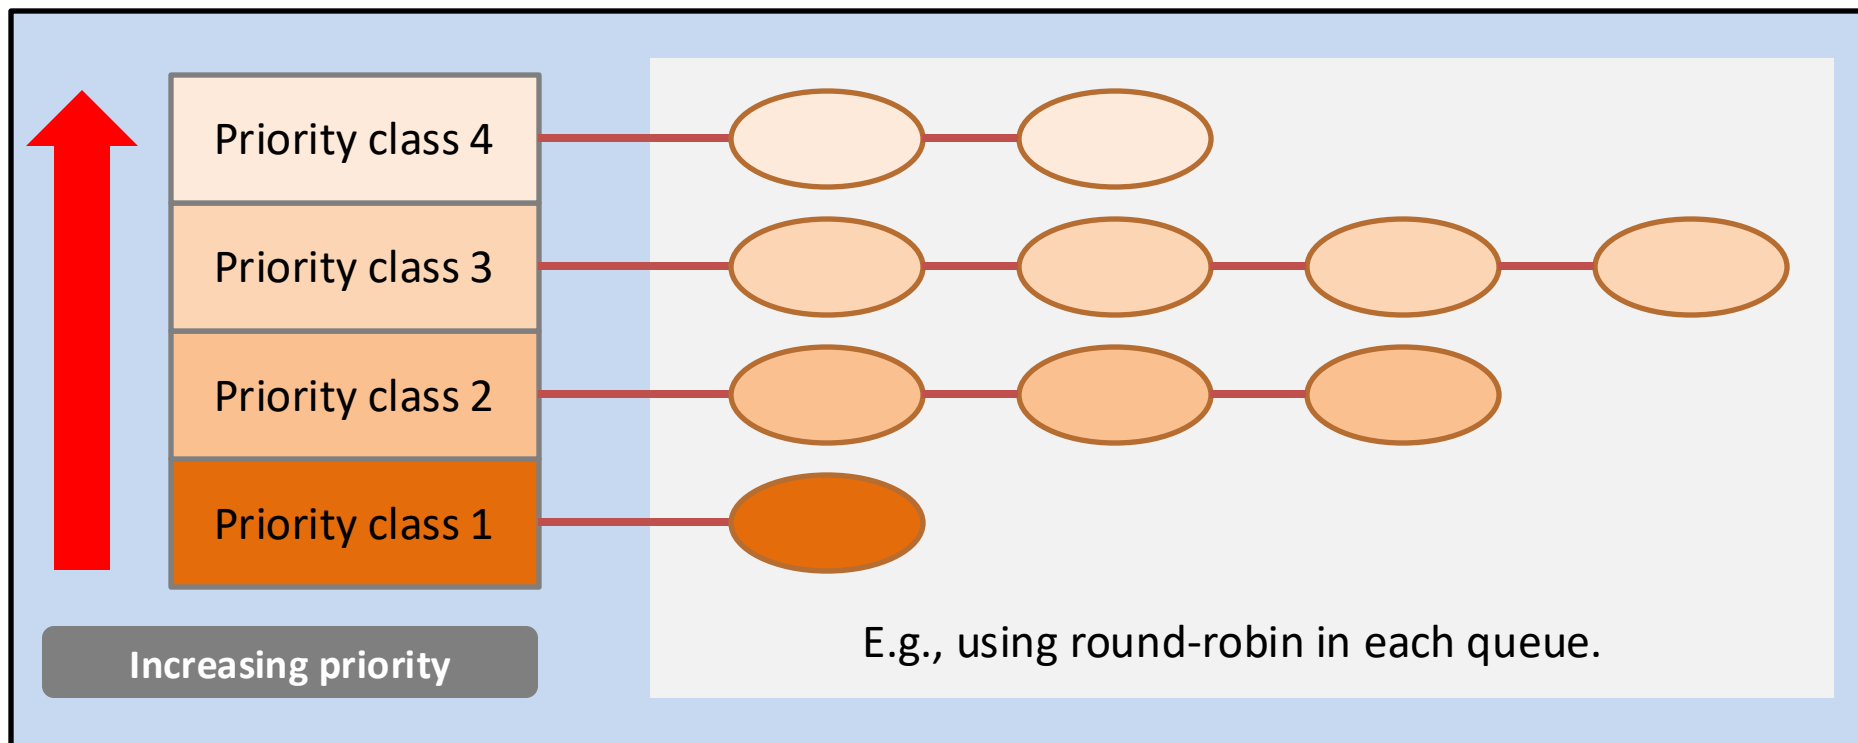
2 Classes	
Static priority	Dynamic priority
Every task is given a fixed priority.	Every task is given an initial priority.
The priority is <u>fixed</u> throughout the life of the task.	The priority is <u>changing</u> throughout the life of the task.

If a task is preempted in the middle

- Note:
 - it has already been dequeued
 - Re-enqueue back to the queue
 - Quantum preserved / recharge?
 - Depends
 - Preserved: need more book keeping
 - Recharge: easy (assumed in this course)

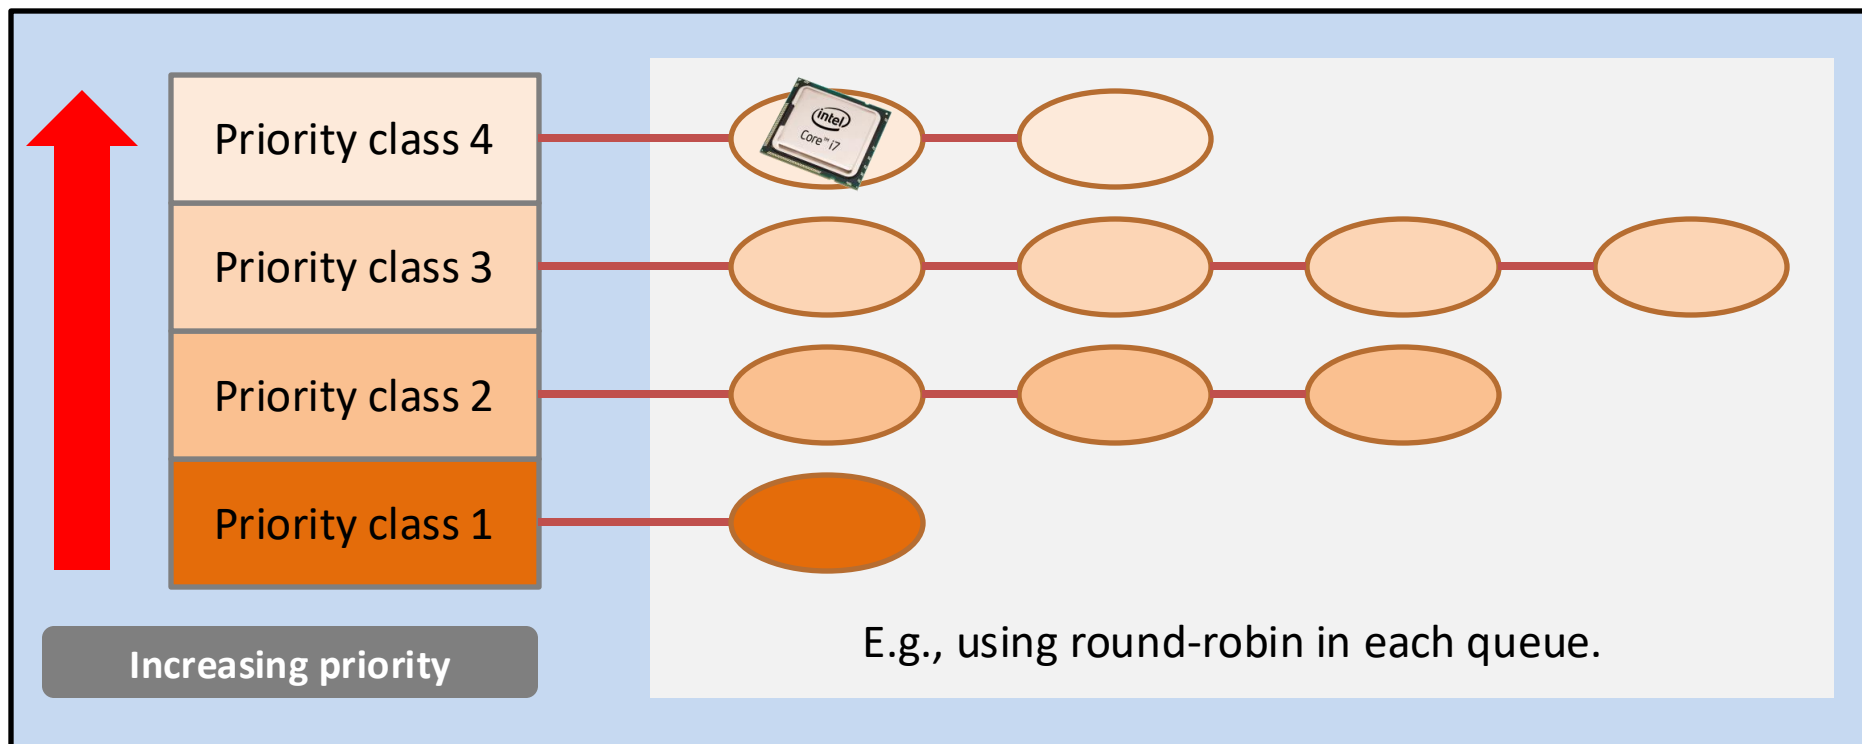
Static priority scheduling – an example

- **Properties:** process is assigned a fix priority when they are submitted to the system.
 - E.g., Linux kernel 2.6 has 100 priority classes, [0-99].



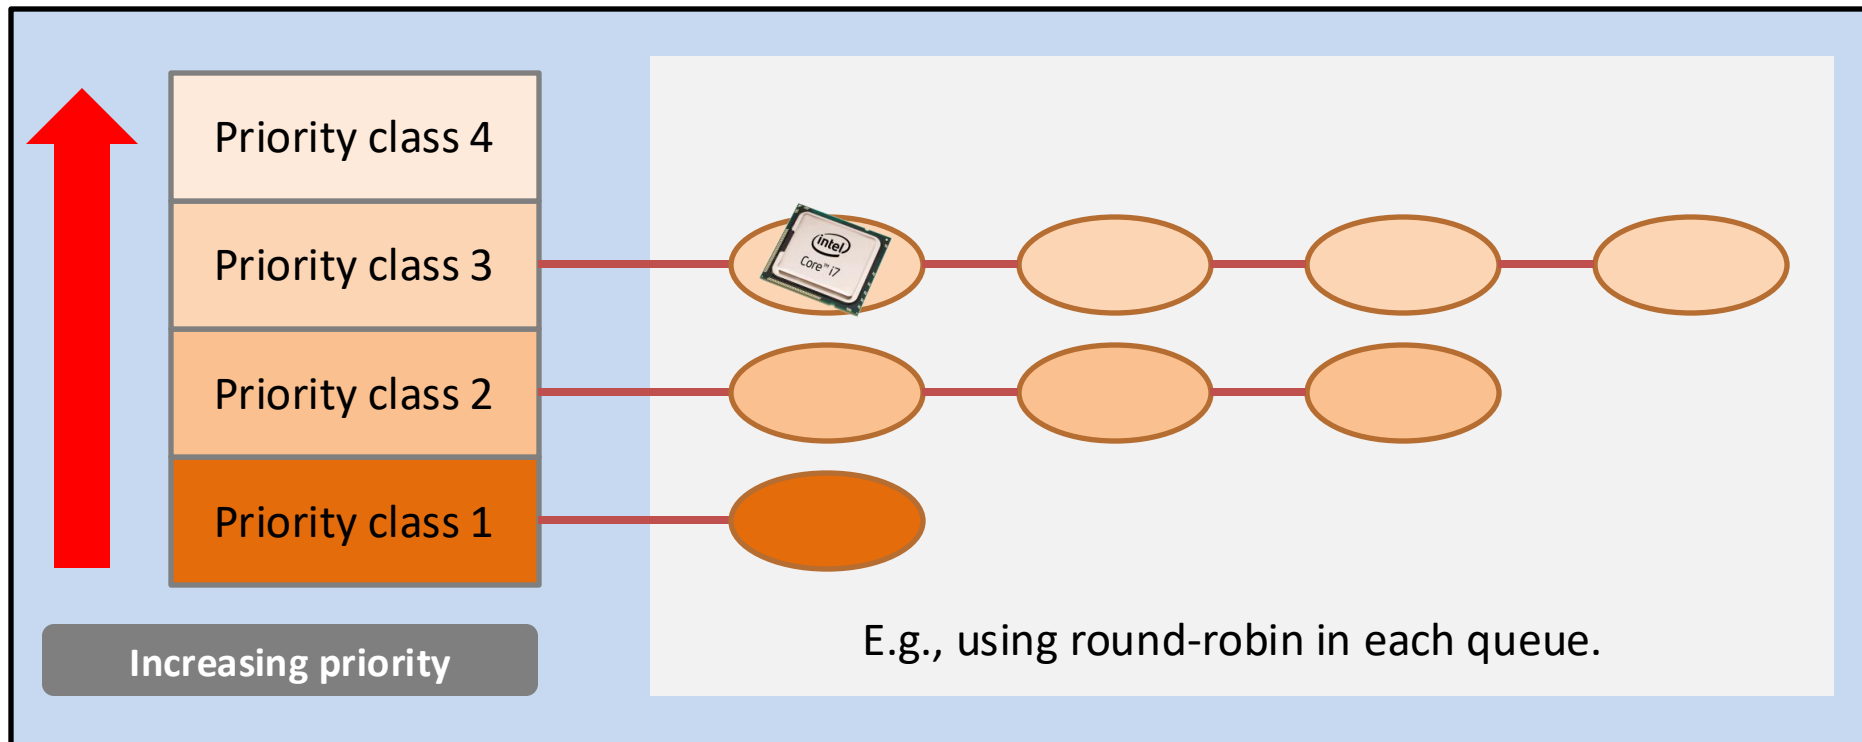
Static priority scheduling – an example

- The highest priority class will be selected.
 - The tasks are usually short-lived, but important;
 - To prevent high-priority tasks from running indefinitely.



Static priority scheduling – an example

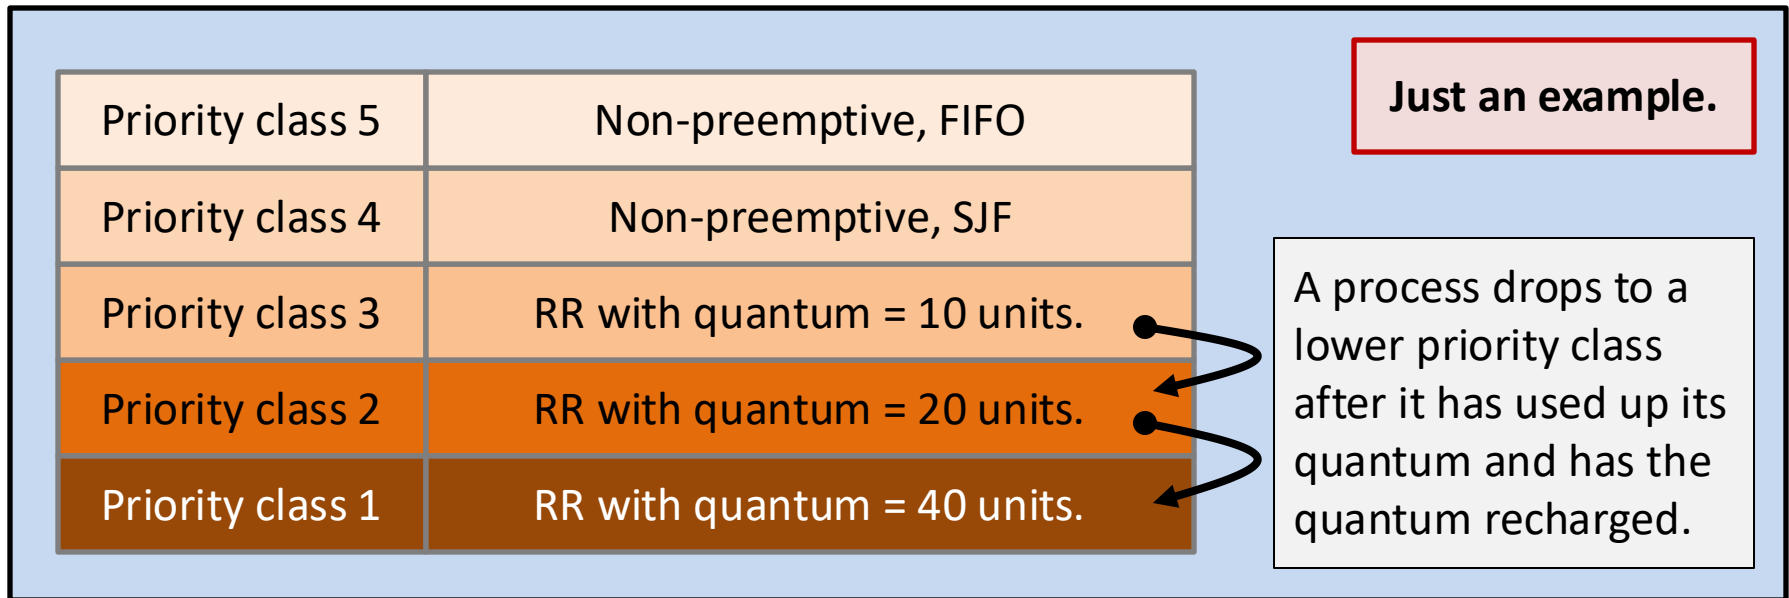
- Lower priority classes will be scheduled only when the upper priority classes has no tasks.



Multiple queue priority scheduling

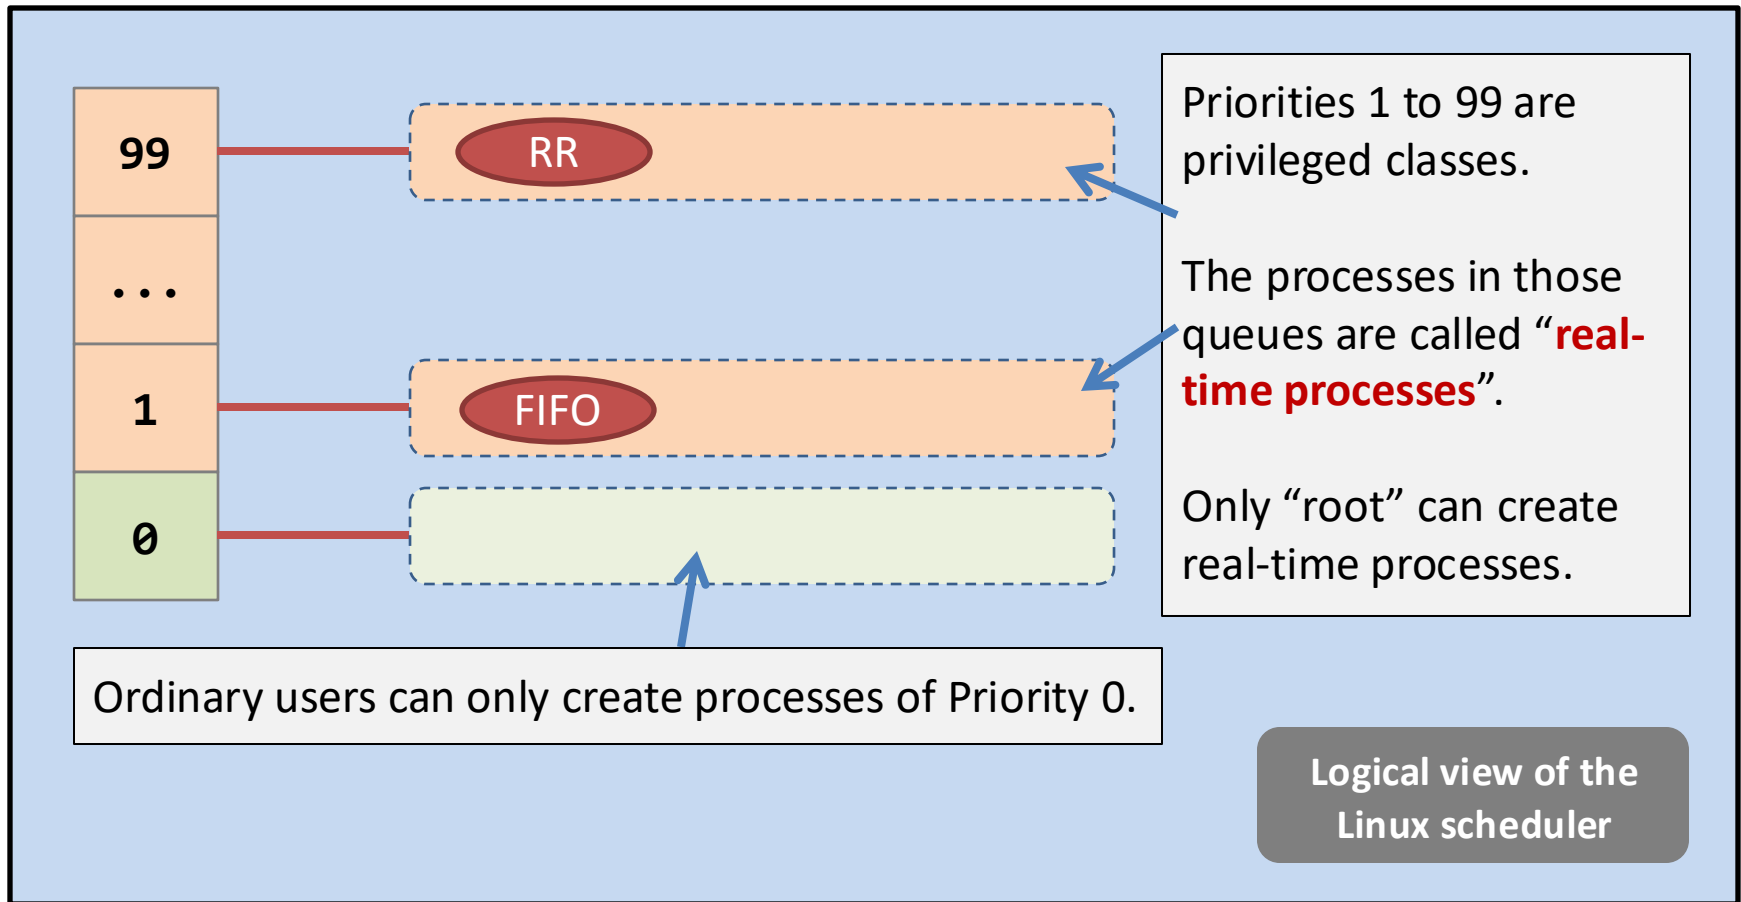
- **Definitions.**

- It is still a priority scheduler.
- But, at each priority class, **different schedulers** may be deployed.
- The priority can be a mix of static and dynamic.



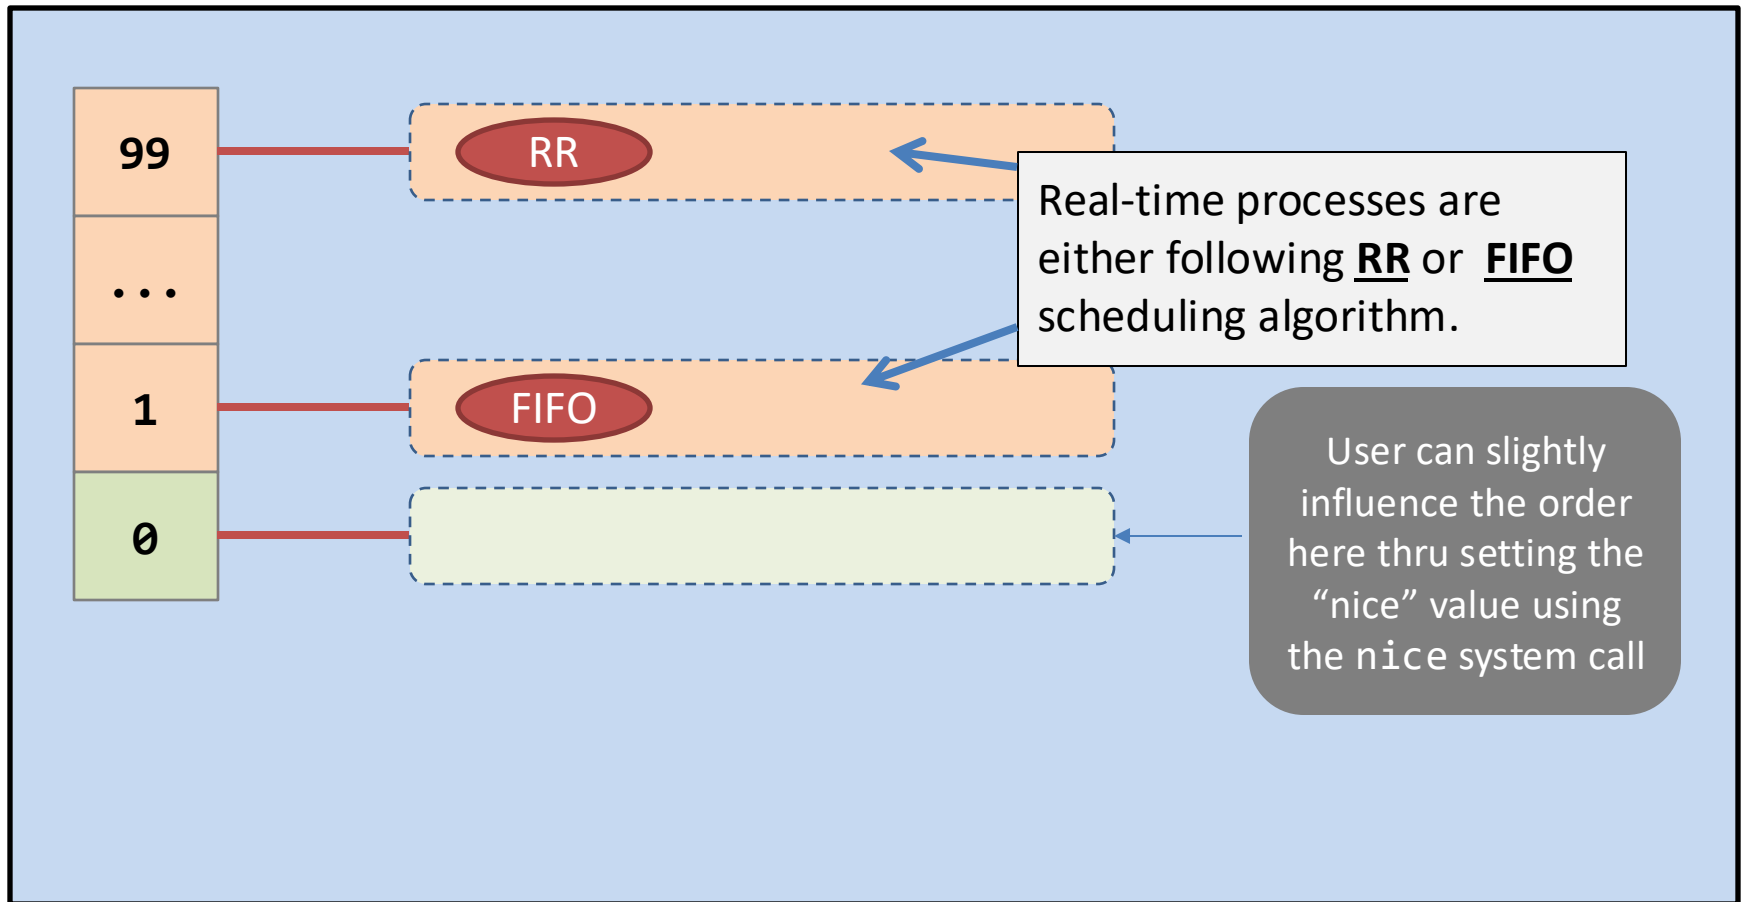
Multiple queue priority scheduling

- **Real example, the Linux Scheduler.**
 - A multiple queue, (kind of) static priority scheduler.



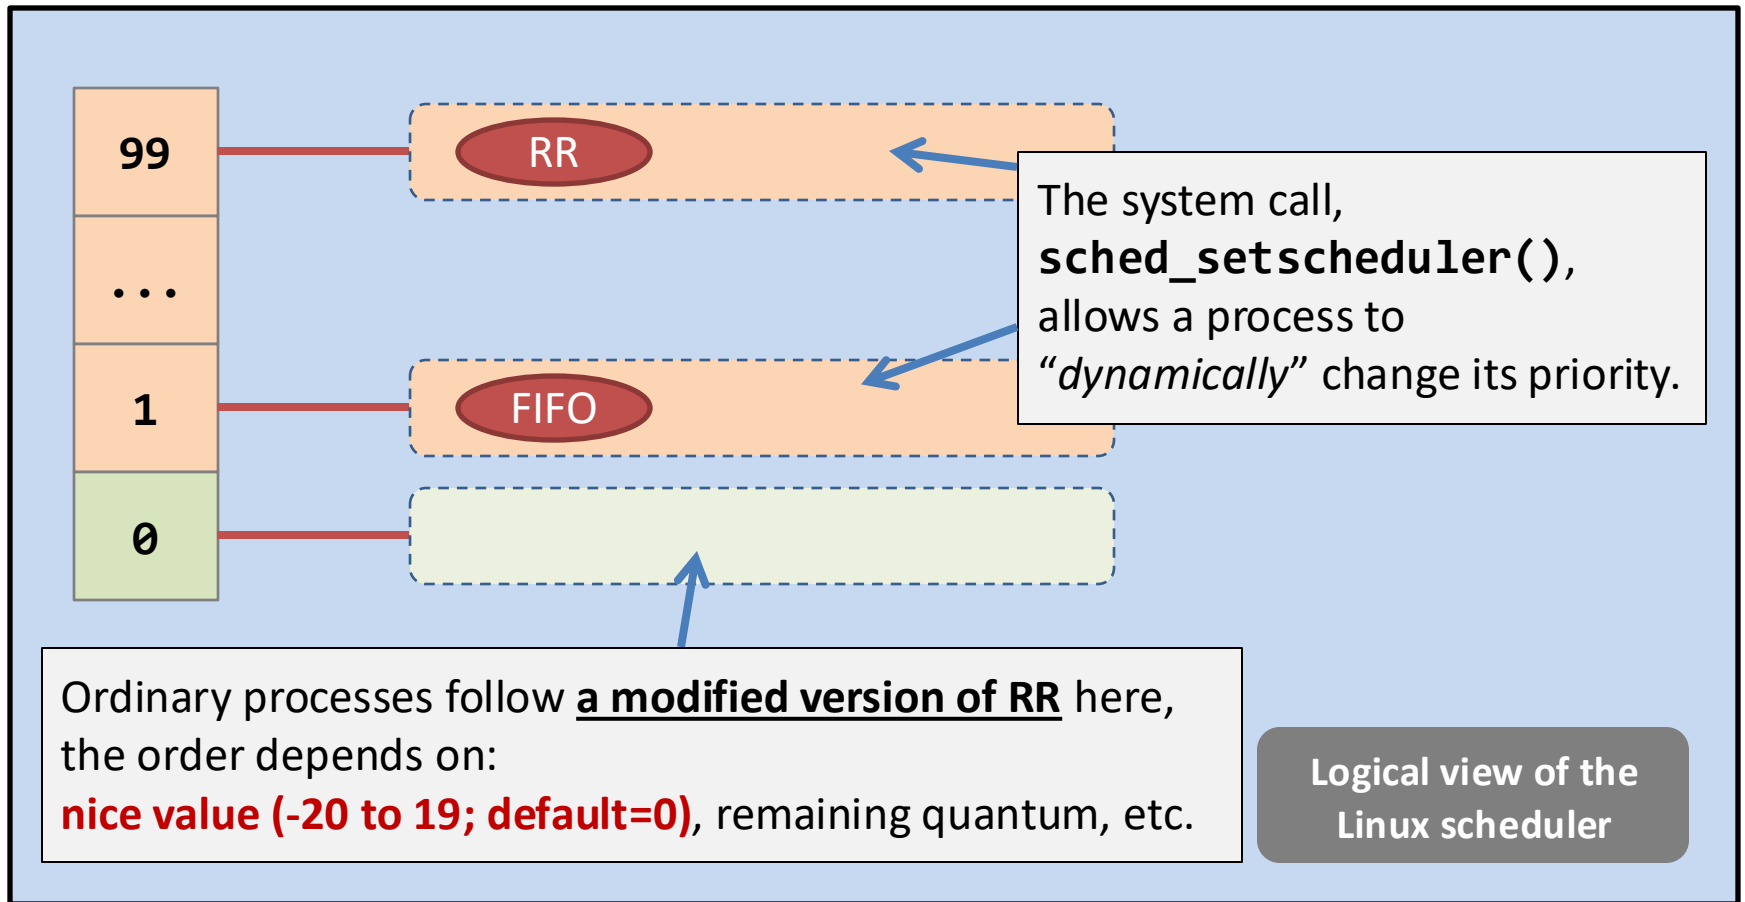
Multiple queue priority scheduling

- **Real example, the Linux Scheduler.**
 - A multiple queue, (kind of) static priority scheduler.



Multiple queue priority scheduling

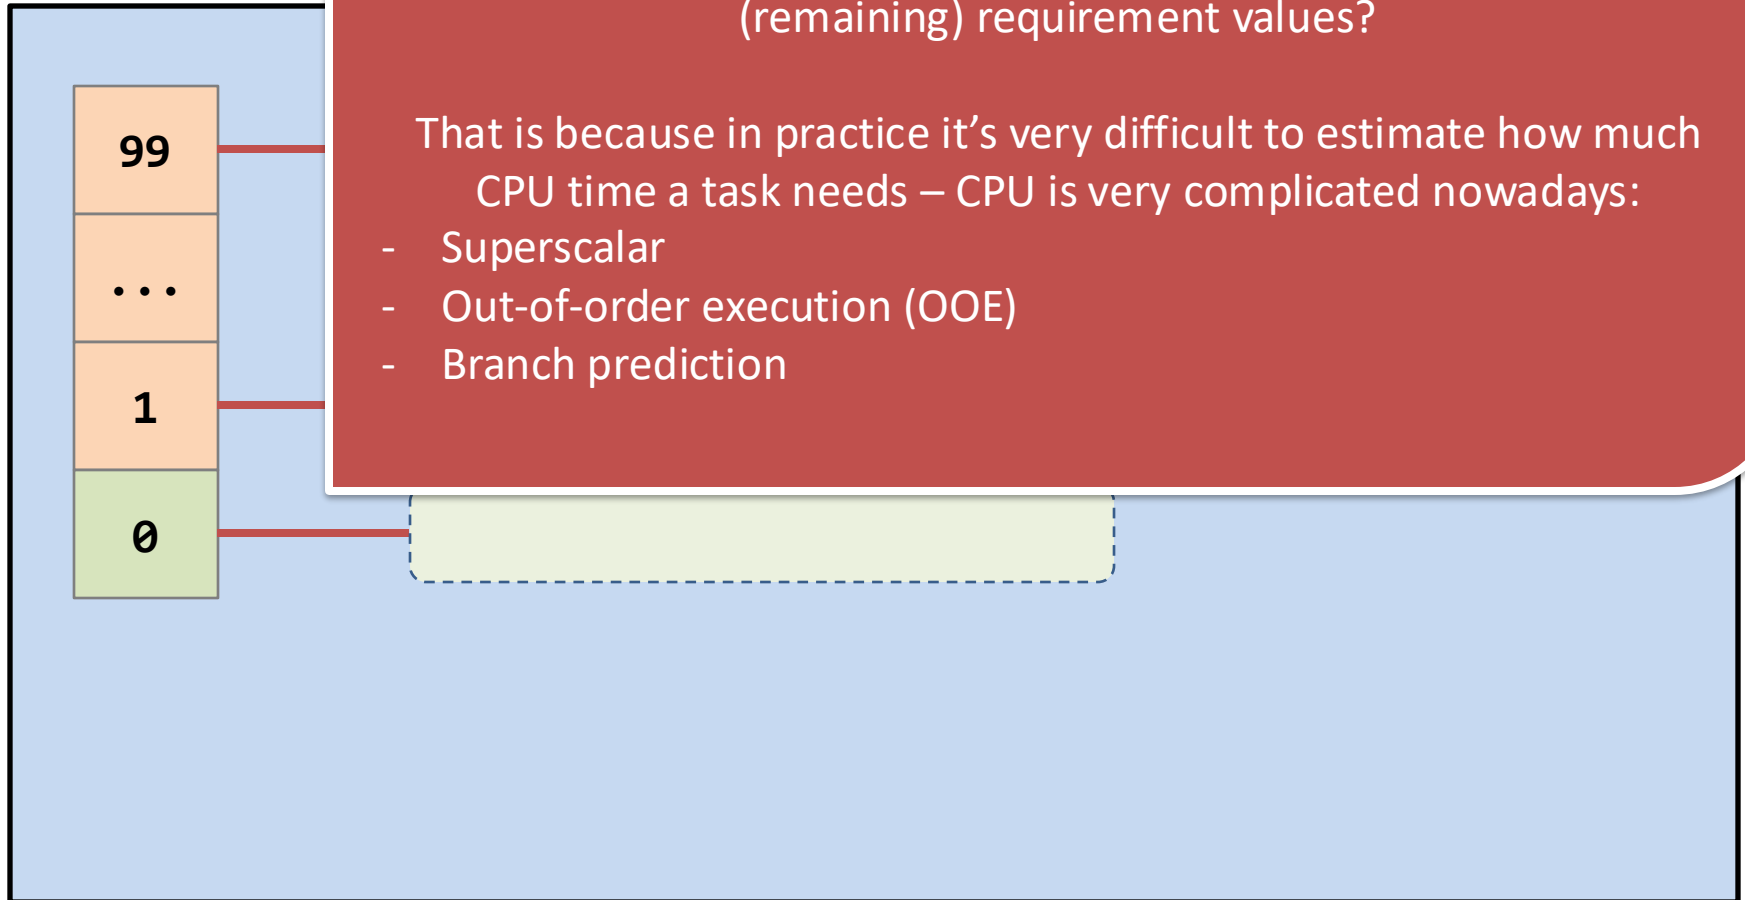
- **Real example, the Linux Scheduler.**
 - A multiple queue, (kind of) static priority scheduler.



Multiple queue priority scheduling

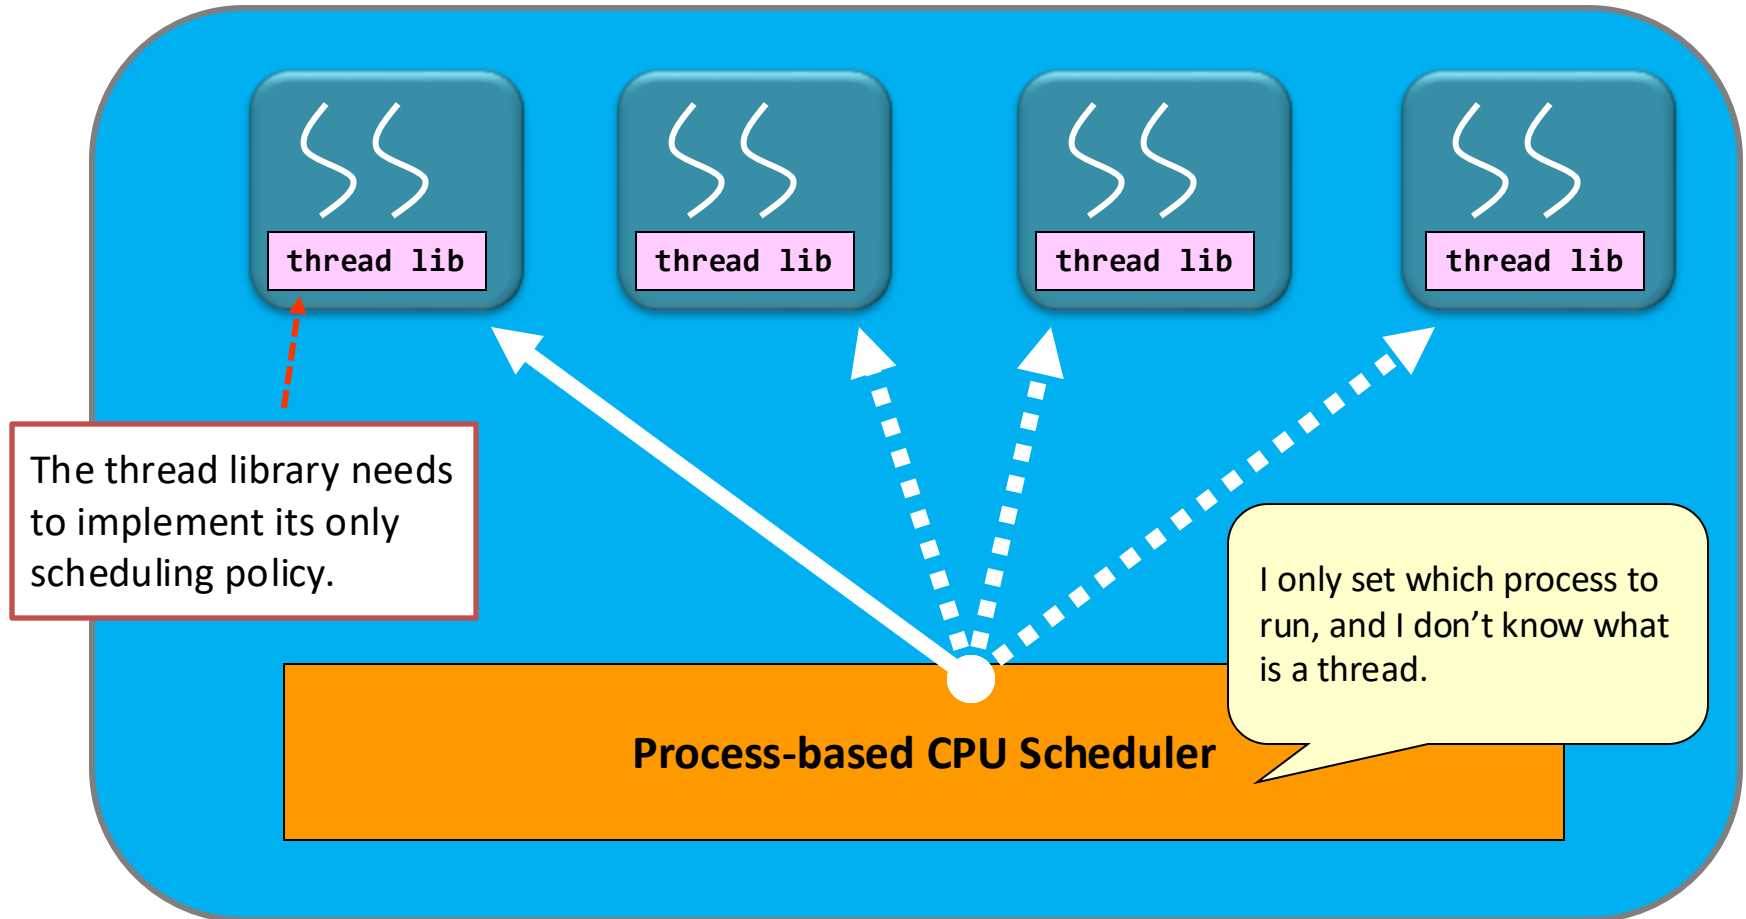
- **Real example the Linux Scheduler**

- A multiprocessor



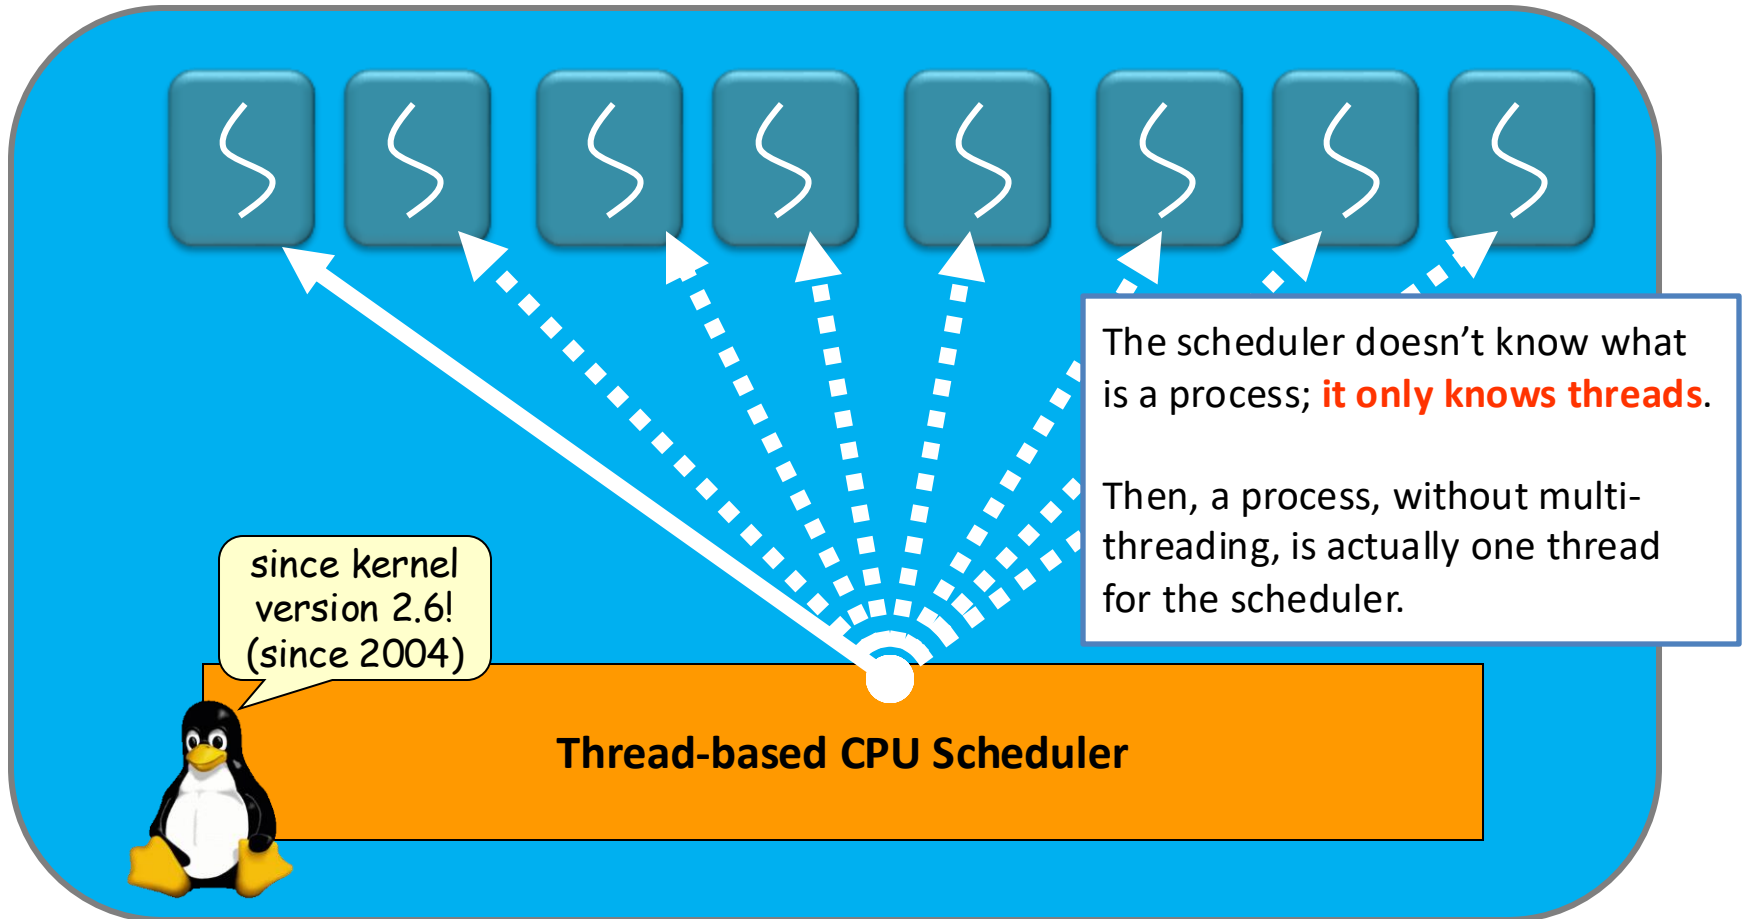
Thread Scheduling

- If a scheduler only interests in **processes**...



Thread Scheduling

- If a scheduler only interests in **threads**...



Scheduling for modern servers

- Symmetric multiprocessing (SMP)
 - All processes/threads **share a common ready queue**
 - ➔ **Race Condition**
 - Scheduling:
 - Each processor/core scheduler examines the common ready queue and selects a process to execute
 - Affinity Scheduling:
 - Soft affinity: **attempt** to keep the same process/thread on the same core
 - vs. Hard affinity: use `pthread_attr_setaffinity_np()` to manually bind a thread to a specific core in your pThread program
 - Especially important when we are reaching NUMA world