


Operating Systems

Eric Lo

14 – I/O

I/O device controller

- Block devices vs. Character devices
- E.g., Character device controller's tasks
 - Raw serial bit stream \leftrightarrow bytes for controller's registers
 - Checksum error handling
 - Built-in registers and data buffer for communicating with the CPU

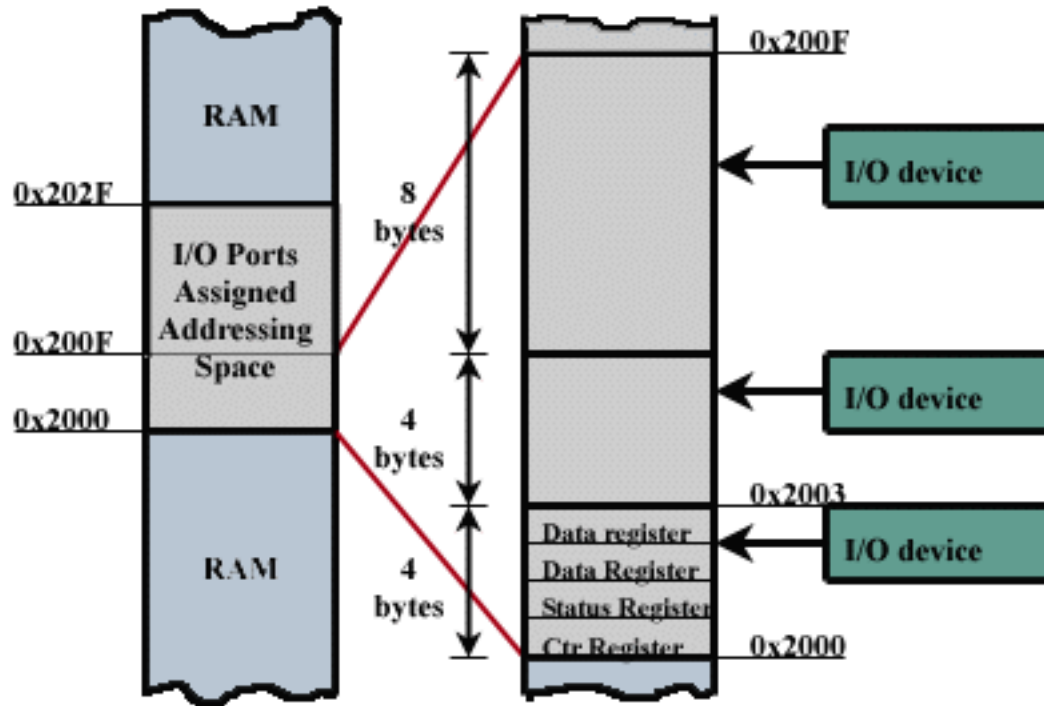


I/O device controller

- Built-in registers and data buffer for communicating with the CPU
 - Direct I/O
 - CPU places data into (and read data from) register/data buffer directly through instructions like
 - **IN REG, PORT** //read the value at PORT of the device to CPU register REG
 - Memory-mapped I/O
 - Map the device's registers and data buffer to the memory space
 - » Both kernel (driver) and user-space can do that

Memory-mapped I/O

- Kernel Space



Ack: McGill

Device Drivers

- Follow the standard programming interface offered by the OS
 - After all, the programming interface for all printers are the same, right?
 - The device's manufacturer follows the interface to implement a driver that includes code to read and write the device's registers
- E.g., Microsoft Windows Interface:
 - `printer.print(Byte b);`
- Canon-C123Driver implements `printer.print(Byte b){`
 `MOV CanonPrinterRegister b //none of the OS's business here`
}
- HP-T125Driver implements `printer.print(Byte b){`
 `MOV EpsonPrinterRegister b //none of the OS's business here`
}

I/O Communication protocol

- **Polling**
 - CPU puts one byte to (e.g., printer) device's **DATA** register
 - CPU keeps polling the device's **READY** register in order to put the next byte
- **Interrupt**
 - CPU waits for interrupt and does something else in between
- **Direct Memory Access (DMA)**
 - DMA controller on system bus
 - Offloading the per-byte **polling/interrupt** job from CPU to DMA controller

DMA

- OS sets up a memory-mapped region (i.e., memory-mapped I/O) in the kernel space
- OS pins those memory-mapped pages in memory
 - So that those pages are kept for the I/O communication
 - Won't let paging swap them away

Direct Memory Access

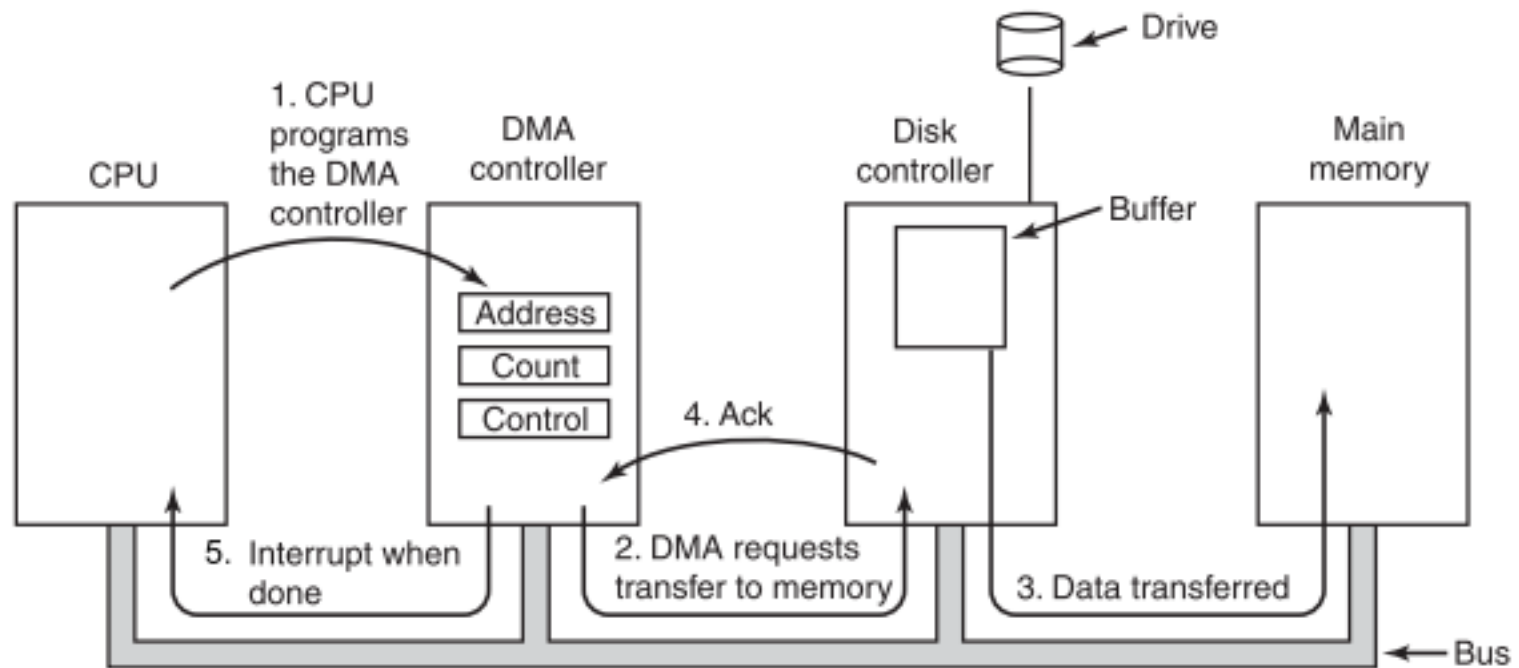
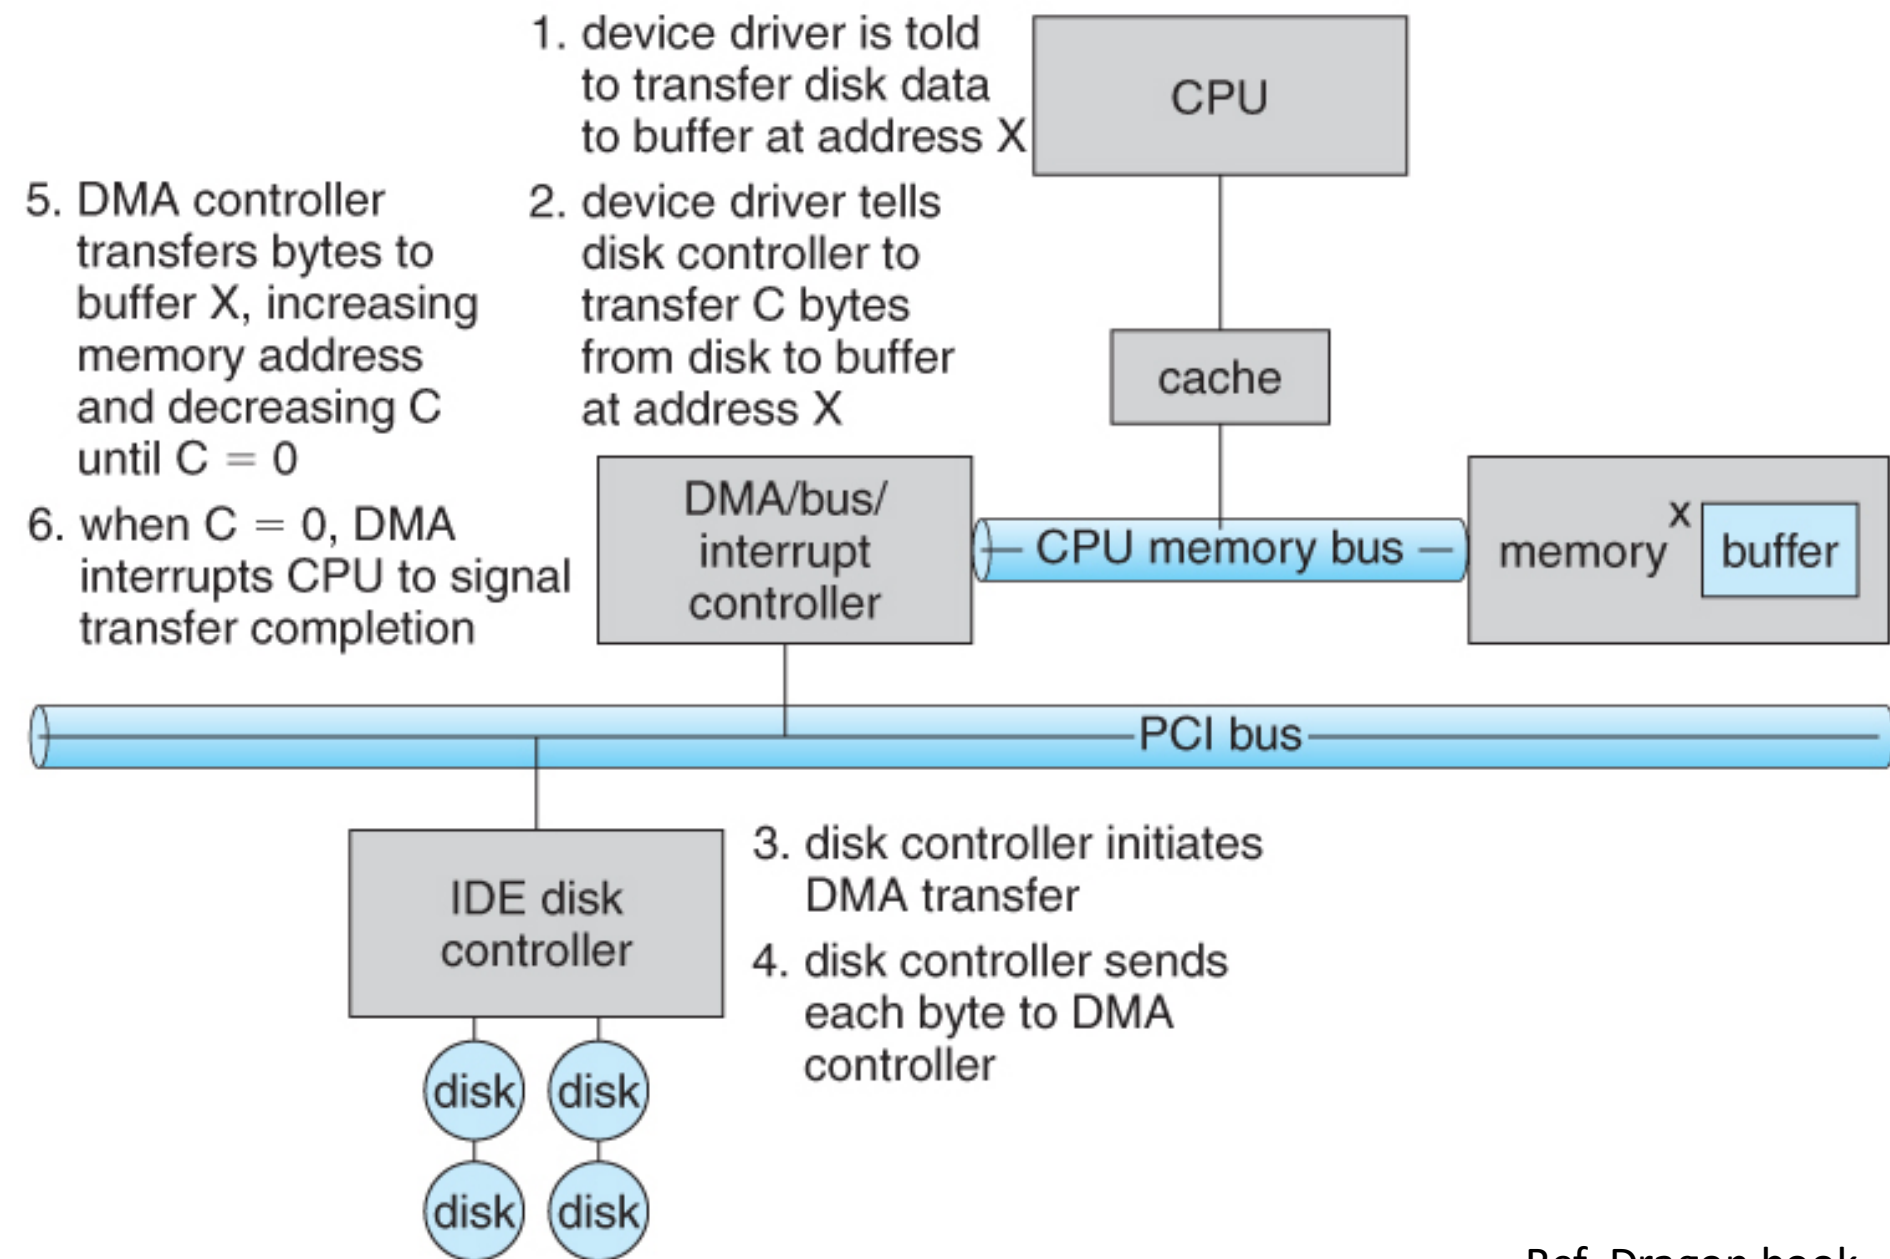


Figure 5-4. Operation of a DMA transfer.

Direct Memory Access



Print Spooling

- A printer obviously shouldn't be a concurrent sharable "file" (character special file)
 - No >1 process can print to the printer together
- A user process exclusively opens the printer but does nothing for hours?
 - No other processes can print!
- Create a root level printer **daemon** process, and a **spooling** directory
 - To print a file, a user process has to
 - generate the entire file to be printed
 - put it to the spooling directory
 - Only the daemon can access the printer device file

Wiki:

- documents formatted for printing are stored in a queue at the speed of the computer, then retrieved and printed at the speed of the printer.
- Multiple processes can write documents to the spool without waiting, and can then perform other tasks, while the "spooler" process operates the printer

BTW, all daemon background processes have a naming convention

inetd, httpd, nfsd, sshd, named, and lpd

Memory-Mapped I/O on user-process

- Can also be used in user-space
- Set up an mmap to link to a file
- Write to an address = write to a file
- Advantage?

