

Operating Systems

Eric Lo

15 – Synchronization II

In Synchronization I

Achieving Mutual Exclusion – Locking approach

- Locking by Disabling Interrupt
 - This facility is not available to user level program though
 - Doable for kernel level but
 - Multi-core complication
 - It is possible that another core modifying the shared object in the memory
 - Multi-core CPU bundles with **atomic instructions** to make sure one core's modification is atomic
 - All subsequent discussions assume all instructions are atomic
 - **e.g., x++ is atomic**
 - Details of atomic instruction will be covered in topic "Synchronization II"

Atomic Instructions

- E.g., `test_and_set()` & `compare_and_swap()`
- E.g., Two `test_and_set(x)`s are concurrently executed by **different CPUs/cores** on the same memory location `x`
 - `x` will get updated by one by one
 - Can be used through newer versions of C/C++
 - E.g., C11 or above

p.s. C++11 has a `Thread` library, even easier to write multi-threading; that library is written using `Pthread`; programs using that are less portable than using `Pthread` directly

Achieving Mutual Exclusion - Revisit

- Lock-based
 - Spin-based lock
 - E.g., `pthread_spin_lock`
 - What is inside?
 - » Spin on 1 shared variable + **atomic instruction** to update it
 - » Spin on 2 shared variables + **atomic instruction** to update them (~Peterson)
 - » Spin on shared variables + **atomic instruction** to update them in an even smarter way
 - (=Ticket, MCS algorithm, etc.)
 - Sleep-based lock
 - E.g., POSIX semaphore, `pthread_mutex_lock`
 - What is inside?
 - » `yield()` + **atomic instructions** to update the extra shared variables
 - » Need to work with kernel's scheduler -> context switch to kernel space -> expensive ☹
 - » No waste the CPU to spin but let others use it (if any) ☺
 - Lock-free
 - No such a call like `pthread_lock_free`
 - Let the threads update the shared variables freely
 - Right before leaving the CS
 - use atomic-instruction to test if there is any race condition
 - No? set a new value for the shared variable
 - Yes? Restart and Retry

Basic Spinlock implementation using `test_and_set()`

```
Initialize Atomic lock = 0;
```

```
while (TestAndSet(&lock) == 1)
    ; //do nothing, spinning; this loop quits when 0==1 //spin_lock
```

when reaching this point, =you get the lock

lock = 1

*But others should
be spinning*

```
<critical
.....
section>
```

```
lock = 0; //release the lock //spin_unlock
```

```
<remainder
section>
```

```
//pseudocode (indeed a single atomic instruction)
boolean test_and_set(boolean *lock){
    boolean initial = *lock; //original value
    *lock = 1; //set
    return initial; //return the original value!
}
```

compare_and_swap()

- Maurice Herlihy was awarded the Edsger W. Dijkstra Prize for his seminal 1991 paper "Wait-Free Synchronization"
 - the incapability of TAS when synchronizing >2 threads

```
//pseudocode (indeed a single atomic instruction)
boolean cas(int *value, int expected, int new) {
    if (*value != expected) {
        return false;
    }
    value = new;
    return true;
}
```

Basic Spinlock implementation using CAS

- lock – a shared variable

```
Initialize _Atomic lock = 0;
thread-function() {
while (!compare_and_swap(&lock, 0, 1))
    { ; }//
<critical
    .....
section>
lock = 0; //release the lock
<remainder
section>
}
```

This can be wrapped as 'spin-lock()'



//spin lock

//unlock

```
//pseudocode (indeed a single atomic instruction)
boolean cas(int *value, int expected, int new) {
    if (*value != expected) {
        return false;
    }
    *value = new;
    return true;
}
```

Check your knowledge

- Do you see pthread_mutex_lock?
- Is it lock-free?

```
Initialize _Atomic lock = 0;
thread-function() {
    while (!compare_and_swap(&lock, 0, 1))
        { ; }// //spin lock
    <critical
        .....
    section>
    lock = 0; //release the lock
    <remainder
        section>
}
```

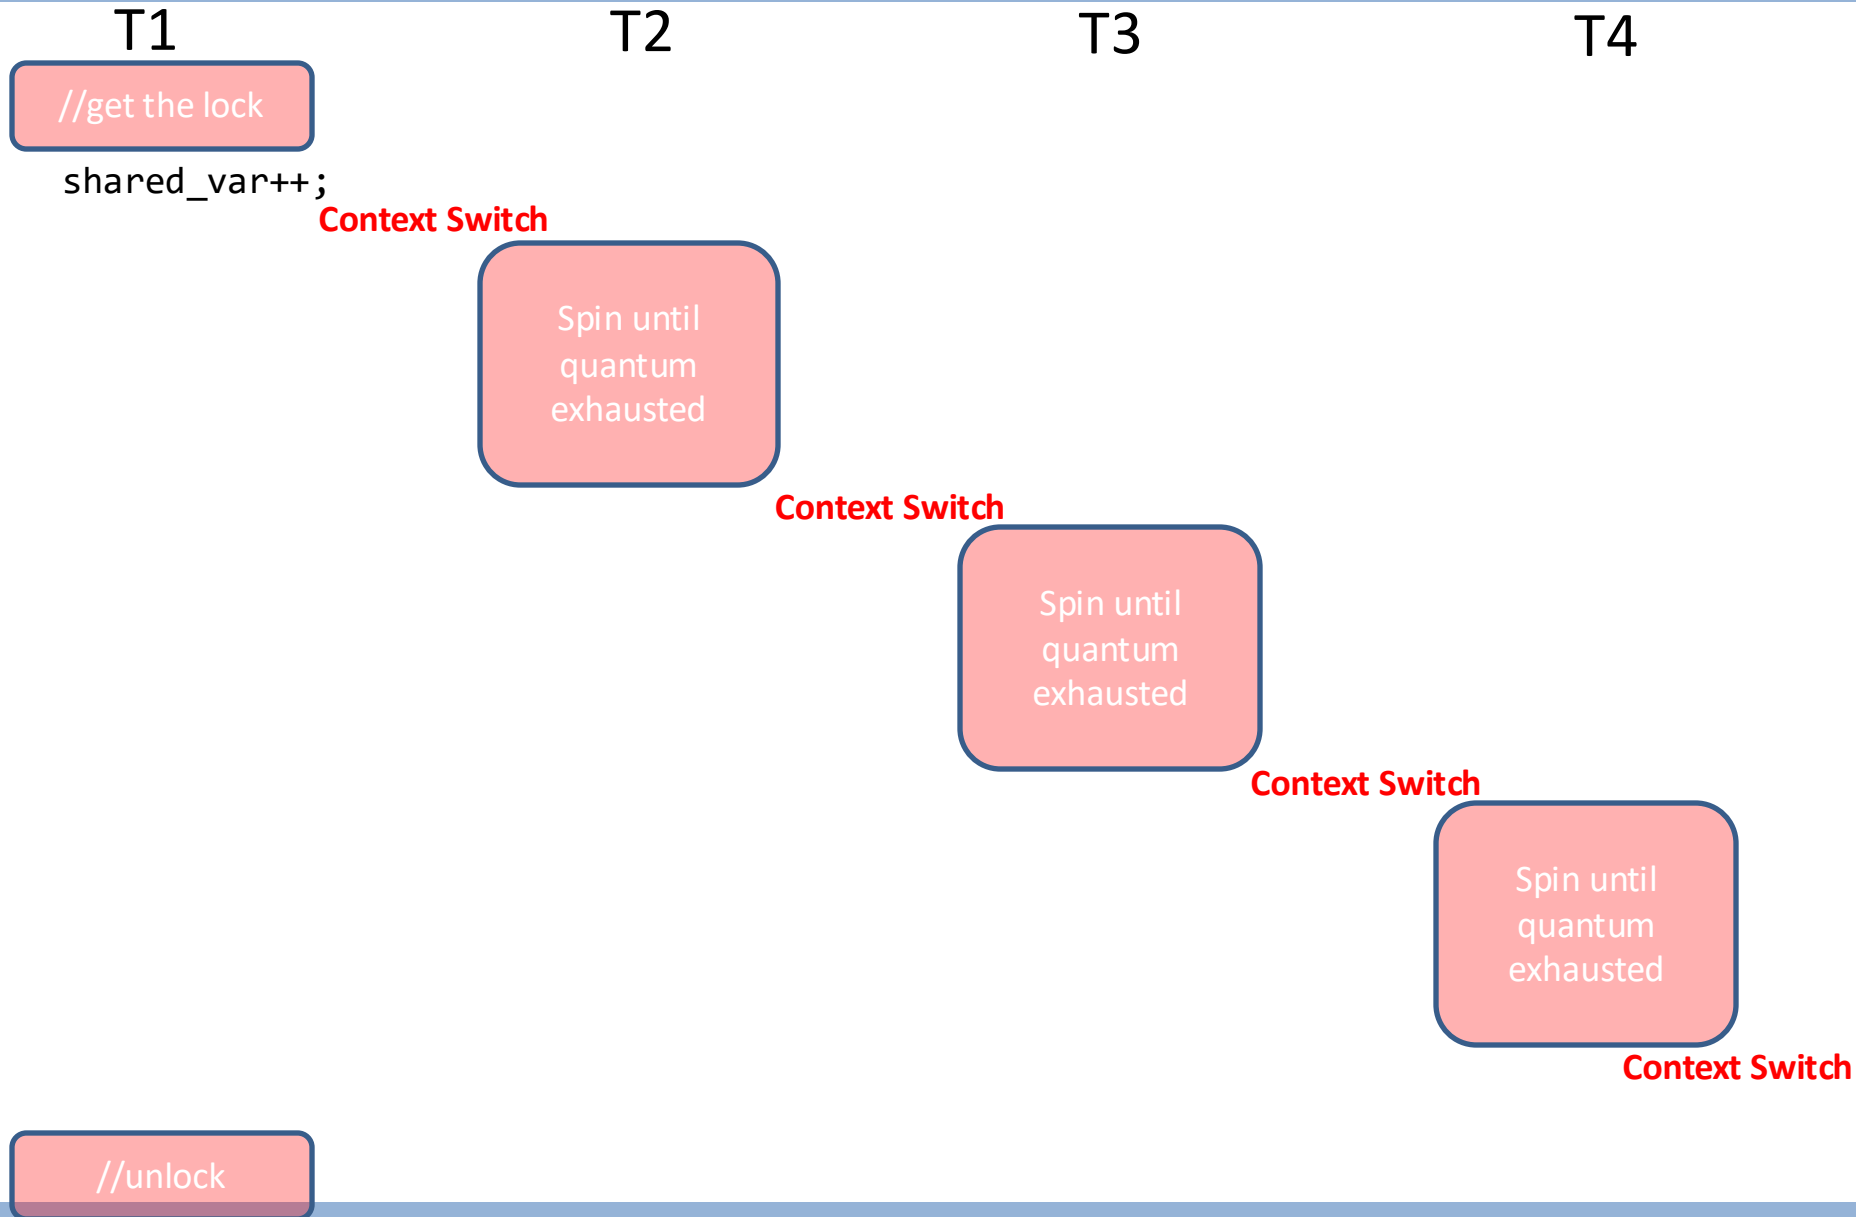

Counting: Lock-based vs. Lock-free

```
Initialize _Atomic lock = 0;
thread-function() {
    while (!cas(&lock, 0, 1)) //spin lock
        { ; }//
    <critical
        shared_var++;
    section>
    lock = 0; //release the lock
    <remainder
        section>
}
```

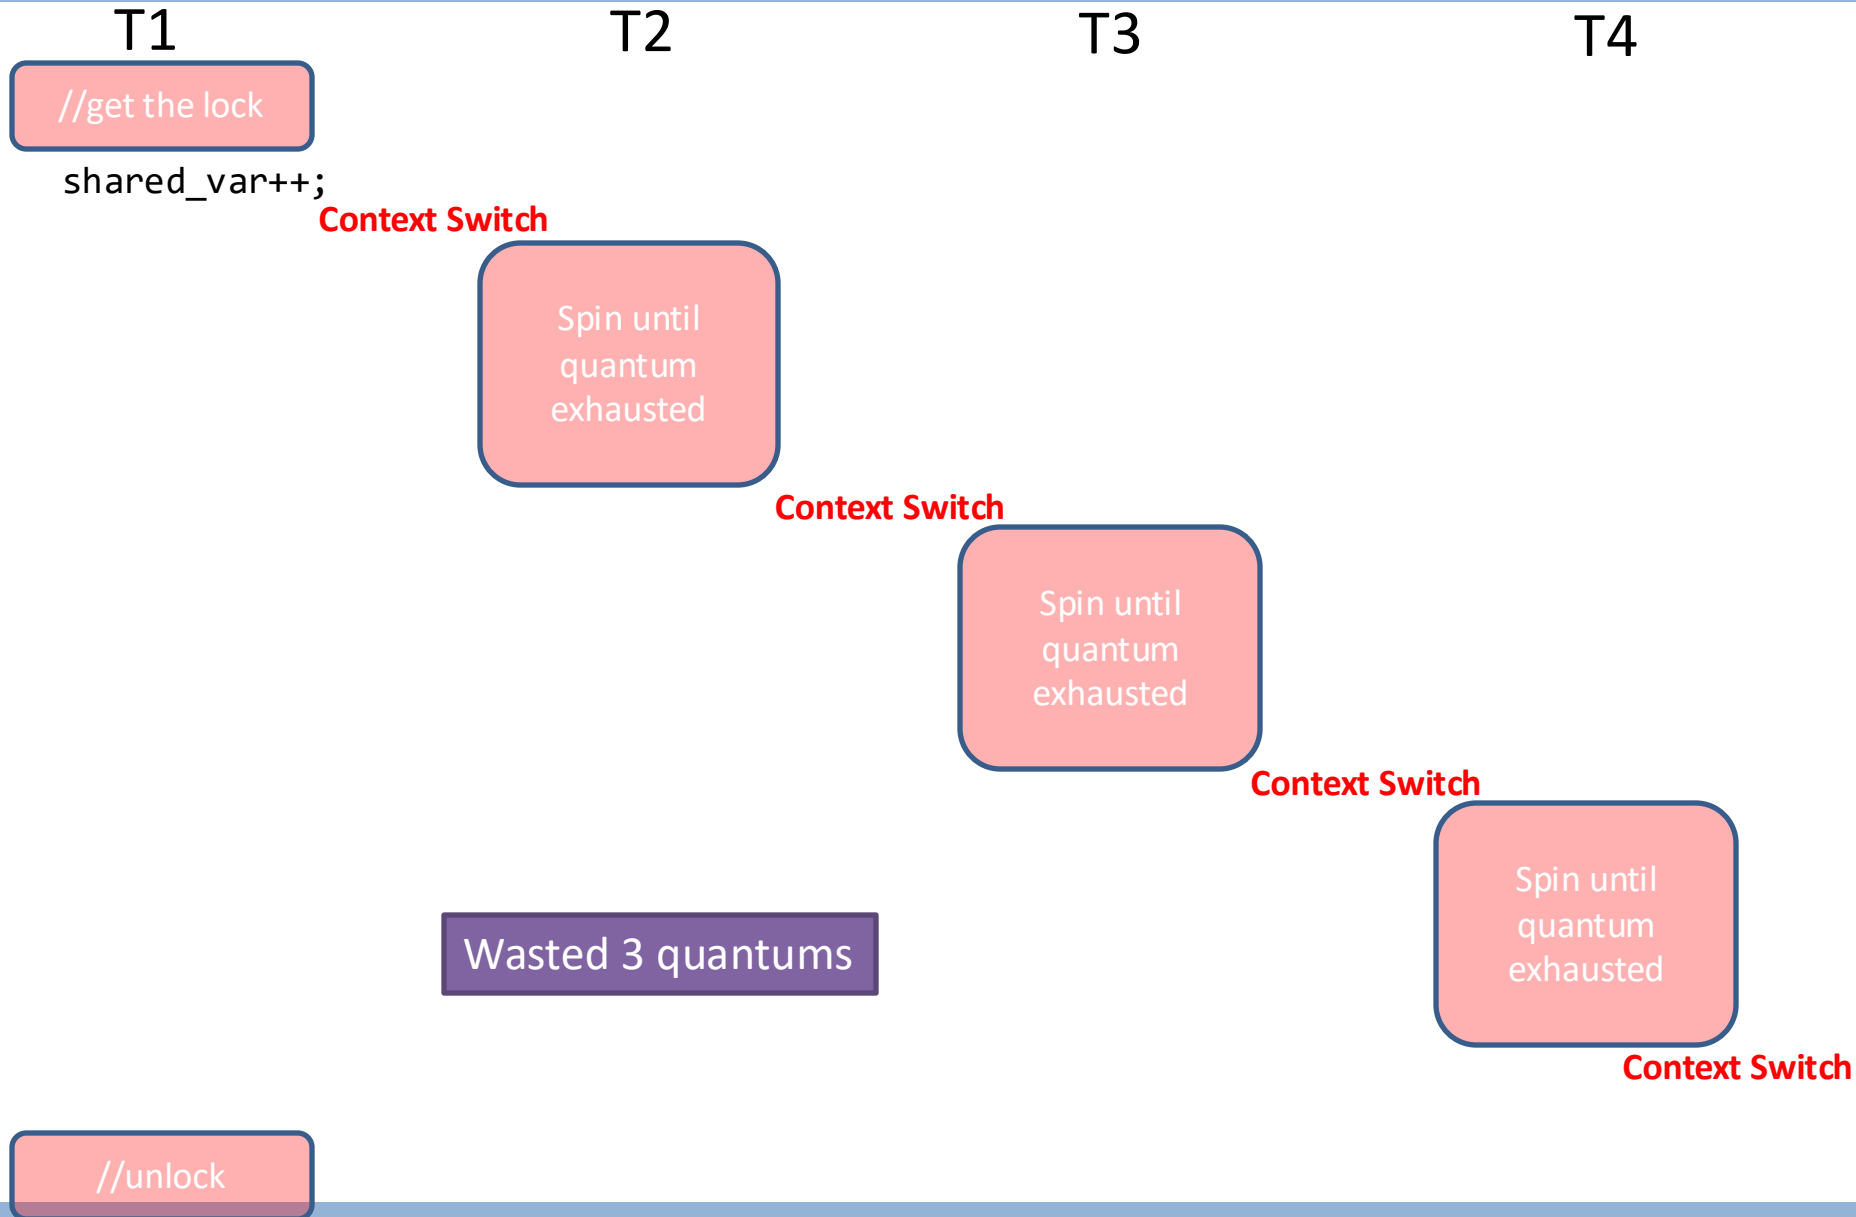
```
Initialize _Atomic lock = 0;
thread-function() {
    do {
        <critical
            pure = shared_var; //optimistic
                                //do the CS first
        section>
        while (!cas(&shared_var, pure, pure++))
            //redo if contaminated
        <remainder
            section>
    }
```

```
boolean cas(int *value, int expected, int new) {
    if (*value != expected) {
        return false;
    }
    value = new;
    return true;
}
```

Counting: Lock-based (4 threads; 1 core)



Counting: Lock-based (4 threads; 1 core)



Counting: Lock-free (4 threads; 1 core)

T1

T2

T3

T4

pure=10

Context Switch

pure=10
shared_var=11

Context Switch

pure=11
shared_var=12

Context Switch

pure=12
shared_var=13

Context Switch

cas(13, 10, pure++);
Contaminated! cas failed

RE-DO!

```
do {  
  <critical  
    pure = shared_var;           //optimistic  
                                //redo if contaminated  
  section>  
  while (!cas(&shared_var, pure, pure++))
```

Counting: Lock-free (4 threads; 1 core)

T1

T2

T3

T4

pure=10

Context Switch

pure=10
shared_var=11

Context Switch

pure=11
shared_var=12

Context Switch

pure=12
shared_var=13

Context Switch

cas(13, 10, pure++);
Contaminated! cas failed

RE-DO!

Wasted 1 quantum

```
do {  
  <critical  
    pure = shared_var;           //optimistic  
                                //redo if contaminated  
  section>  
  while (!cas(&shared_var, pure, pure++))
```

That's too good to be true...

- It could also be like this:

Counting: Lock-free (4 threads; 1 core)

T1

T2

T3

T4

pure=10

Context Switch

pure=10
shared_var=11
(success!)

Context Switch

pure=11

Context Switch

pure=11

Context Switch

cas(11, 10, pure++);
Contaminated! cas failed

RE-DO! Wasted 1 quantum

shared_var=12
(success!)

cas(12, 11, pure++);
Contaminated! cas failed

RE-DO! Wasted 1 quantum

cas(12, 11, pure++);
Contaminated! cas failed

RE-DO! Wasted 1 quantum

Lock-free vs Lock-based

- Lock-based incurs problems of:
 - Lock-holder sleeping
 - Lock-holder dies
 - System dies with it
 - Deadlock
 - System needs to detect and resolve it
- Lock-free is free of the above
 - But lock-free also has its problems, e.g., ABA
 - Very difficult to code
 - Usually the code added to deal with ABA problem introduces another instance of ABA problem ...



Lock-free needs to take care of the ABA problem

Time

Thread 1:
Pop(*lfstack){
 //Originally, lfstack->head=0x01
 orig = atomic_load(lfstack); //orig.head=0x01
 next.head = orig.head->next; //next.head=0x02

Context switch




Thread 2:
Pop(*lfstack) //free 0x01
Pop(*lfstack) //free 0x02
 //now, lfstack->head=0x03

Context switch

Thread 3:
Push(*lfstack, 4){
 Node *node = malloc(sizeof(Node)); //allocator reuses the same memory
 // block 0x01, so node=0x01
 ... //when thread 3 pushes **node** into the stack, it let
 //lfstack->head point to **node**, then, lfstack->head=0x01;
}

Context switch

Thread 1:
 CAS(lfstack, &orig, next); //here, it compares the content points to by
 //lfstack with the content of orig, that is, lfstack->head and orig.head, for equality.
 //For lfstack->head, which is 0x01, while for orig.head, which is also 0x01, they are
 //the same, then CAS succeeds! So it sets lfstack->head=next.head=0x02! Then,
 //lfstack->head will point to a freed memory! Error!
}



One solution to ABA problem

- Add a 'counter' to indicate how many people have modified it before
- So, even if the content is the same, we can still detect inconsistency

Lock-free vs Wait-free

Lock-free	Wait-free
<ul style="list-style-type: none">• Guarantee progress for some threads,<ul style="list-style-type: none">• Regardless the state/speed of other threads• Guarantee system-wise progress	<ul style="list-style-type: none">• Guarantee progress for every thread,<ul style="list-style-type: none">• Regardless the state/speed of other threads• Guarantee per-thread progress• Sometimes impossible to achieve• Thread1.operation-X() must finish in a finite number of steps<ul style="list-style-type: none">• Regardless of the state/speed of the other thread
E.g., lock-free stack, lock-free hashtable	E.g., The lockless page cache patches to the Linux kernel are an example of a wait-free system.

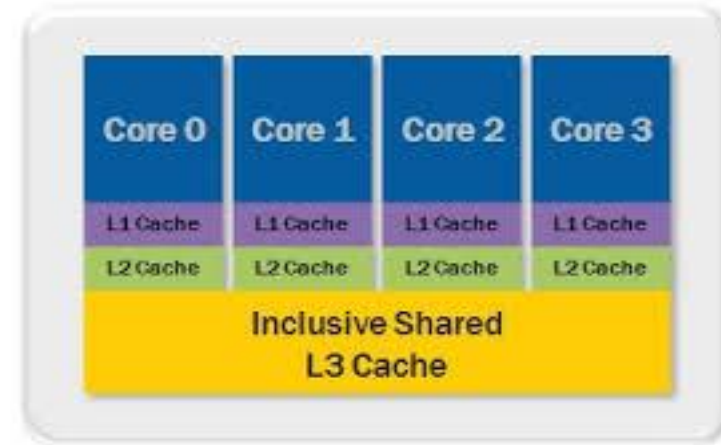
Which one is easier to achieve?

CAS in C11 (programming)

- C11's CAS is not restricted to `int` type only.
 - The more usual one is **pointer-based**. Check out the lab!
 - Lecture: use the `int` version
- C11's CAS has weak and strong version
 - `_Bool atomic_compare_exchange_weak(volatile A *object, C *expected, C desired)`
 - Check out the lab for the meaning of
 - Volatile
 - A: atomic
 - C: non-atomic
 - For our course using x86, no difference

Memory Consistency Model

- Weak and strong refer to memory consistency model
 - At the beginning, Core 1:
 - LOAD memory address 0x888888 (content="old")
 - STORE *0x888888 = "new"
 - Update stills in Core 1's local L1
 - (not yet flush to memory)
 - Then, Core 2:
 - Read memory address 0x888888
 - Whether core 2 reads "old" or "new"
 - Strong consistency memory model (we assume this in the course)
 - x86 (i.e., Intel and AMD) → **cache coherence** → core 2 reads "new"
 - Weak consistency memory model
 - ARM (Advanced RISC Machine) → almost all mobile devices → no cache coherence → core 2 reads "old"
 - To avoid that happens, programmers have to explicit add **memory fence** themselves




Multi-object Synchronization

Inter-process communication (IPC)

- **Classic IPC problems.**
 - Producer-consumer problem
(single object synchronization)
 - Dining philosopher problem
(multi-object synchronization)


Dining philosopher – introduction

- 5 philosophers, 5 plates of spaghetti, and 5 chopsticks.
- The jobs of each philosopher are to think and to eat
- They **need exactly two chopsticks** in order to eat the spaghetti.
- Question: how to construct a synchronization protocol such that they
 - will not **starve to death**, and
 - will not result in any **deadlock scenarios**?
 - A waits for B's chopstick
 - B waits for C's chopstick
 - C waits for A's chopstick




It's a multi-object synchronization problem



Dining philosopher – introduction



Philosophers



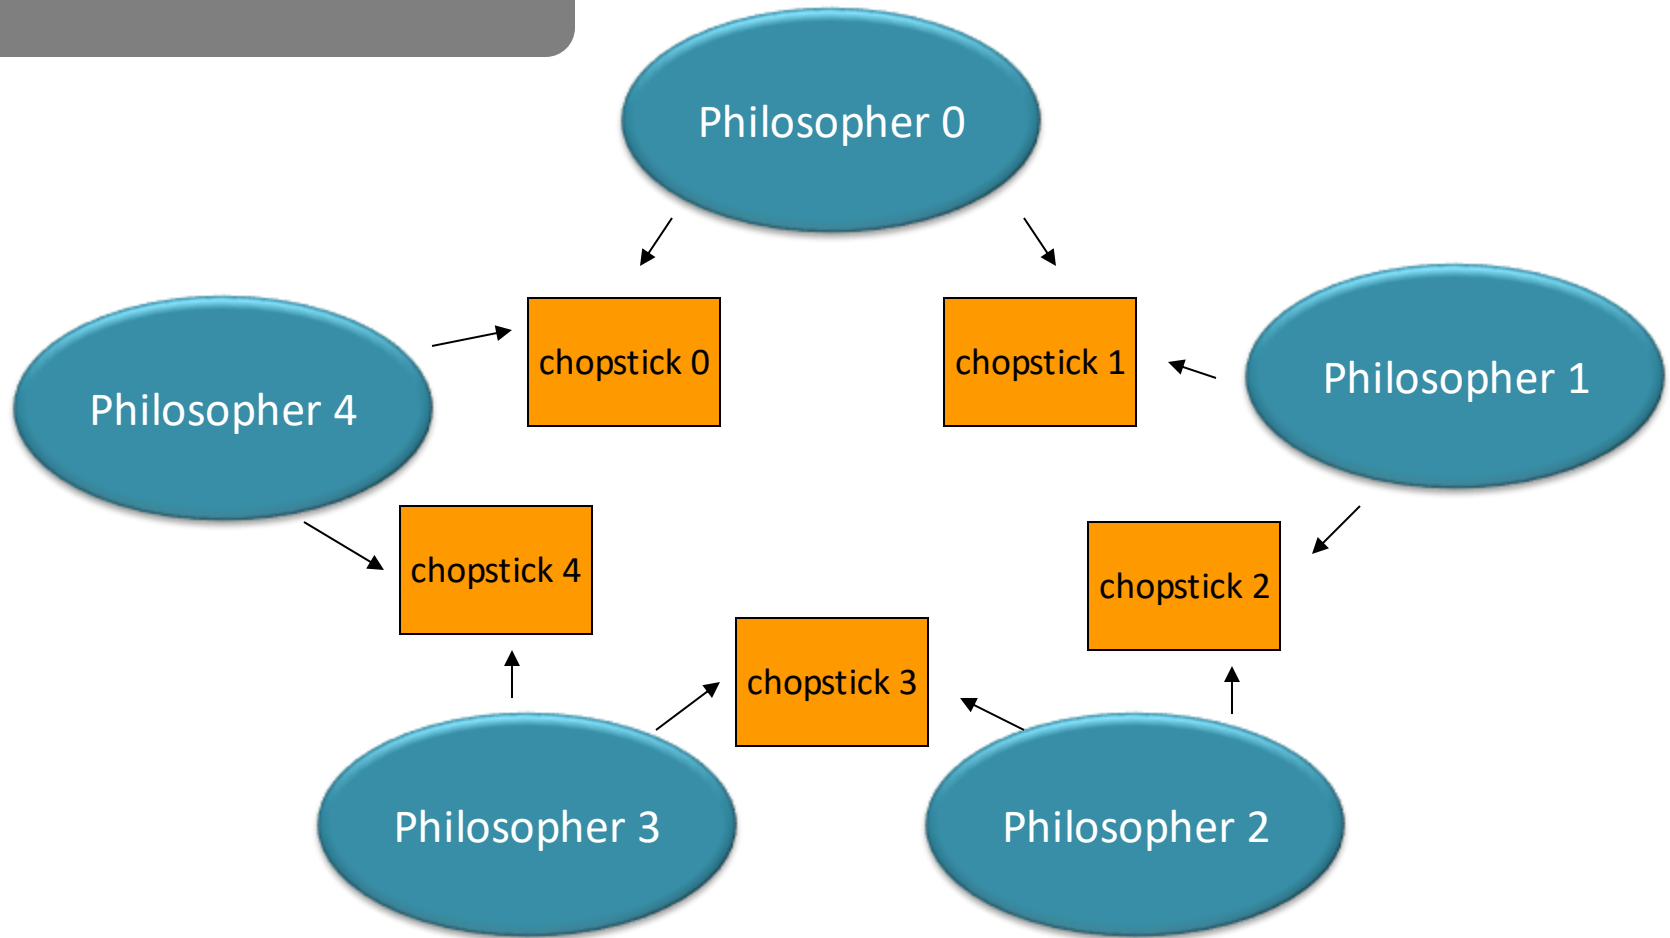
Chopsticks



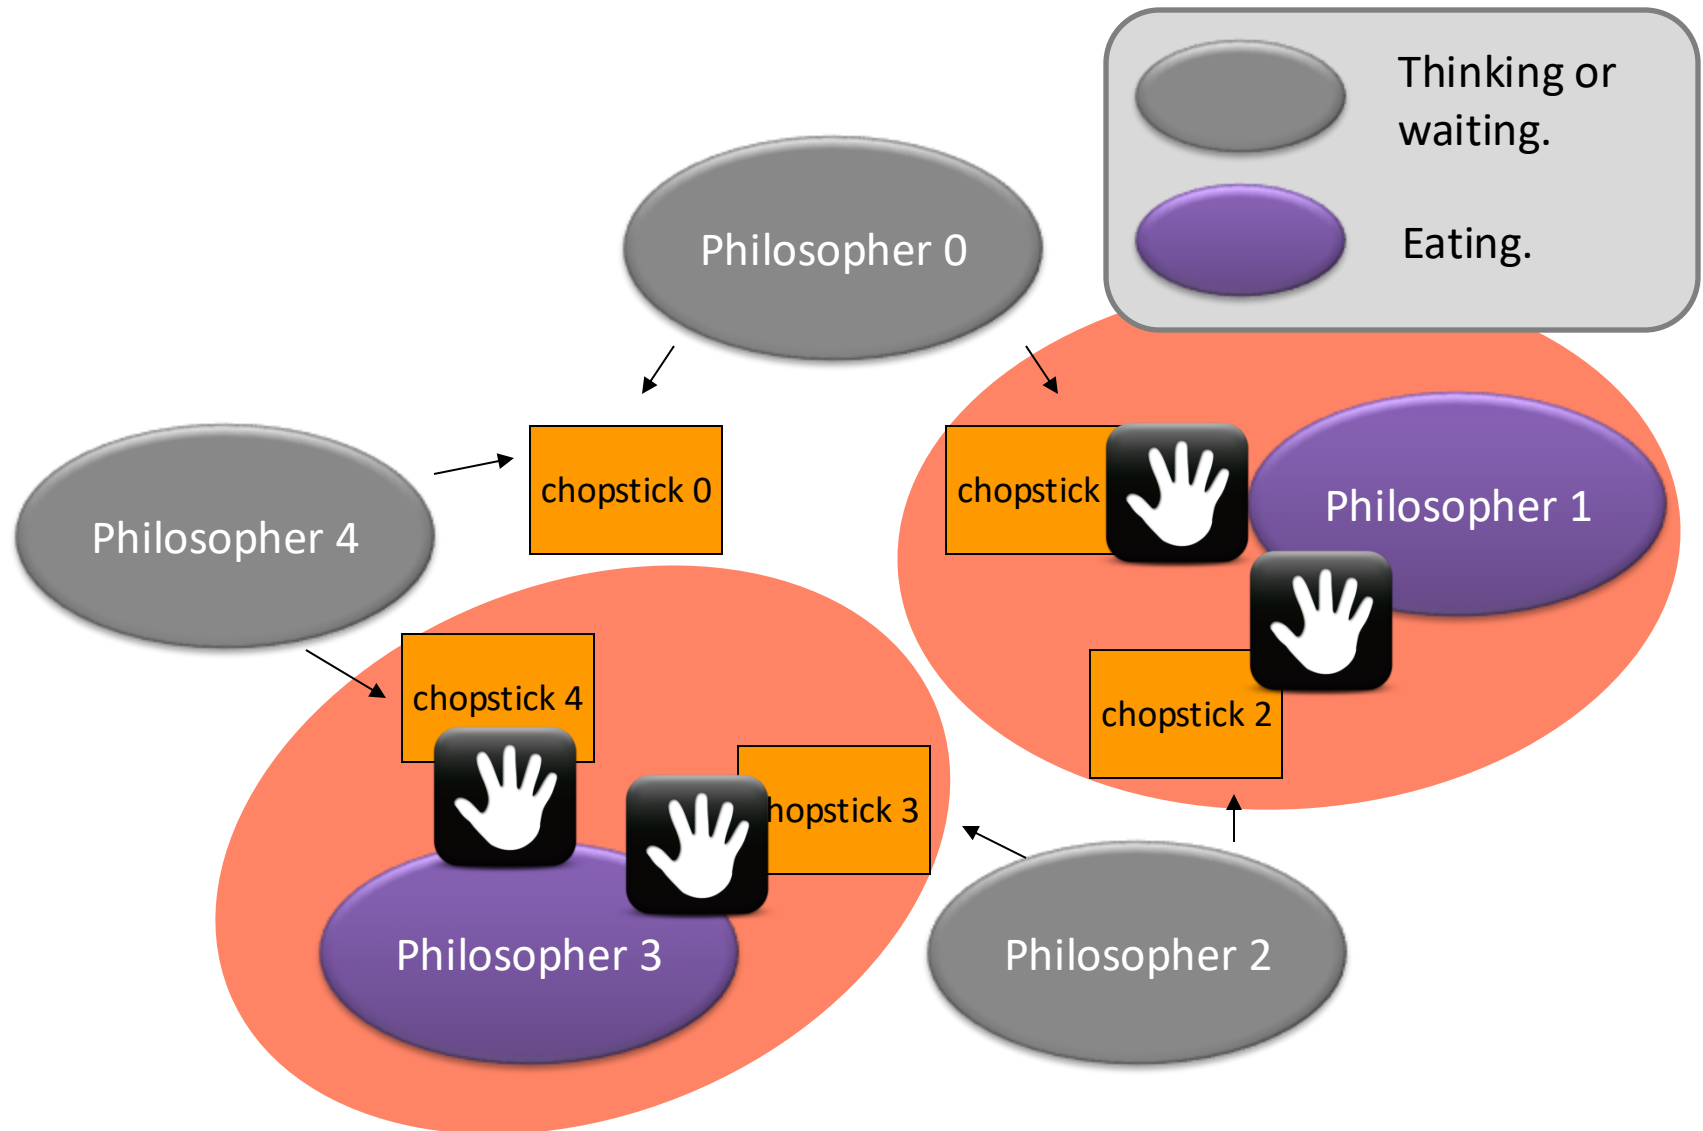
A process needs two shared resources in order to do some work

Dining philosopher – introduction

Imagine the problem as



Dining philosopher – introduction



Dining philosopher – requirement #1

- **Mutual exclusion**
 - While you are eating, people can't steal your chopstick
 - Two persons can't hold the same chopstick
- Let's propose the following solution:
 - When you are hungry, you have to check if anyone is using the chopsticks that you need.
 - If yes, you wait.
 - If no, **seize chopstick-left, chopstick-right.**
 - After eating, put down all your chopsticks.

Dining philosopher – meeting requirement #1?

Shared object

```
#define N 5  
semaphore chopstick[N];
```

Five binary semaphores

Helper Functions

```
void take_chopstick(int i)  
{  
    swait(&chopstick[i]);  
}
```

```
void put_chopstick(int i) {  
    spost(&chopstick[i]);  
}
```

Section
Entry

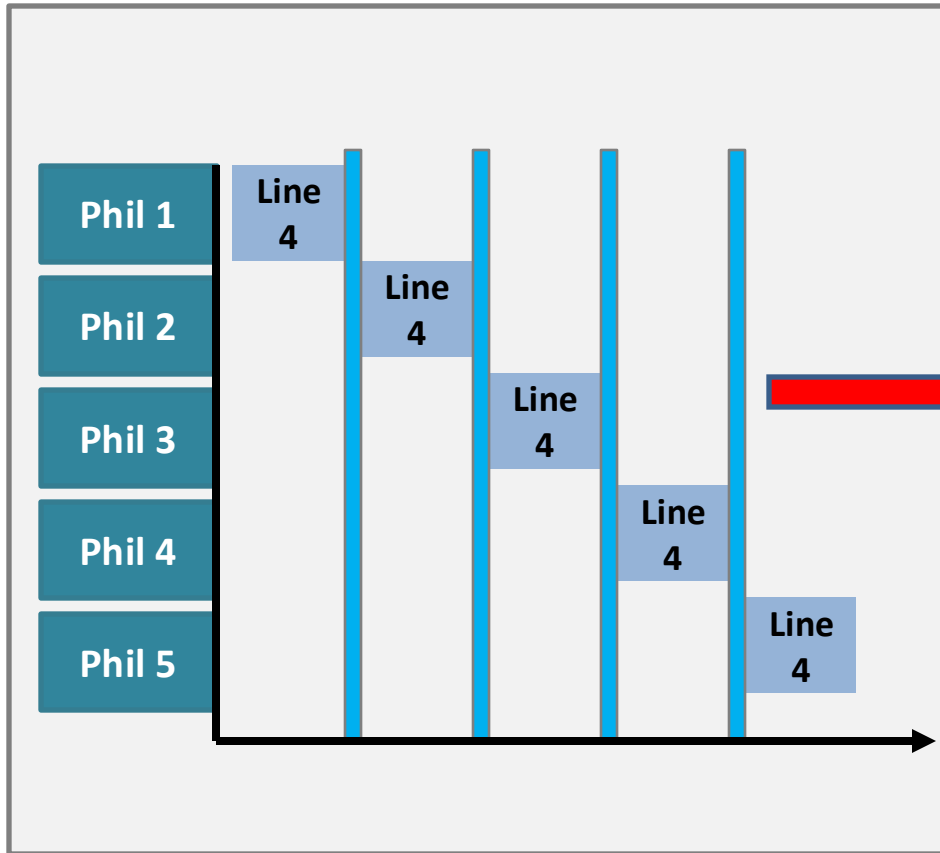
Critical
Section

Section
Exit

Main Function

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take_chopstick(i);  
5         take_chopstick((i+1) % N);  
6         eat();  
7         put_chopstick(i);  
8         put_chopstick((i+1) % N);  
9     }  
10 }
```

Dining philosopher – deadlock



Main Function	
1	void philosopher(int i) {
2	while (TRUE) {
3	think();
4	take_chopstick(i);
5	take_chopstick((i+1) % N);
6	eat();
7	put_chopstick(i);
8	put_chopstick((i+1) % N);
9	}
10	}

Dining philosopher – requirement #2

- **Synchronization**
 - **Deadlock free**
- How about the following suggestions:
 - First, a philosopher **takes a chopstick**.
 - If a philosopher finds that she cannot take the second chopstick, then she should **put it down**.
 - Then, the philosopher **goes to sleep** for a while.
 - When wake up, she retries
 - Loop until both chopsticks are seized.

Dining philosopher – meeting requirement #2?

Potential Problem: Philosophers are all busy
(no deadlock), but no progress
(starvation)

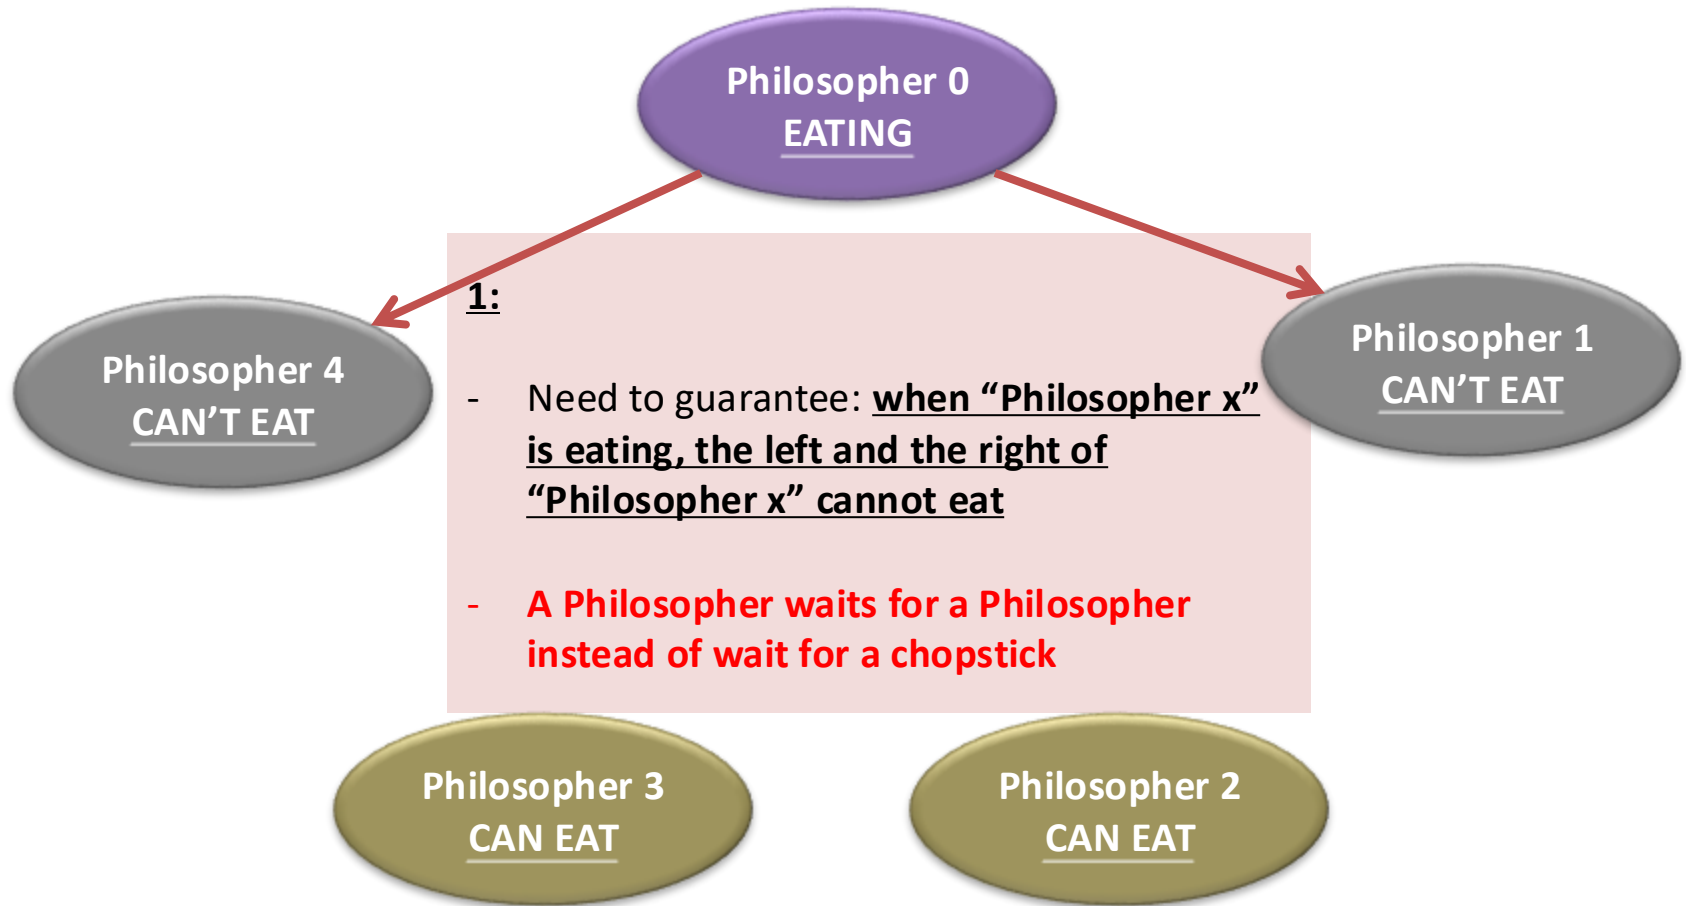
Imagine:

- all pick up their left chopsticks,
- seeing their right chopsticks unavailable (because P1's right chopstick is taken by P2 as her left chopstick) and then putting down their left chopsticks,
- all sleep for a while
- all pick up their left chopsticks,

What's the problem?

- Requirement:
 - Atomic Transaction on **Multiple** Objects
 - All (2 chopsticks) or nothing
- Yet:
 - We took two chopsticks non-atomically

Dining philosopher – one solution idea



2: Add a captain (coordinator) for help

Dining philosopher – a solution.

Shared object

```
#define N 5
#define LEFT  ((i+N-1) % N)
#define RIGHT  ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore p[N] = 0;
```

Identifying i's "left" and "right"

The states of the philosophers, including "EATING", "THINKING", and "HUNGRY".

Remember, this is a shared array.

To guarantee mutual exclusive access to the "**state[N]**" array.

To model a philosopher (instead of a chopstick)

Dining philosopher – a solution.

Shared object

```
#define N 5
#define LEFT ((i+N-1) % N)
#define RIGHT ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore p[N] = 0;
```

Main function

```
1 void philosopher(int i) {
2     think();
3     take_chopsticks(i);
4     eat();
5     put_chopsticks(i);
6 }
```

```
void swait(semaphore *s) {
    disable_interrupt();
    *s = *s - 1;
    if ( *s < 0 ) {
        enable_interrupt();
        sleep();
        disable_interrupt();
    }
    enable_interrupt();
}
```

Section entry

```
1 void take_chopsticks(int i) {
2     swait(&mutex);
3     state[i] = HUNGRY;
4     captain(i);
5     spost(&mutex);
6     swait(&p[i]);
7 }
```

Section exit

```
1 void put_chopsticks(int i) {
2     swait(&mutex);
3     state[i] = THINKING;
4     captain(LEFT);
5     captain(RIGHT);
6     spost(&mutex);
7 }
```

```
void spost(semaphore *s) {
    disable_interrupt();
    *s = *s + 1;
    if ( *s <= 0 )
        wakeup();
    enable_interrupt();
}
```

Extremely important helper function

```
1 void captain(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         spost(&p[i]);
5     }
6 }
```

Dining philosopher – Hungry

Tell the captain that you are hungry

If one of your neighbors is eating, the captain just does nothing for you and returns

Then, you **wait** for your chopsticks (**later, the captain will notify you when chopsticks are available**)

Section entry

```
1 void take_chopsticks(int i) {  
2     swait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     spost(&mutex);  
6     swait(&p[i]);  
7 }
```

Critical Section

Extremely important helper function

```
1 void captain(int i) { //atomic captain  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         spost(&p[i]); //s.t. p[i] don't need to wait in line 6 of take_chopsticks  
5     }  
6 }
```

Dining philosopher – Finish eating

Tell the captain:
Try to let your **left neighbor** to eat.

Tell the captain:
Try to let your **right neighbor** to eat.

Section exit

```
1 void put_chopsticks(int i)
{
2     swait(&mutex);
3     state[i] = THINKING;
4     captain(LEFT);
5     captain(RIGHT);
6     spost(&mutex);
7 }
```

Extremely important helper function

```
1 void captain(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         spost(&p[i]);
5     }
6 }
```

Wake up the one who is waiting

Dining philosopher – the final solution.

Don't print

An illustration: How can
Philosopher 1 start eating?

Philosopher 0
THINKING

Philosopher 1
THINKING

Philosopher 4
THINKING

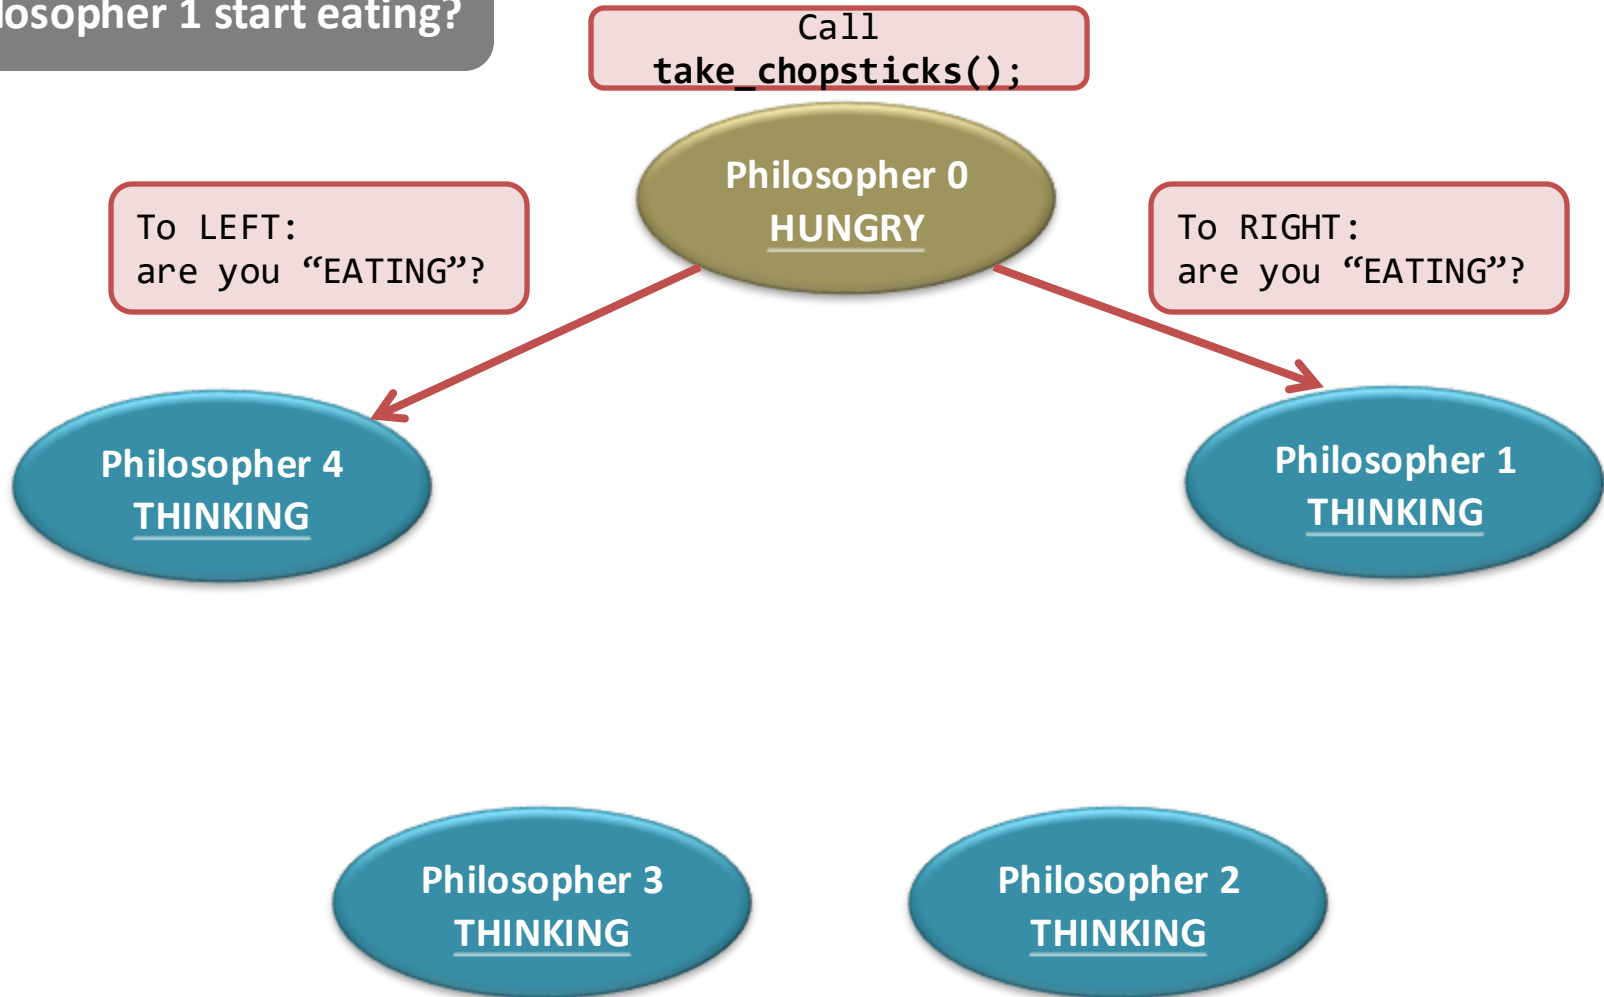
Philosopher 3
THINKING

Philosopher 2
THINKING

Dining philosopher – the final solution.

Don't print

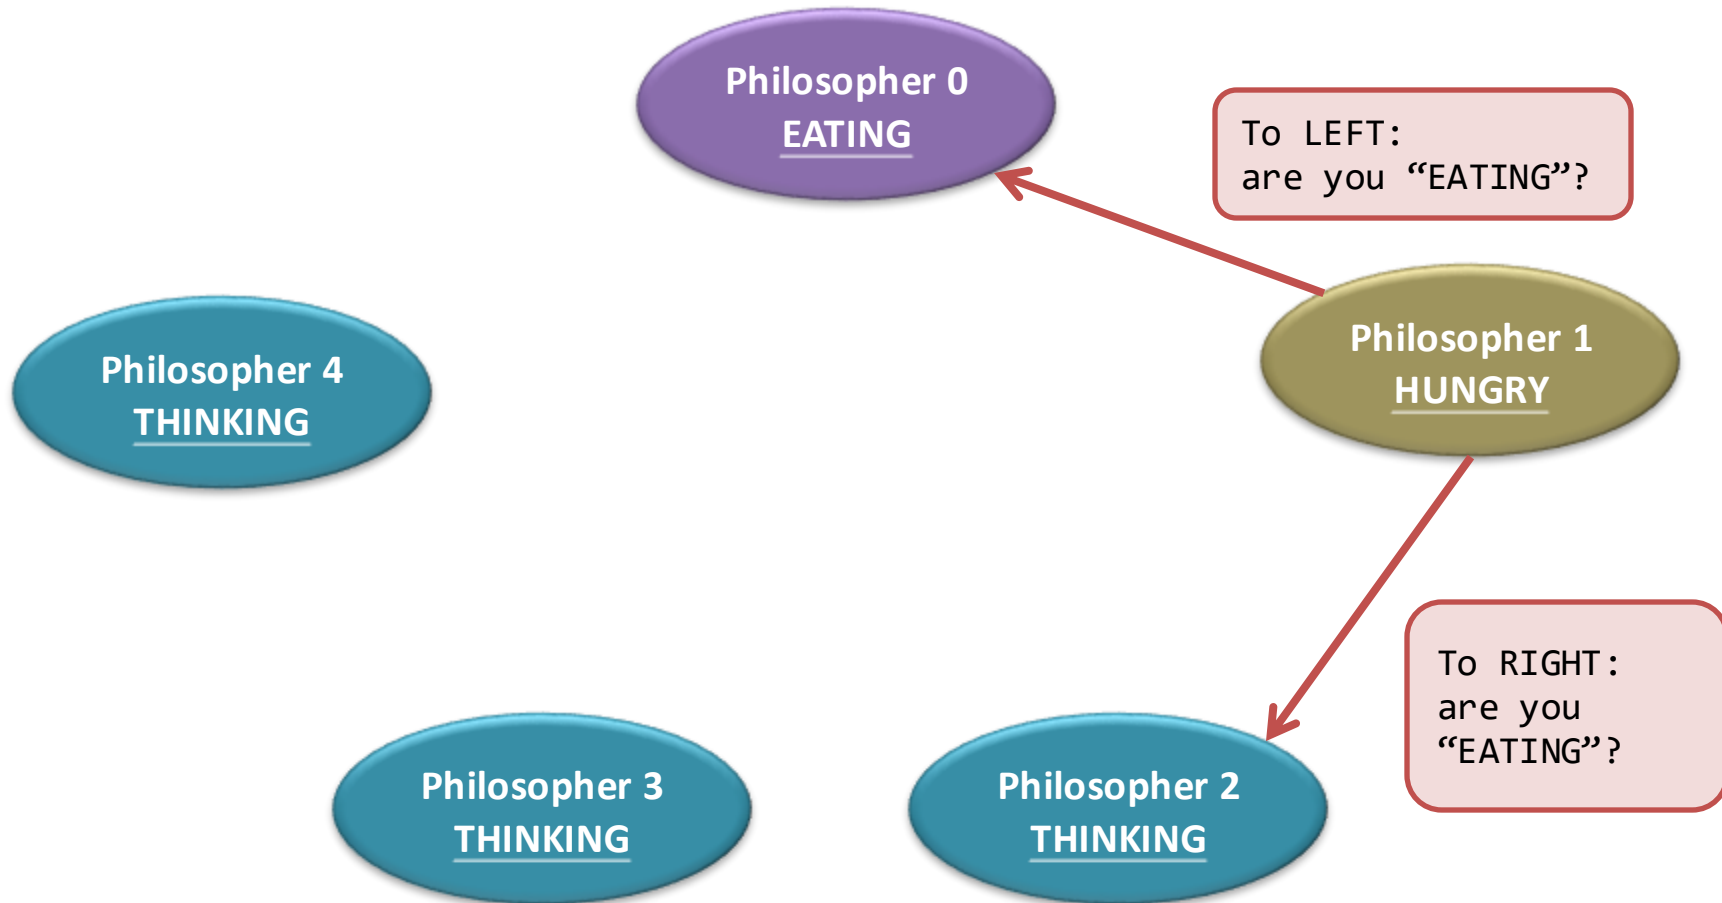
An illustration: How can
Philosopher 1 start eating?



Dining philosopher – the final solution.

Don't print

An illustration: How can
Philosopher 1 start eating?



Dining philosopher – the final solution.

Don't print

An illustration: How can Philosopher 1 start eating?

Philosopher 0
EATING

Philosopher 4
THINKING

To LEFT:
are you
"EATING"?

Philosopher 3
HUNGRY

To RIGHT:
are you
"EATING"?

Philosopher 2
THINKING

Section entry

```
1 void take_chopsticks(int i) {  
2     swait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     spost(&mutex);  
6     swait(&p[i]);  
7 }
```

//as P0 is eating, captain(i) returns
w/o doing anything;
swait(&p[1]);

Philosopher 1
HUNGRY

Dining philosopher – the final solution.

Don't print

An illustration: How can
Philosopher 1 start eating?

Philosopher 0
EATING

Philosopher 4
THINKING

Philosopher 1
HUNGRY

Blocked;
because of
`swait(&p[1]);`

Philosopher 3
EATING

Philosopher 2
THINKING

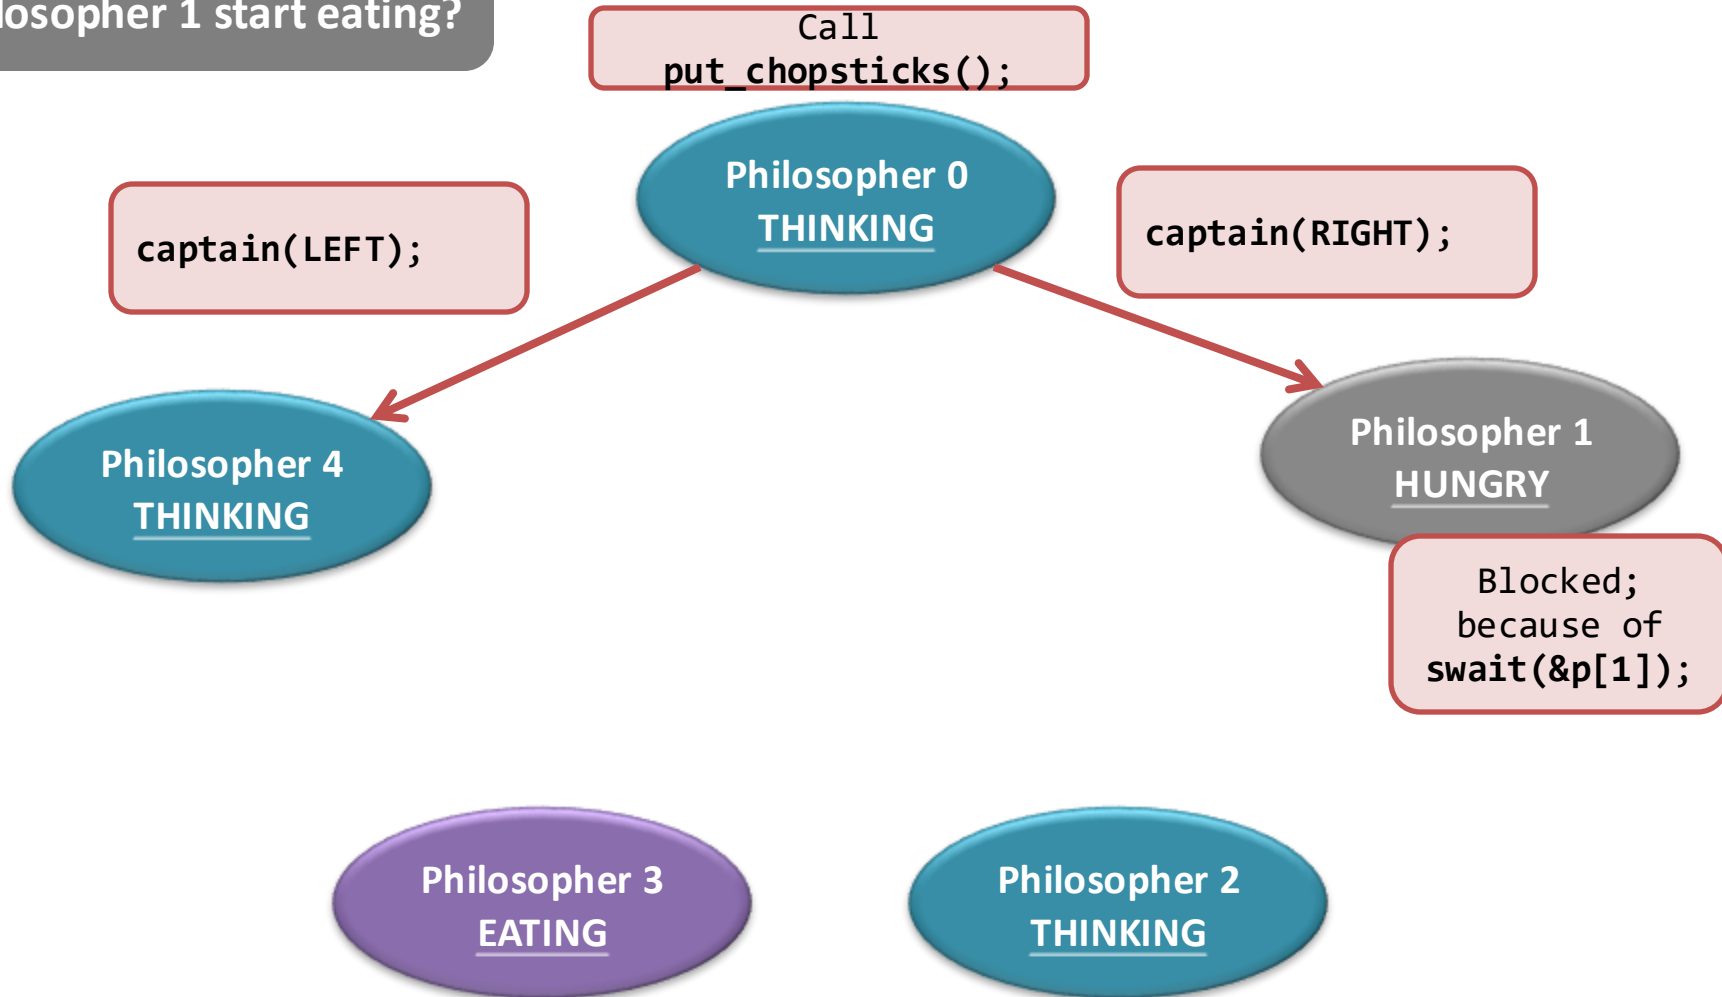
Section entry

```
1 void take_chopsticks(int i) {  
2     swait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     spost(&mutex);  
6     swait(&p[i]);  
7 }
```

Dining philosopher – the final solution.

Don't print

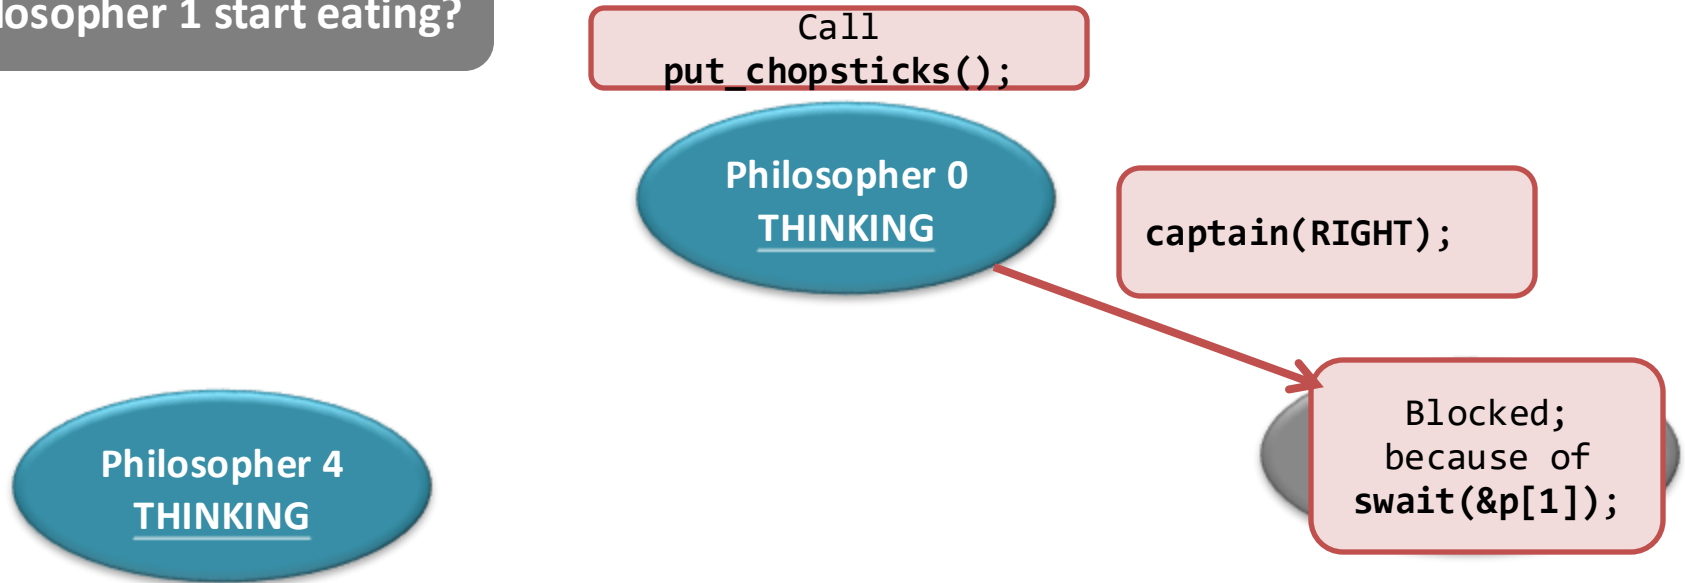
An illustration: How can
Philosopher 1 start eating?



Dining philosopher – the final solution.

Don't print

An illustration: How can Philosopher 1 start eating?



```
1 void captain(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         spost(&p[i]); ← Wake up !  
5     }  
6 }
```

Dining philosopher – the final solution.

Don't print

An illustration: How can
Philosopher 1 start eating?

Philosopher 0
THINKING

Philosopher 4
THINKING

Philosopher 3
EATING

Philosopher 2
THINKING

Section entry

```
1 void take_chopsticks(int i) {  
2     swait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     spost(&mutex);  
6     swait(&p[i]);  
7 }
```

Wake up

Philosopher 1
EATING

Dining philosopher – the core

5 philosophers → ideally how many chopsticks?

how many chopsticks do we have now?

Very common in today's cloud computing multi-tenancy model

Dining philosopher

- Atomic Transaction on **Multiple** Objects
- It is a fundamental problem in database area
 - They have better solutions for it
 - E.g., OCC